

# Lecture 29: Machine Learning

---

Marvin Zhang  
08/10/2015

(Some images borrowed from CS 188.)

# Announcements

---

# Announcements

---

- Project 3 composition revisions due Wednesday night.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.
  - Review session tomorrow, 5-8pm in the Woz.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.
  - Review session tomorrow, 5-8pm in the Woz.
- Tuesday, Wednesday, and Thursday sections canceled.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.
  - Review session tomorrow, 5-8pm in the Woz.
- Tuesday, Wednesday, and Thursday sections canceled.
  - Instead, TAs will lead topic-themed discussions Tuesday and Wednesday.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.
  - Review session tomorrow, 5-8pm in the Woz.
- Tuesday, Wednesday, and Thursday sections canceled.
  - Instead, TAs will lead topic-themed discussions Tuesday and Wednesday.
  - Details will be posted on Piazza.

# Announcements

---

- Project 3 composition revisions due Wednesday night.
- Project 4 due tonight (Monday).
  - Office hours from 4-7pm today in the Woz.
- Homework 11 due tonight - just a survey!
- Project 4 contest due tomorrow (Tuesday) night.
  - Top 3 entries in each category get extra credit! Only one entry so far.
- Final on Thursday, 3-6pm in 2050 VLSB.
  - Review session tomorrow, 5-8pm in the Woz.
- Tuesday, Wednesday, and Thursday sections canceled.
  - Instead, TAs will lead topic-themed discussions Tuesday and Wednesday.
  - Details will be posted on Piazza.
  - Chris and Cale's 9:30-11am labs on Tuesday are NOT canceled.

# What is Machine Learning?

---

# What is Machine Learning?

---

- Natural Language Processing

# What is Machine Learning?

---

- Natural Language Processing



# What is Machine Learning?

---

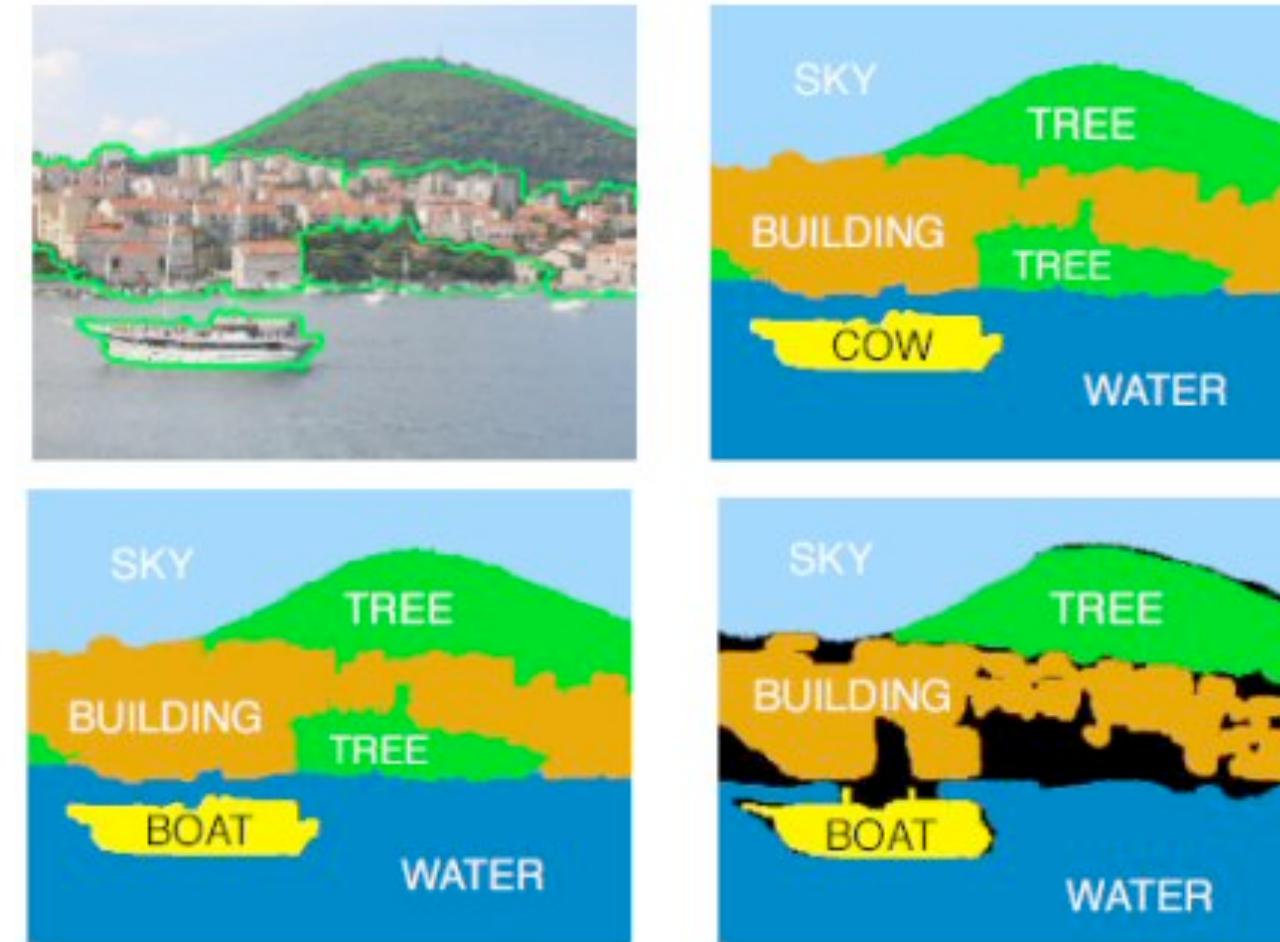
- Natural Language Processing
- Computer Vision



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision
- Robotics



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision
- Robotics



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision
- Robotics



# What is Machine Learning?

---

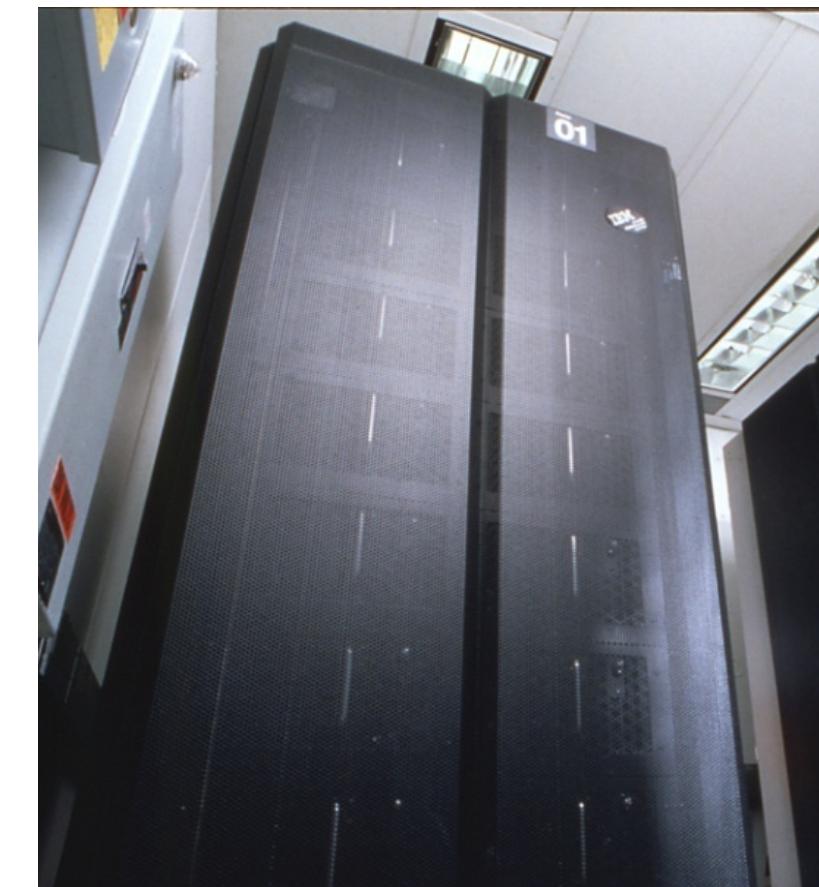
- Natural Language Processing
- Computer Vision
- Robotics
- Game Playing



# What is Machine Learning?

---

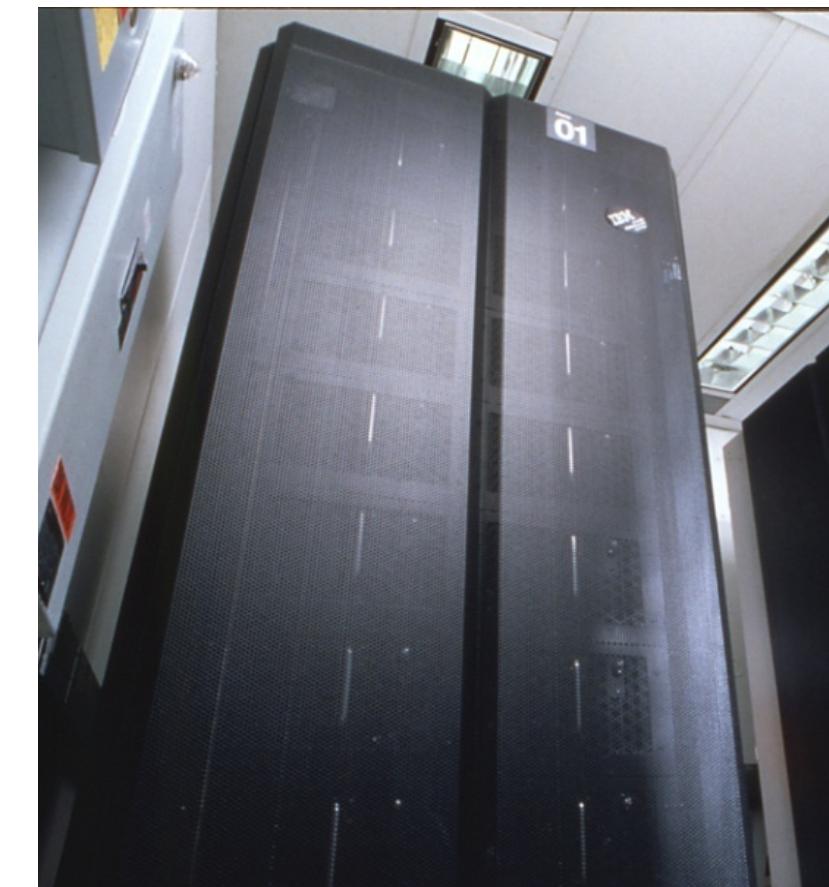
- Natural Language Processing
- Computer Vision
- Robotics
- Game Playing



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision
- Robotics
- Game Playing
- and much more!



# What is Machine Learning?

---

- Natural Language Processing
- Computer Vision
- Robotics
- Game Playing
- and much more!



- What do these all have in common?

# Machine Learning

---

# Machine Learning

---

What is machine learning?

# Machine Learning

---

What is machine learning?

- A subfield of computer science.

# Machine Learning

---

What is machine learning?

- A subfield of computer science.
- The study of algorithms that *analyze data to make decisions*.

# Machine Learning

---

What is machine learning?

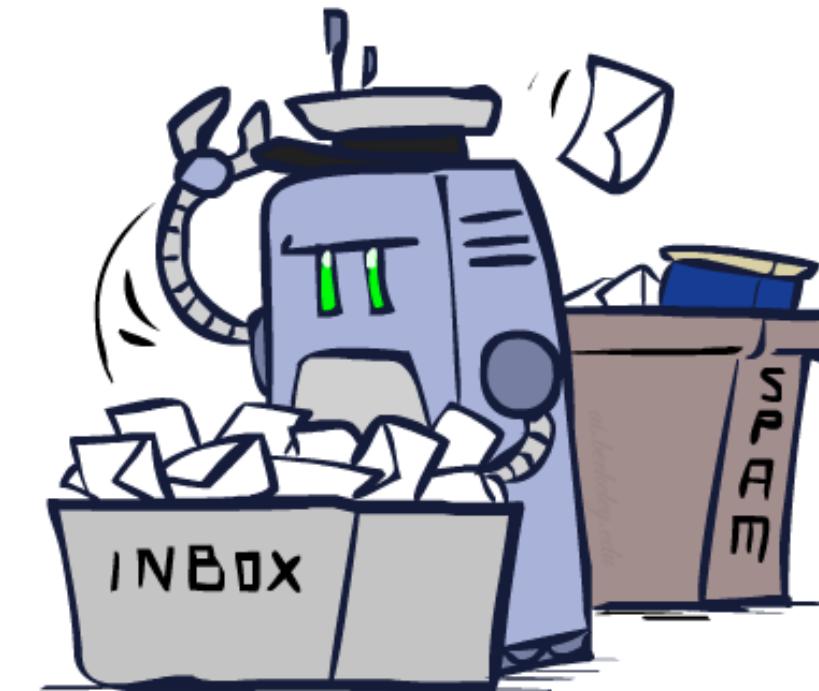
- A subfield of computer science.
- The study of algorithms that *analyze data to make decisions*.
- Examples of decisions:

# Machine Learning

---

What is machine learning?

- A subfield of computer science.
- The study of algorithms that *analyze data to make decisions*.
- Examples of decisions:
  - Is this email ham or spam?



# Machine Learning

---

What is machine learning?

- A subfield of computer science.
- The study of algorithms that *analyze data to make decisions*.
- Examples of decisions:
  - Is this email ham or spam?
  - How do I translate this sentence?

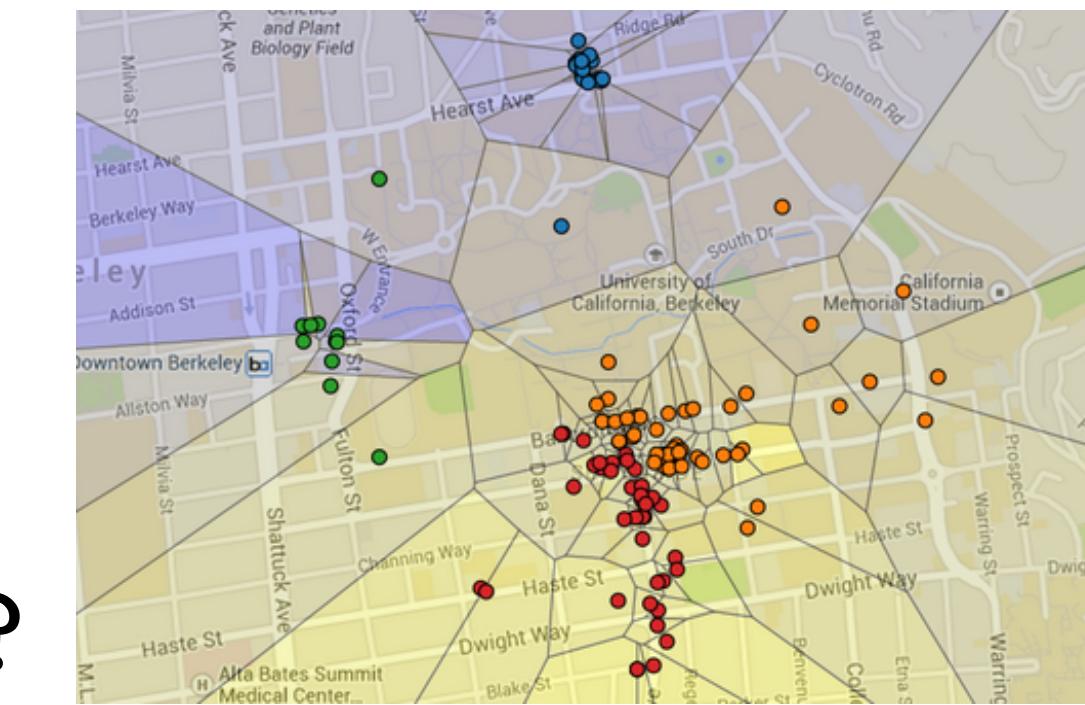


# Machine Learning

---

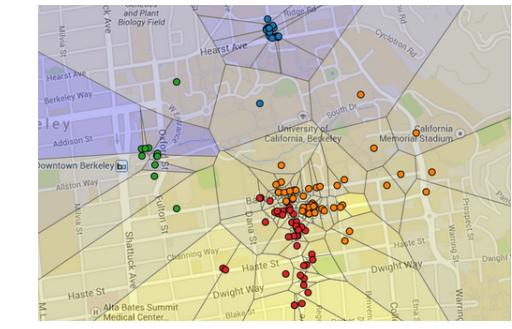
What is machine learning?

- A subfield of computer science.
- The study of algorithms that *analyze data to make decisions*.
- Examples of decisions:
  - Is this email ham or spam?
  - How do I translate this sentence?
  - Will this user like this restaurant?



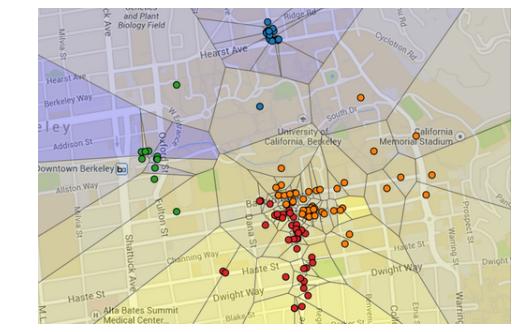
# Machine Learning Example: Maps

---



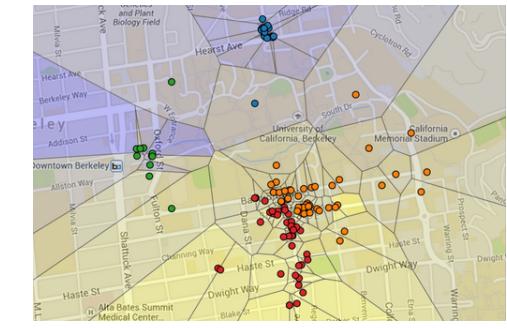
# Machine Learning Example: Maps

---

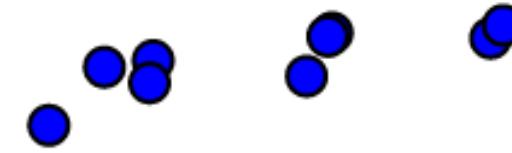
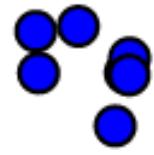
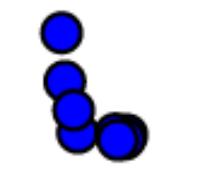


## K-means Clustering

# Machine Learning Example: Maps

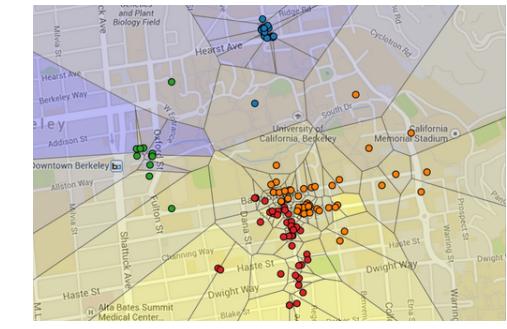


# K-means Clustering



# Machine Learning Example: Maps

---

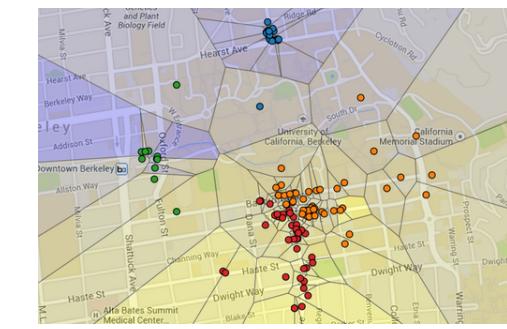


## K-means Clustering



# Machine Learning Example: Maps

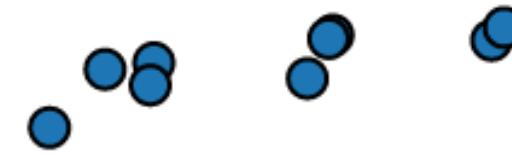
---



## K-means Clustering

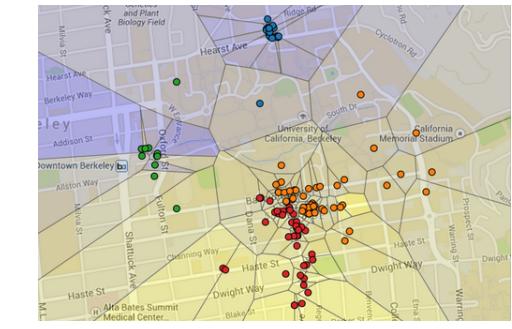


- The data: restaurant locations



# Machine Learning Example: Maps

---



## K-means Clustering

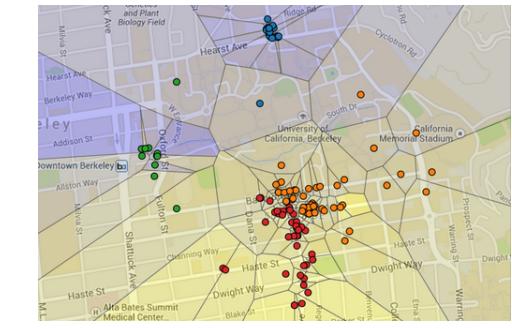


- The data: restaurant locations
- The decision: which cluster does each belong to?



# Machine Learning Example: Maps

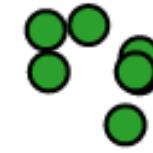
---



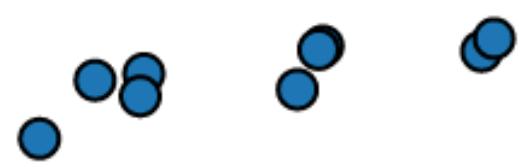
## K-means Clustering



- The data: restaurant locations
- The decision: which cluster does each belong to?

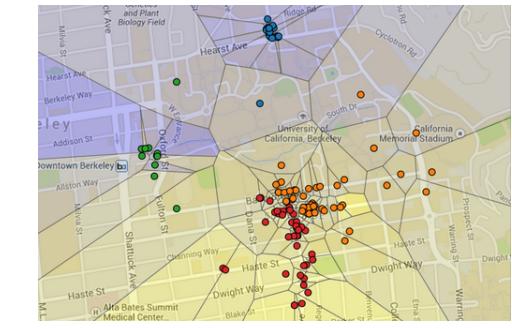


Called *unsupervised Learning*, because no one tells it what the correct decision is.



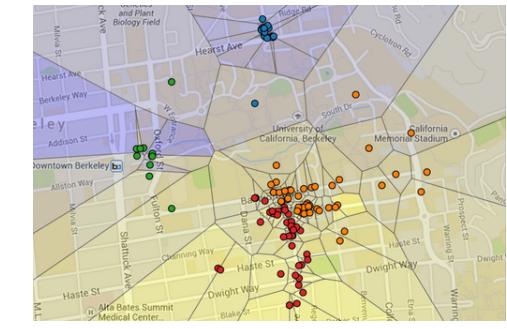
# Machine Learning Example: Maps

---



# Machine Learning Example: Maps

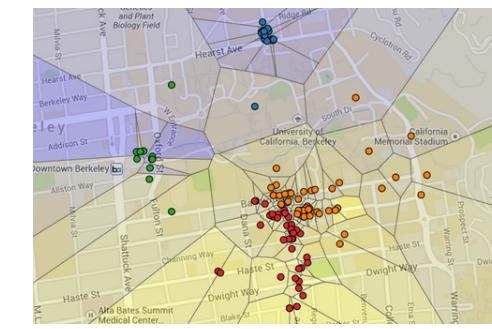
---



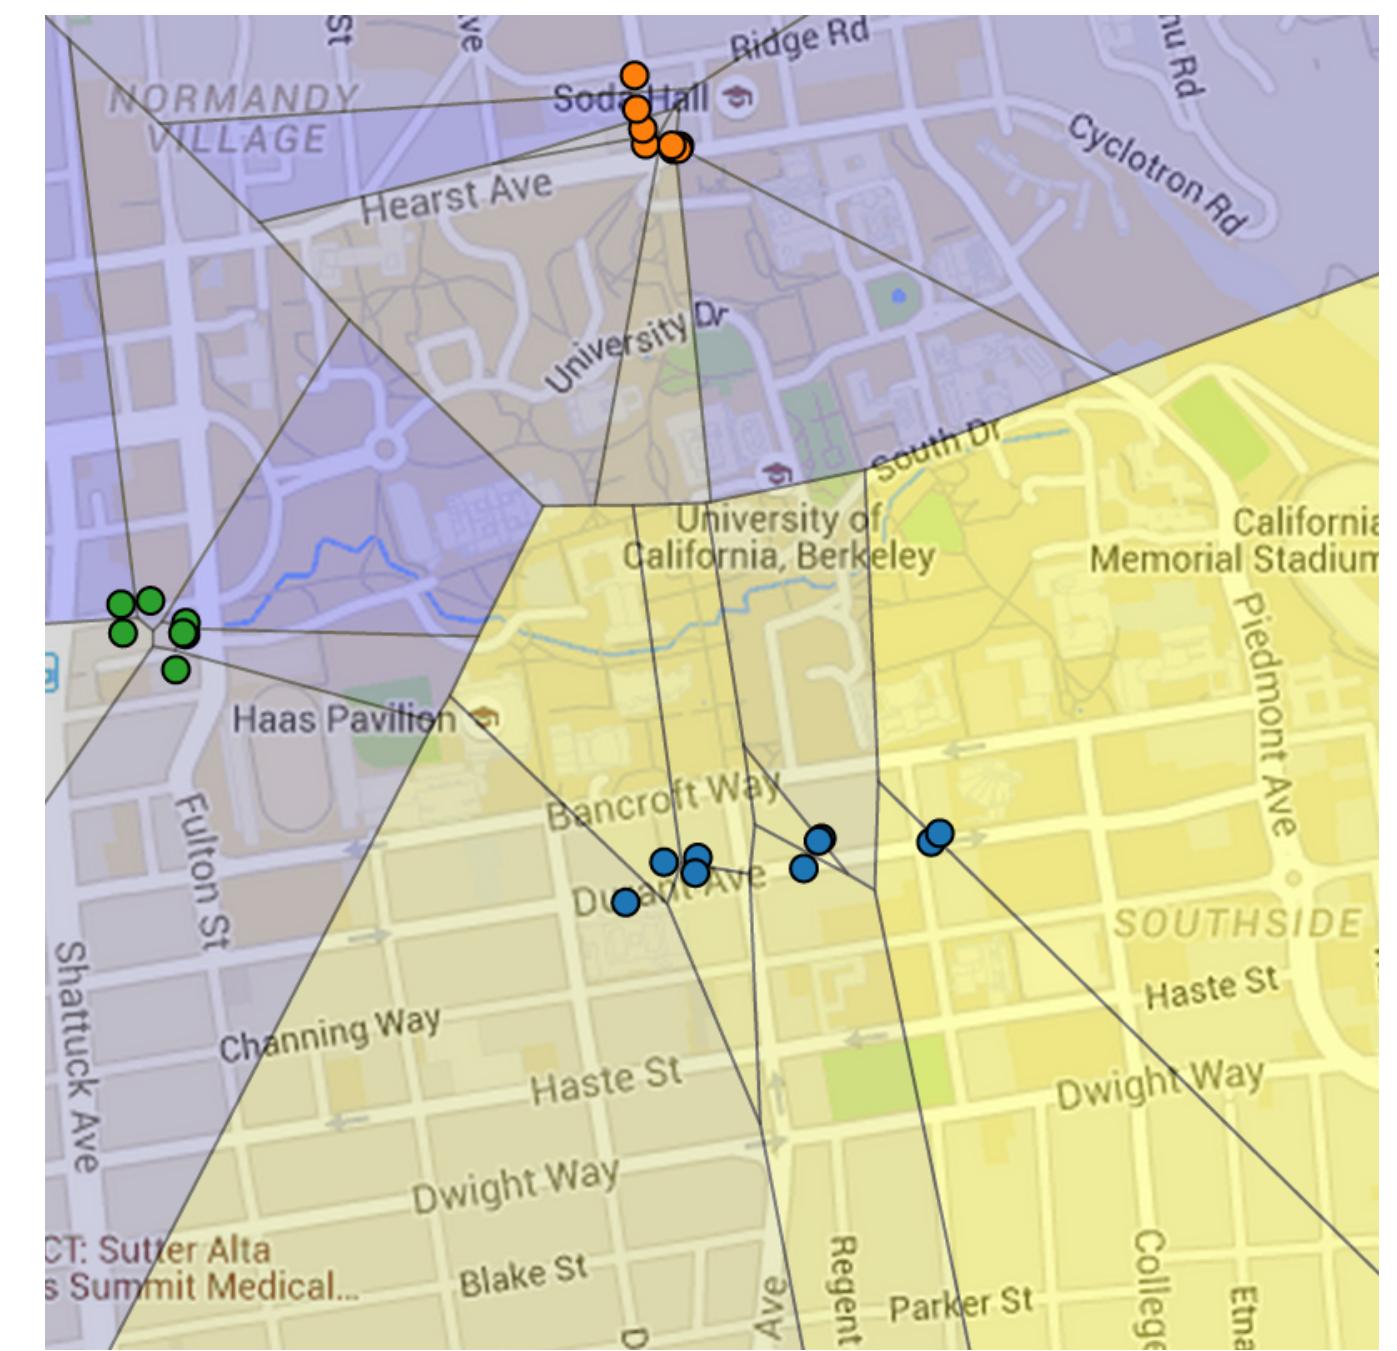
## Linear Regression

# Machine Learning Example: Maps

---

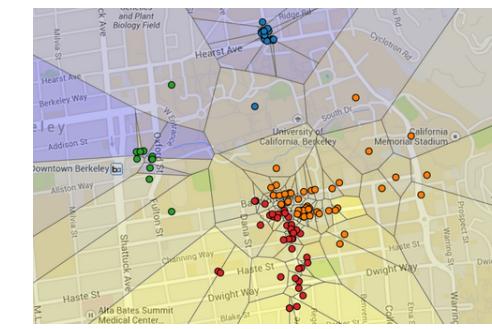


## Linear Regression

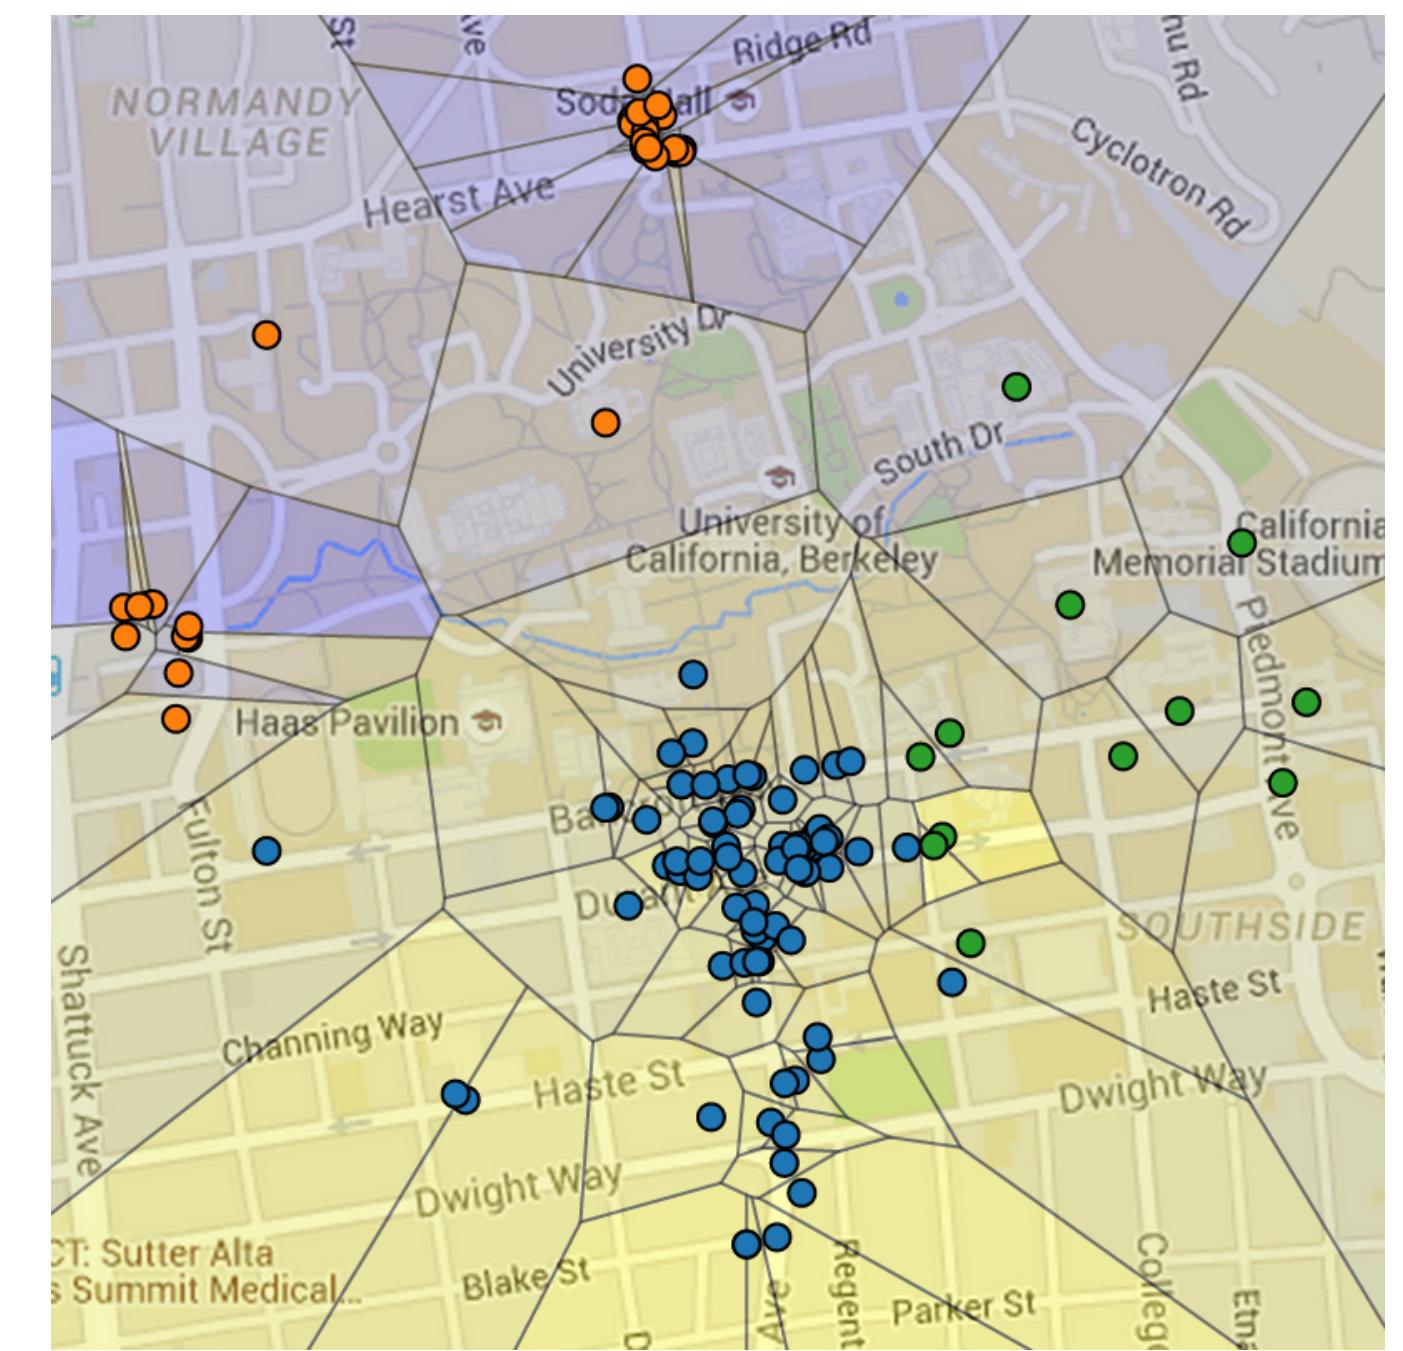


# Machine Learning Example: Maps

---

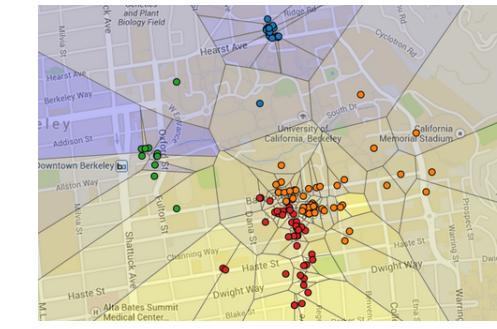


## Linear Regression



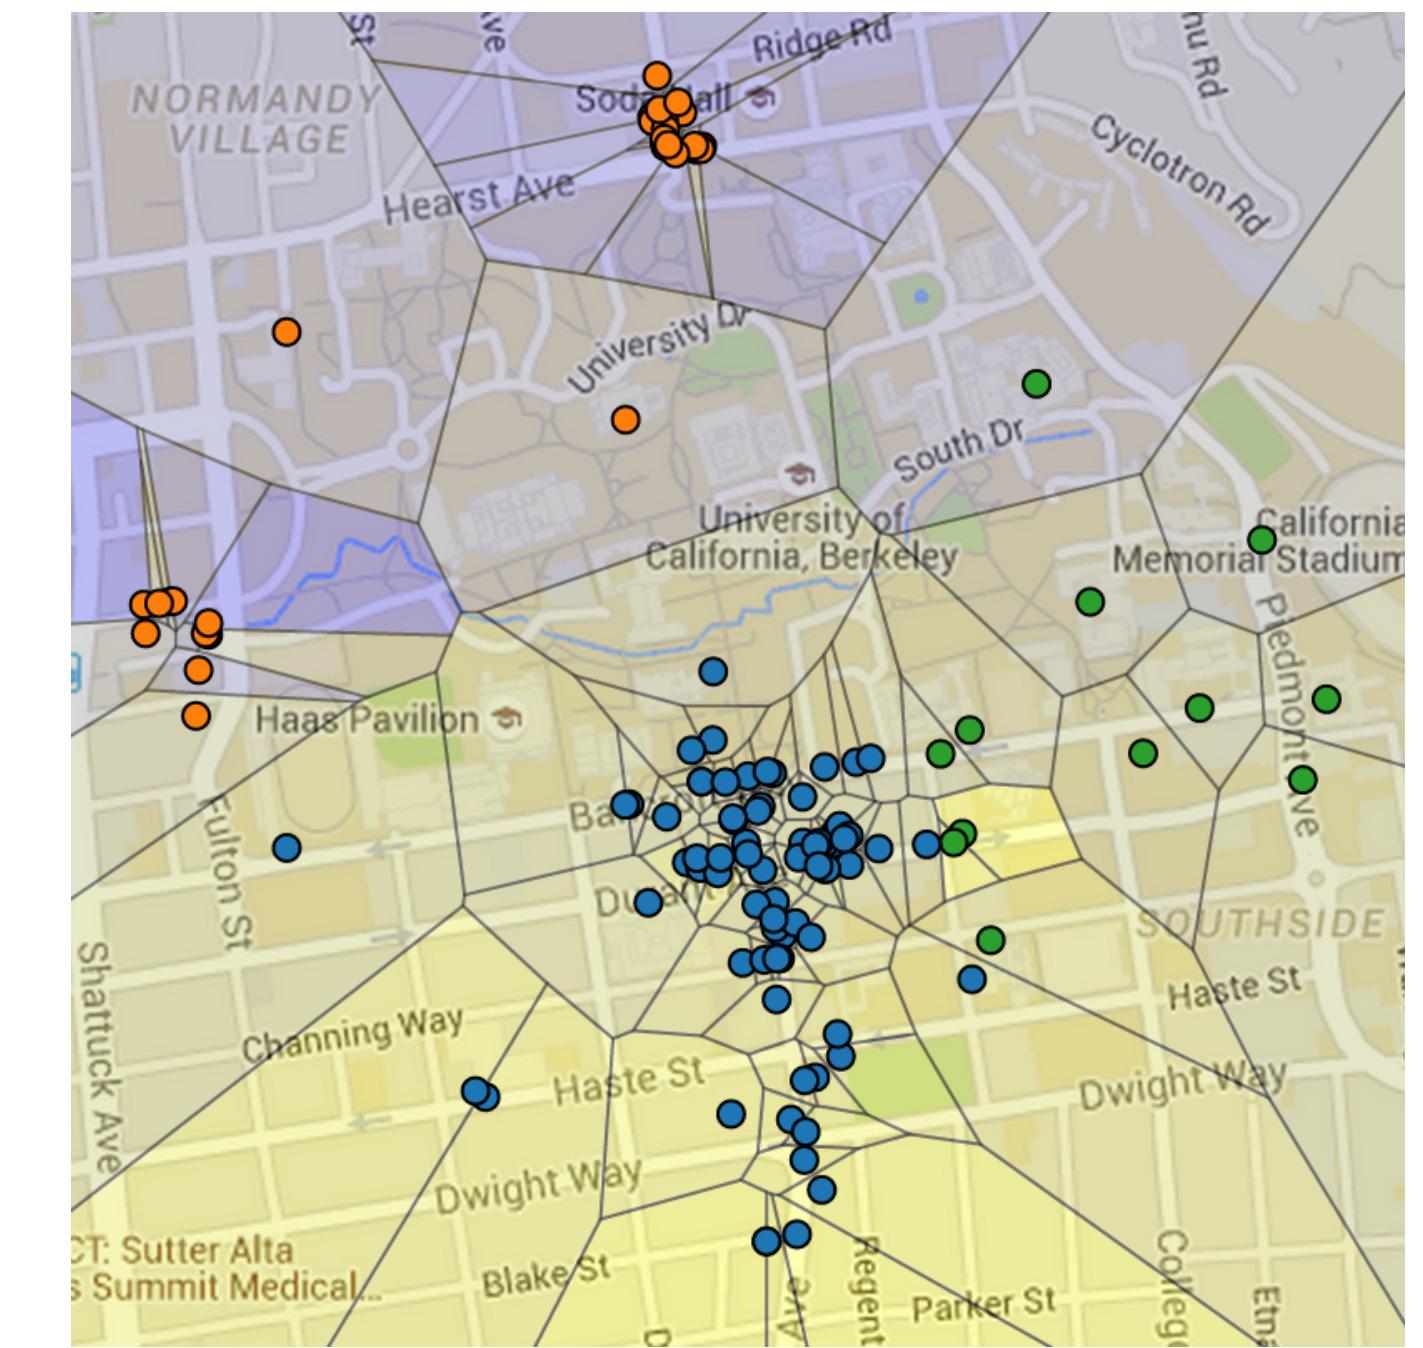
# Machine Learning Example: Maps

---



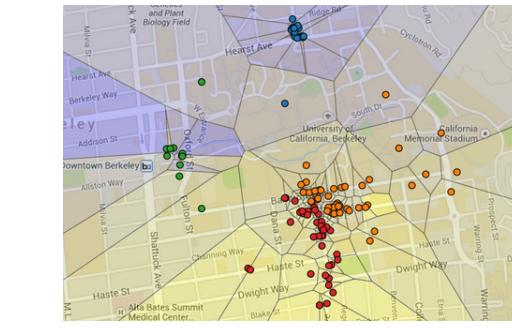
## Linear Regression

- The data: user ratings



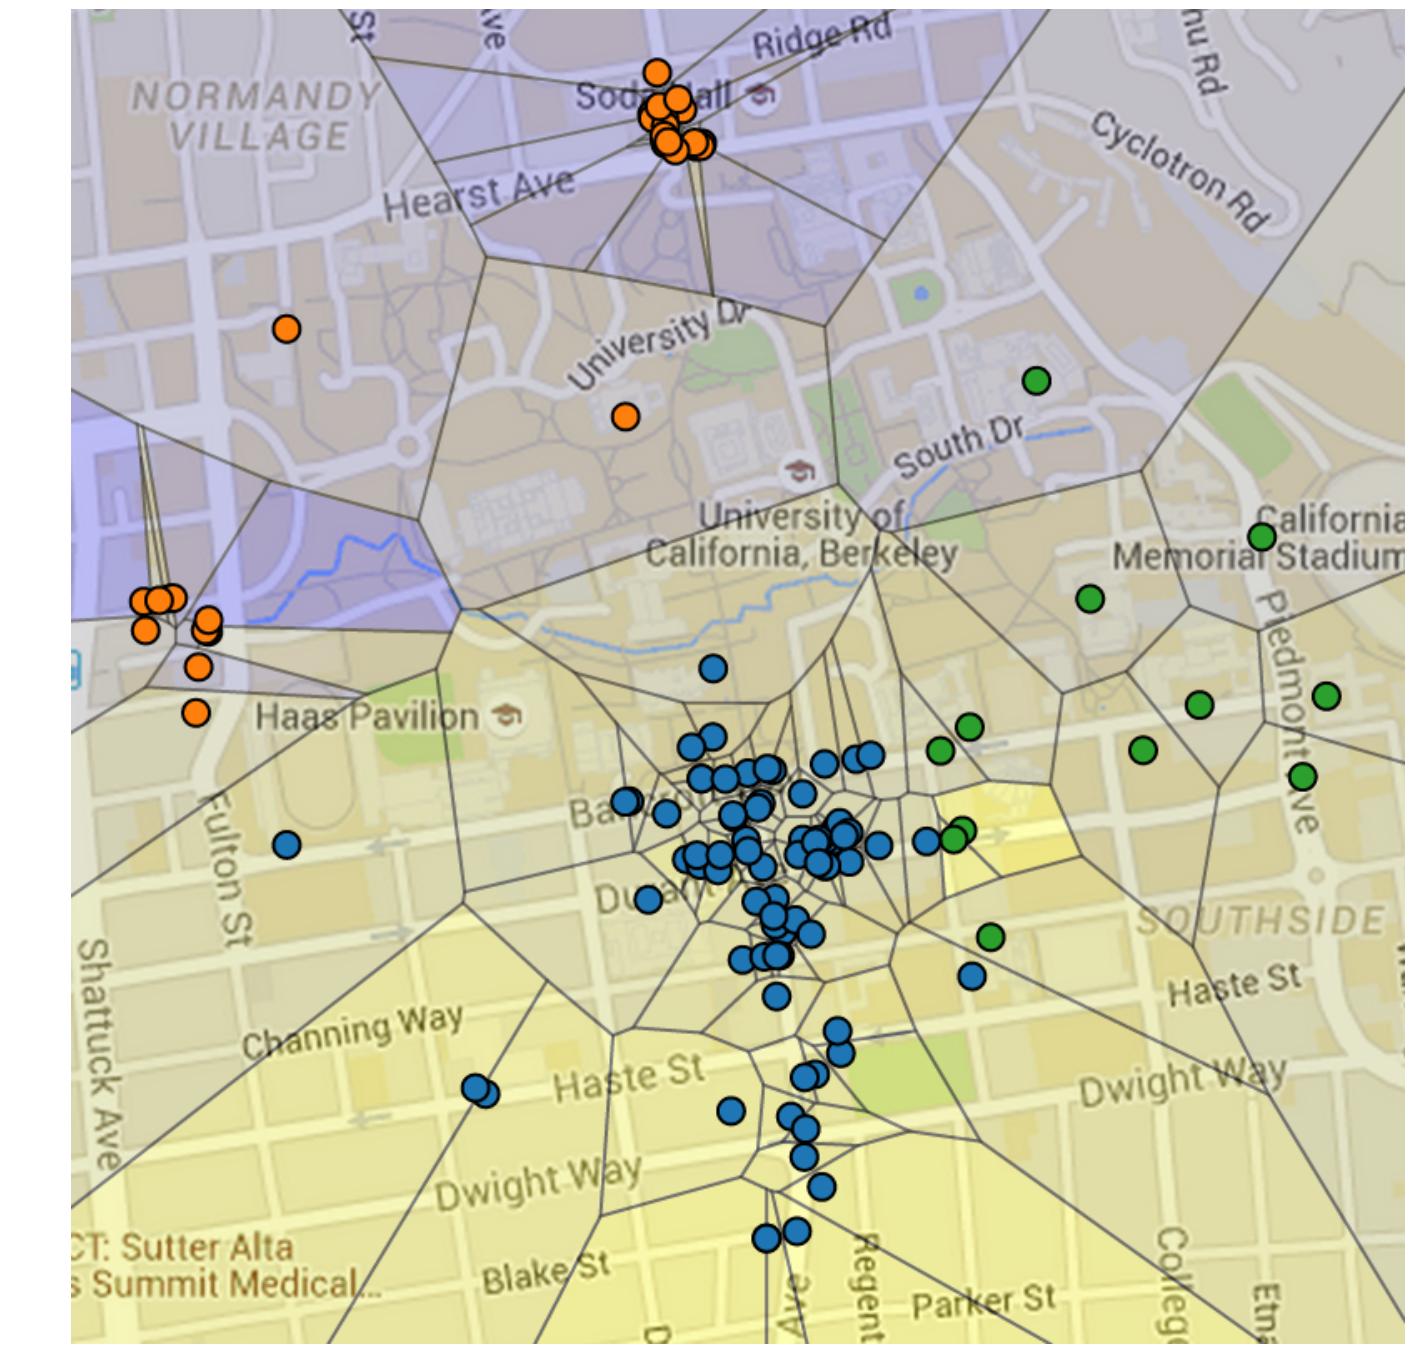
# Machine Learning Example: Maps

---



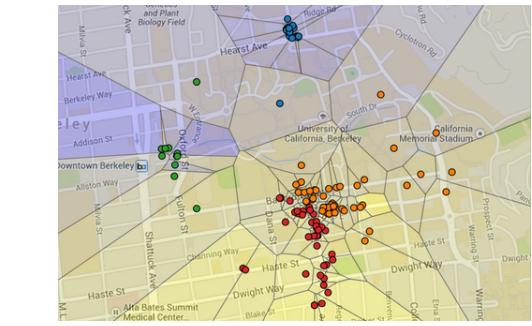
## Linear Regression

- The data: user ratings
- The decision: what rating would the user give a new restaurant?



# Machine Learning Example: Maps

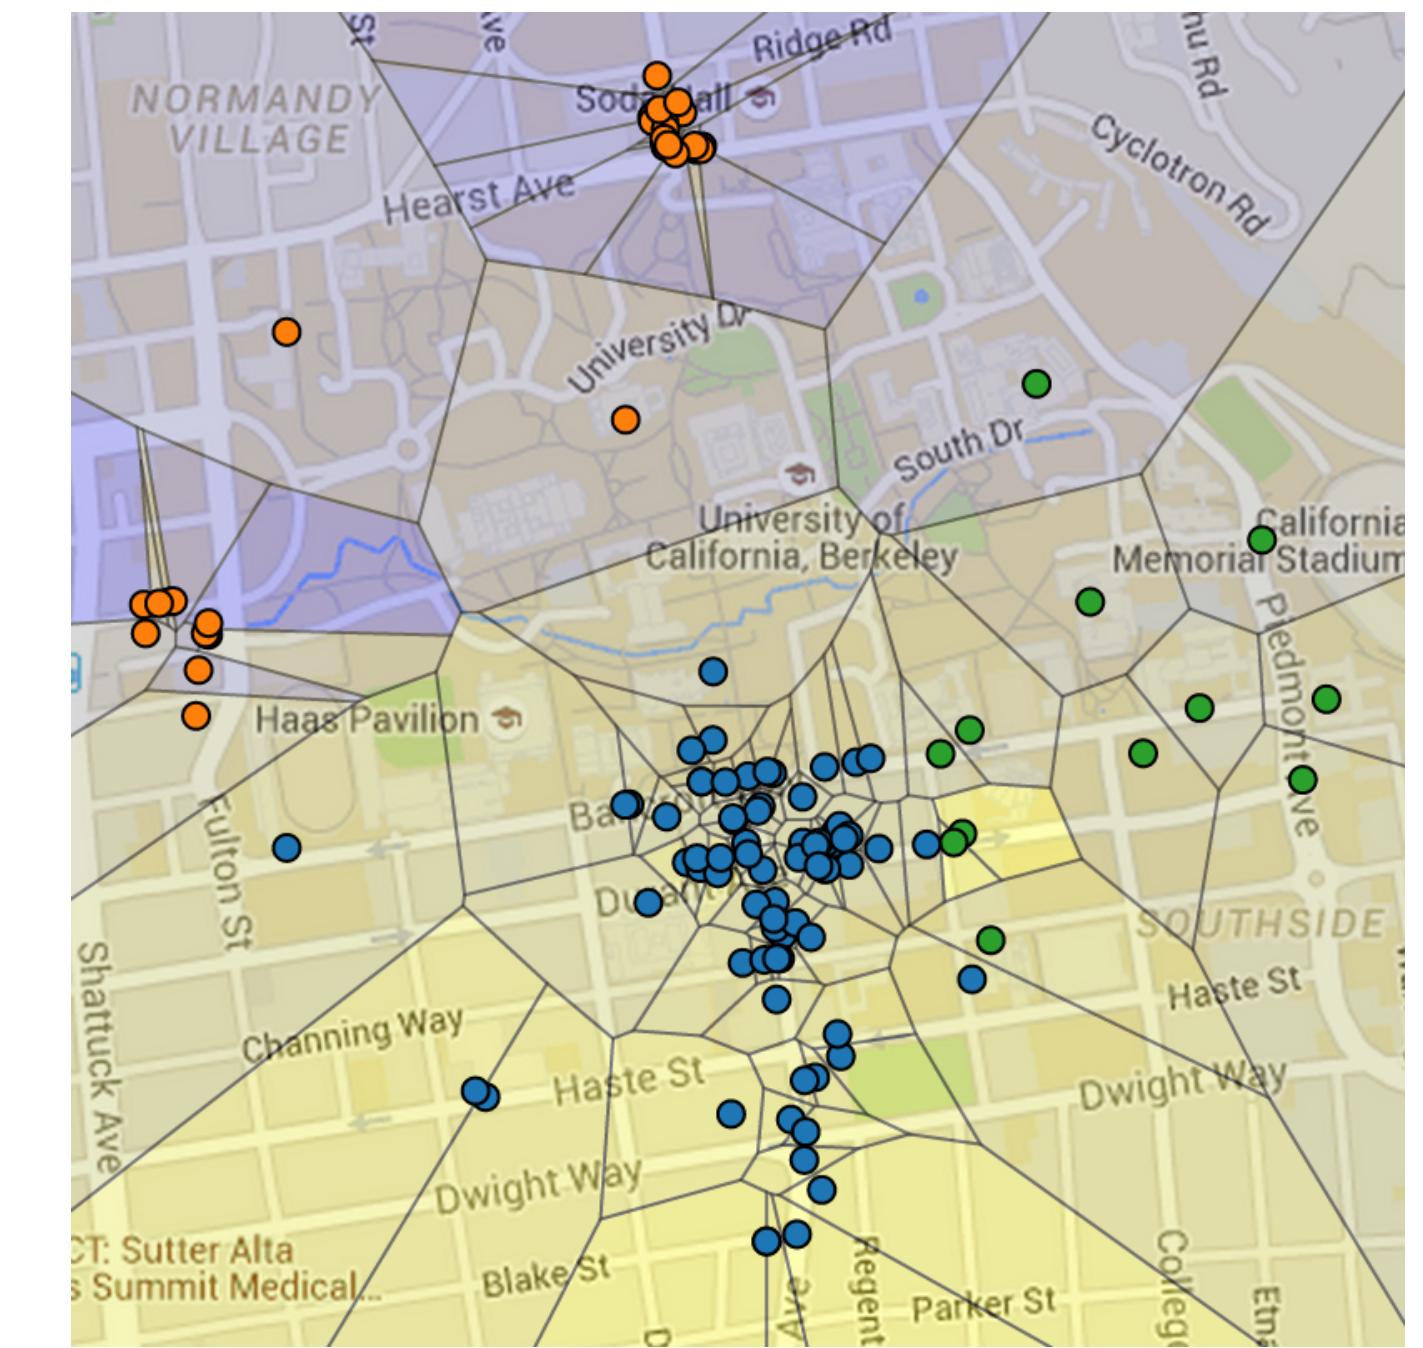
---



## Linear Regression

- The data: user ratings
- The decision: what rating would the user give a new restaurant?

Called *supervised Learning*, because some correct decisions are given.



# Outline

---

# Outline

---

- So far, we've looked at two specific machine learning algorithms from two different domains.

# Outline

---

- So far, we've looked at two specific machine learning algorithms from two different domains.
- Today, we will focus on a subclass of problems in machine learning, known as *reinforcement Learning* problems, and algorithms for these problems.

# Reinforcement Learning

---

# Reinforcement Learning

---

What is reinforcement learning?

# Reinforcement Learning

---

What is reinforcement learning?

- Concerned with *Learning behavior through experience.*

# Reinforcement Learning

---

What is reinforcement learning?

- Concerned with *Learning behavior through experience.*
- Two main components: the *agent* and the *environment*.

# Reinforcement Learning

---

What is reinforcement learning?

- Concerned with *Learning behavior through experience.*
- Two main components: the *agent* and the *environment*.
- The agent lives in and interacts with the environment, and through this experience learns a good pattern of behavior.

# An Analogy

---

# An Analogy

---

Suppose you go on a date with someone.

# An Analogy

---

Suppose you go on a date with someone.



# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.

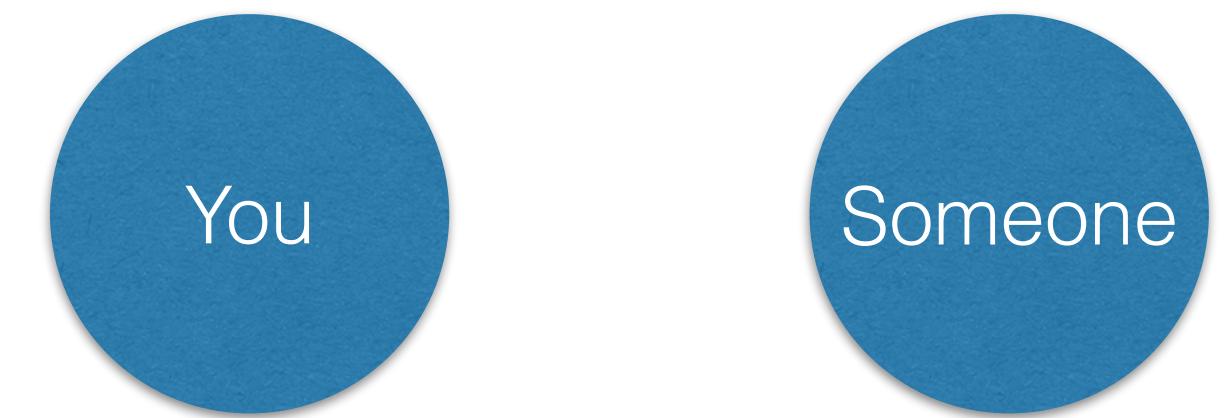


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.



# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.

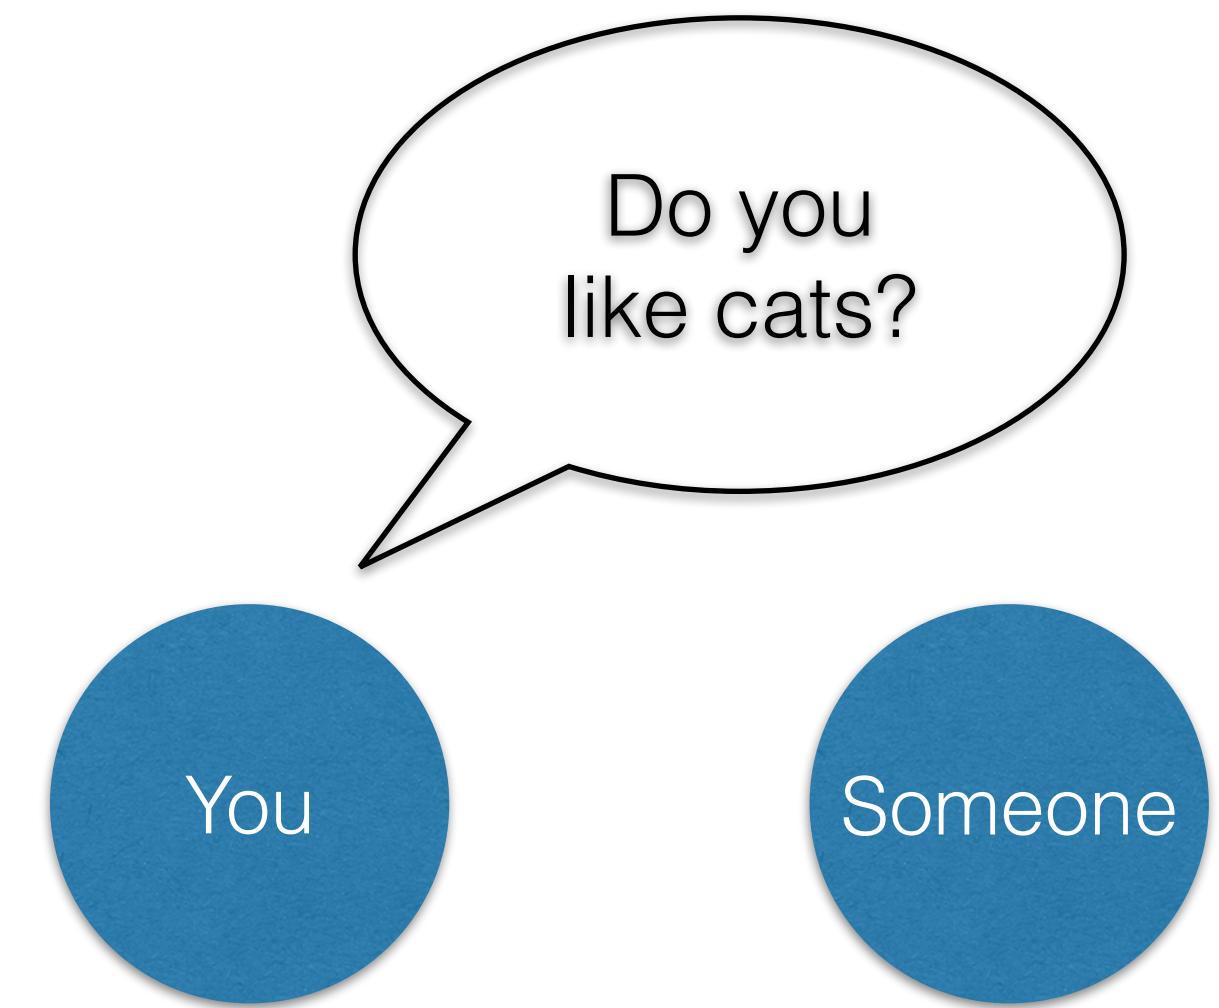


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.

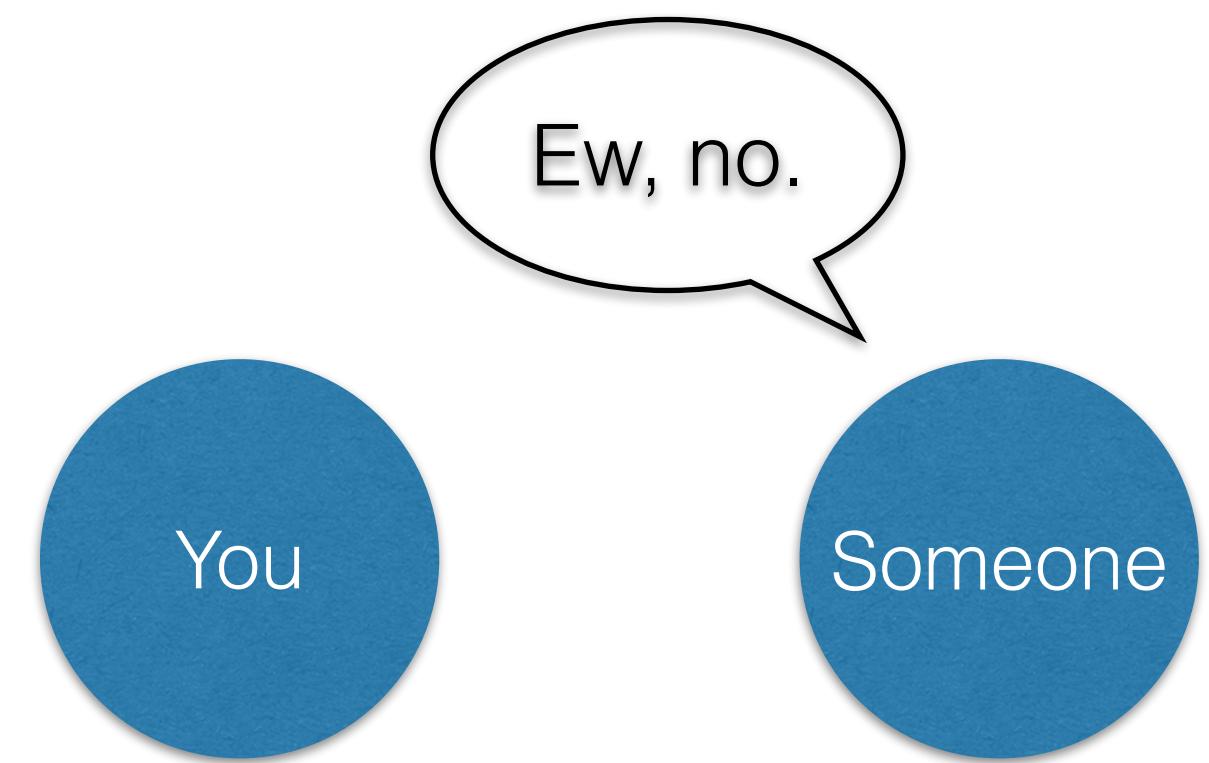


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.

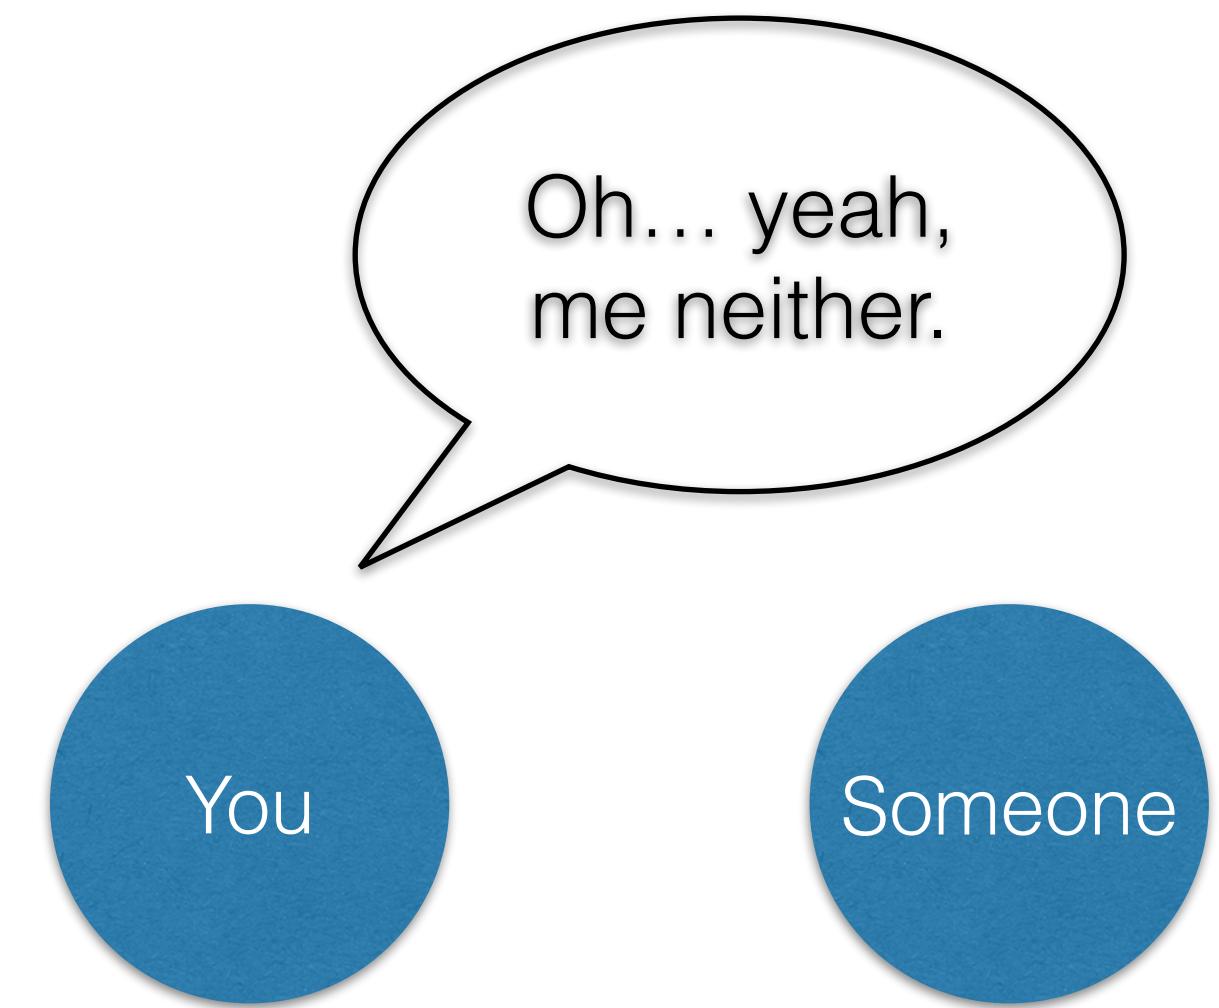


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.

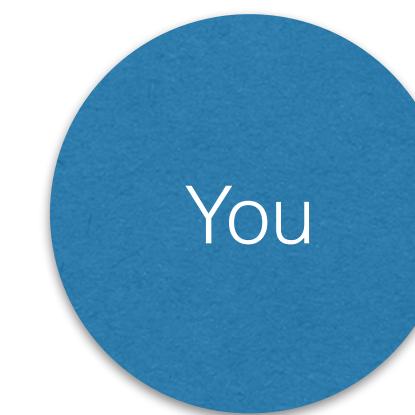


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.

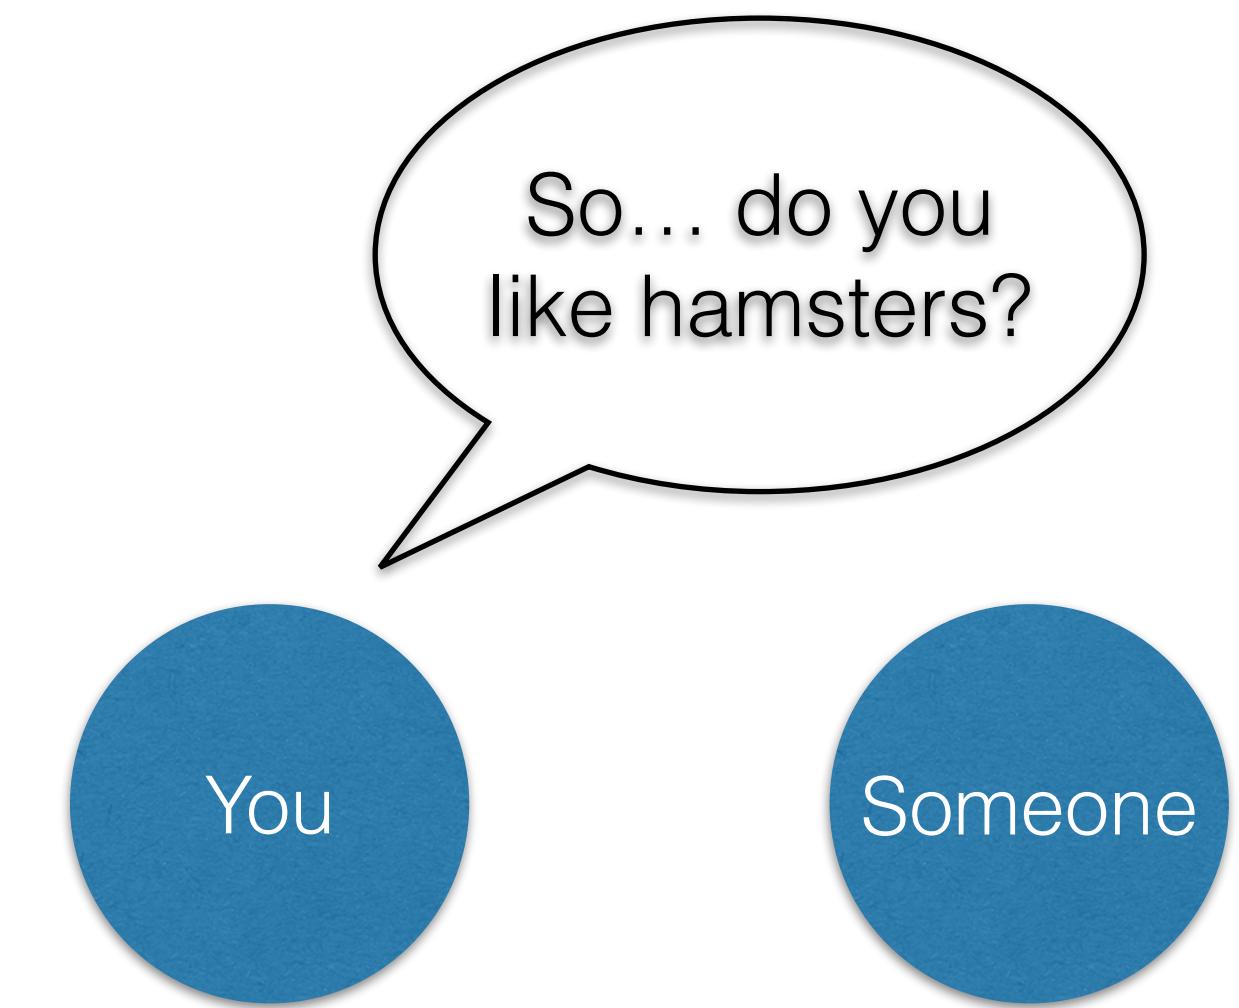


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.

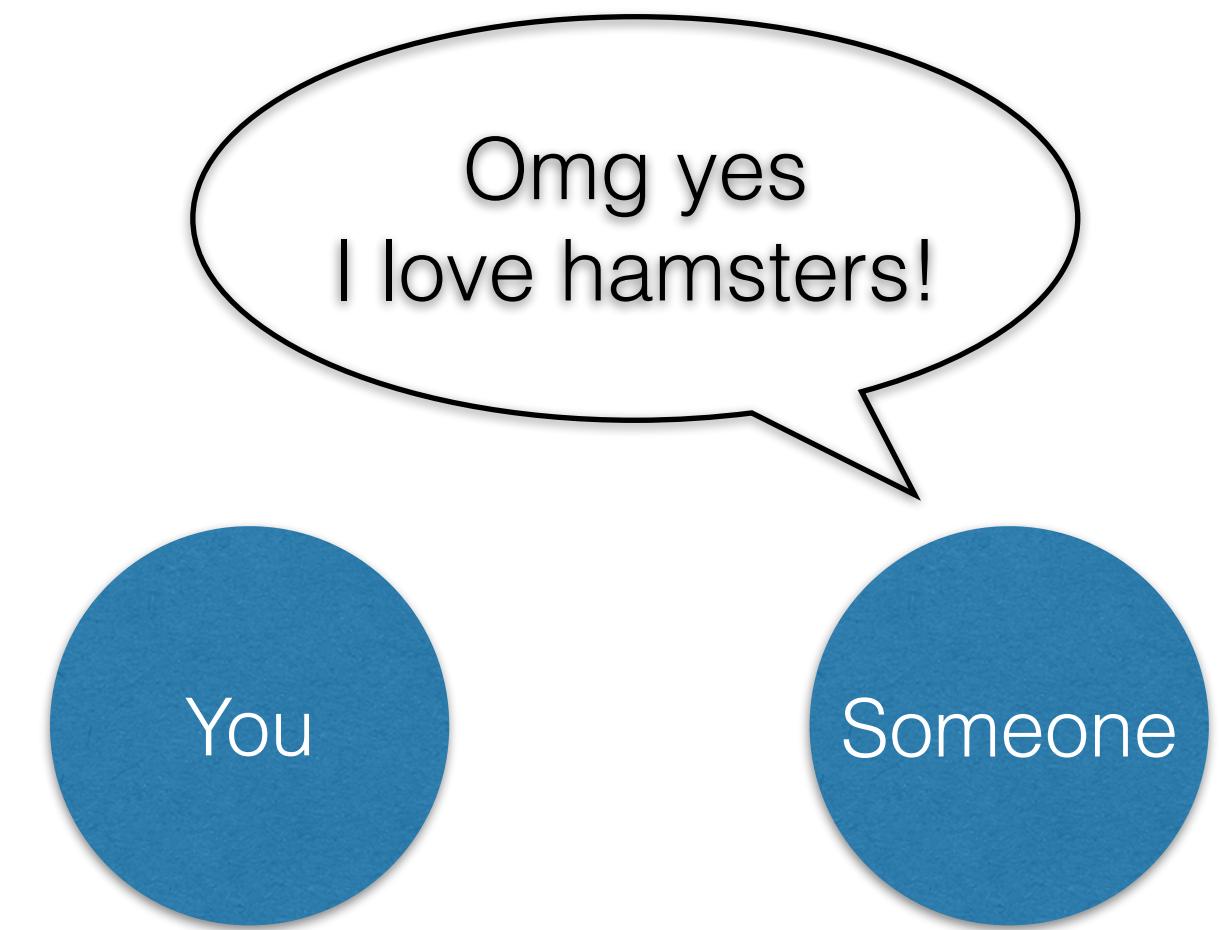


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.

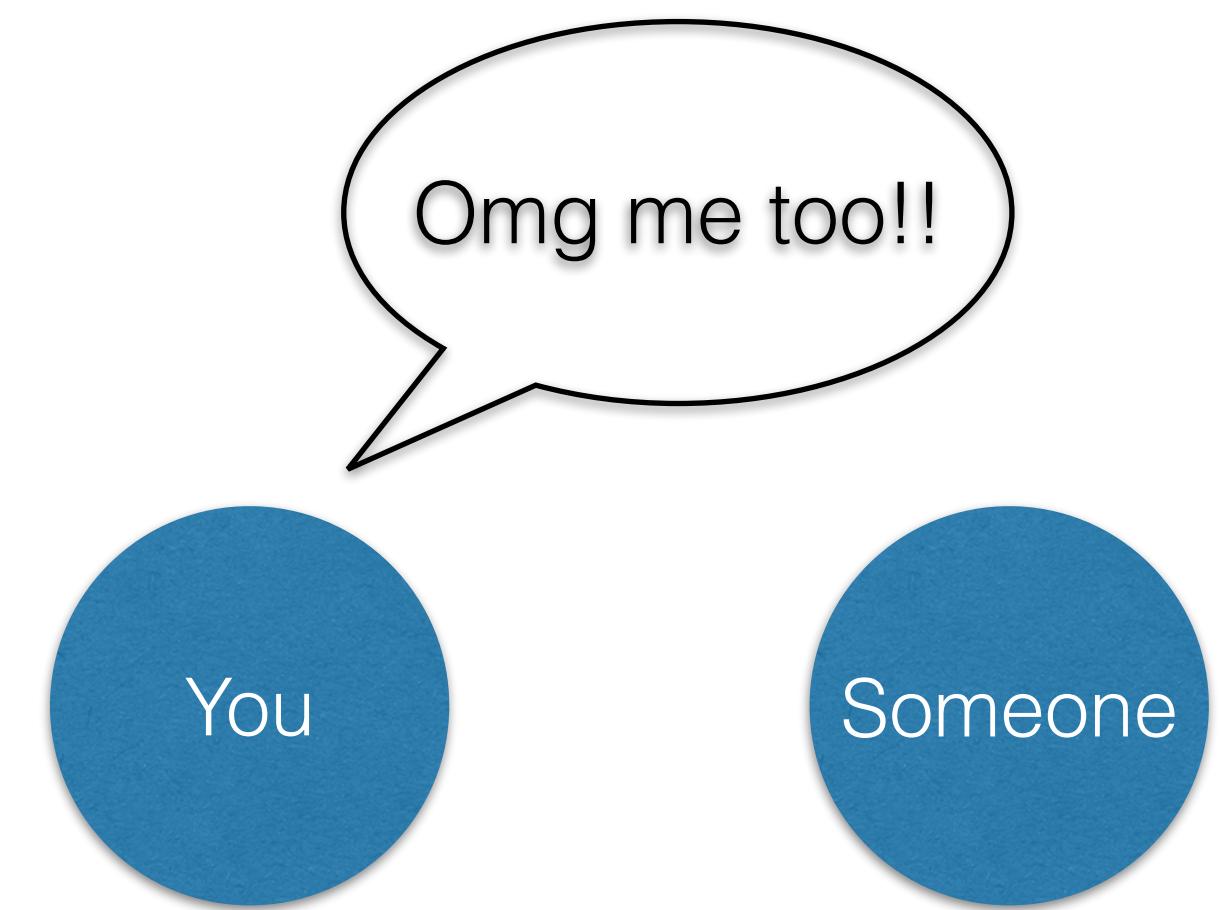


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.



# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.
- If you're a good agent, you may even learn how to behave really, really well!



# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.
- If you're a good agent, you may even learn how to behave really, really well!

45 minutes of talking  
about hamsters later...

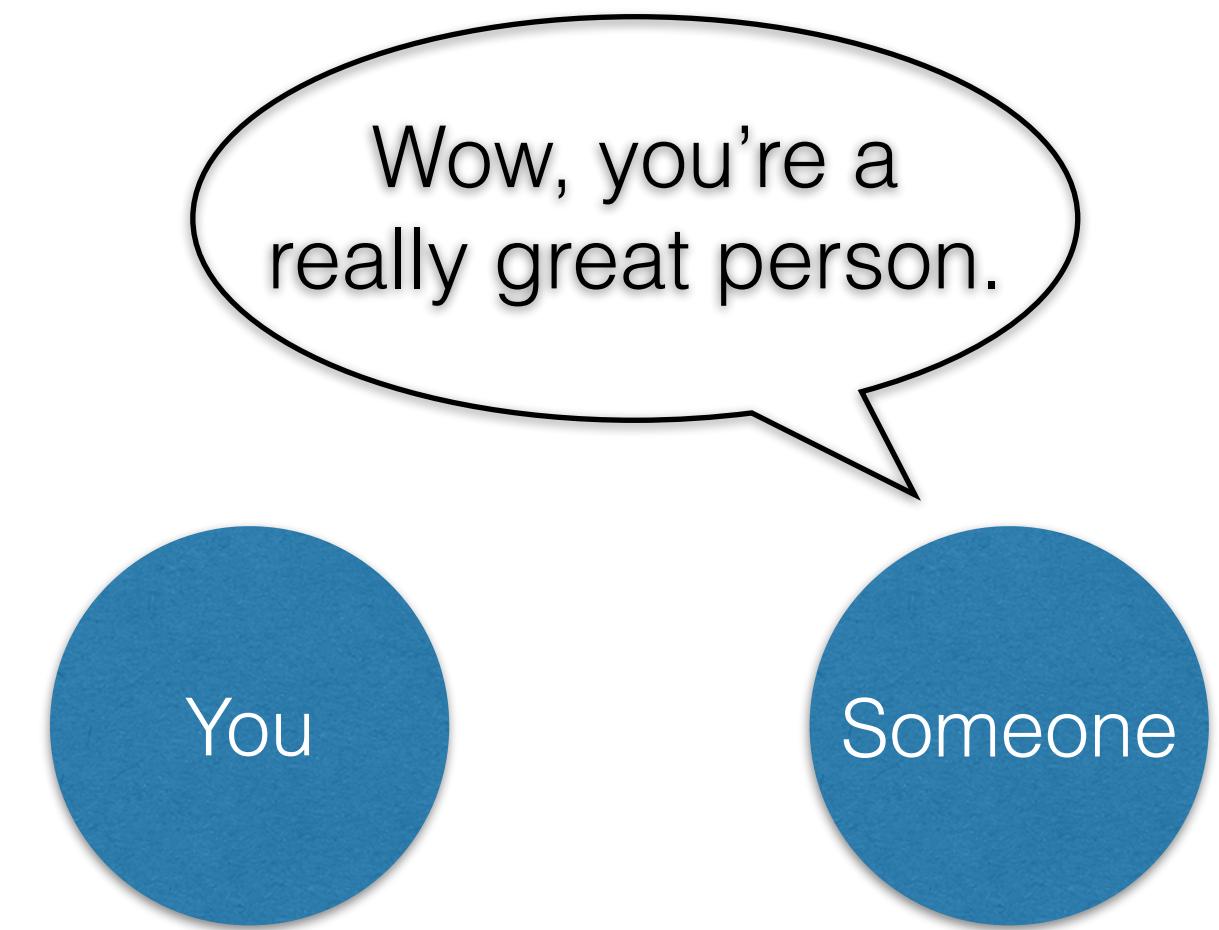


# An Analogy

---

Suppose you go on a date with someone.

- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.
- If you're a good agent, you may even learn how to behave really, really well!



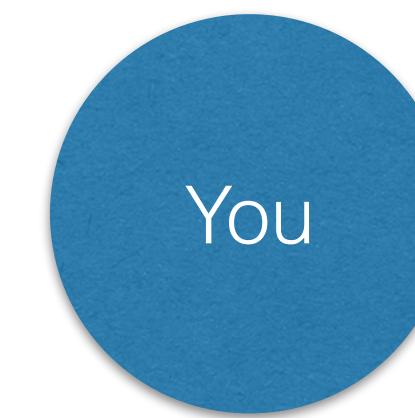
# An Analogy

---

Suppose you go on a date with someone.

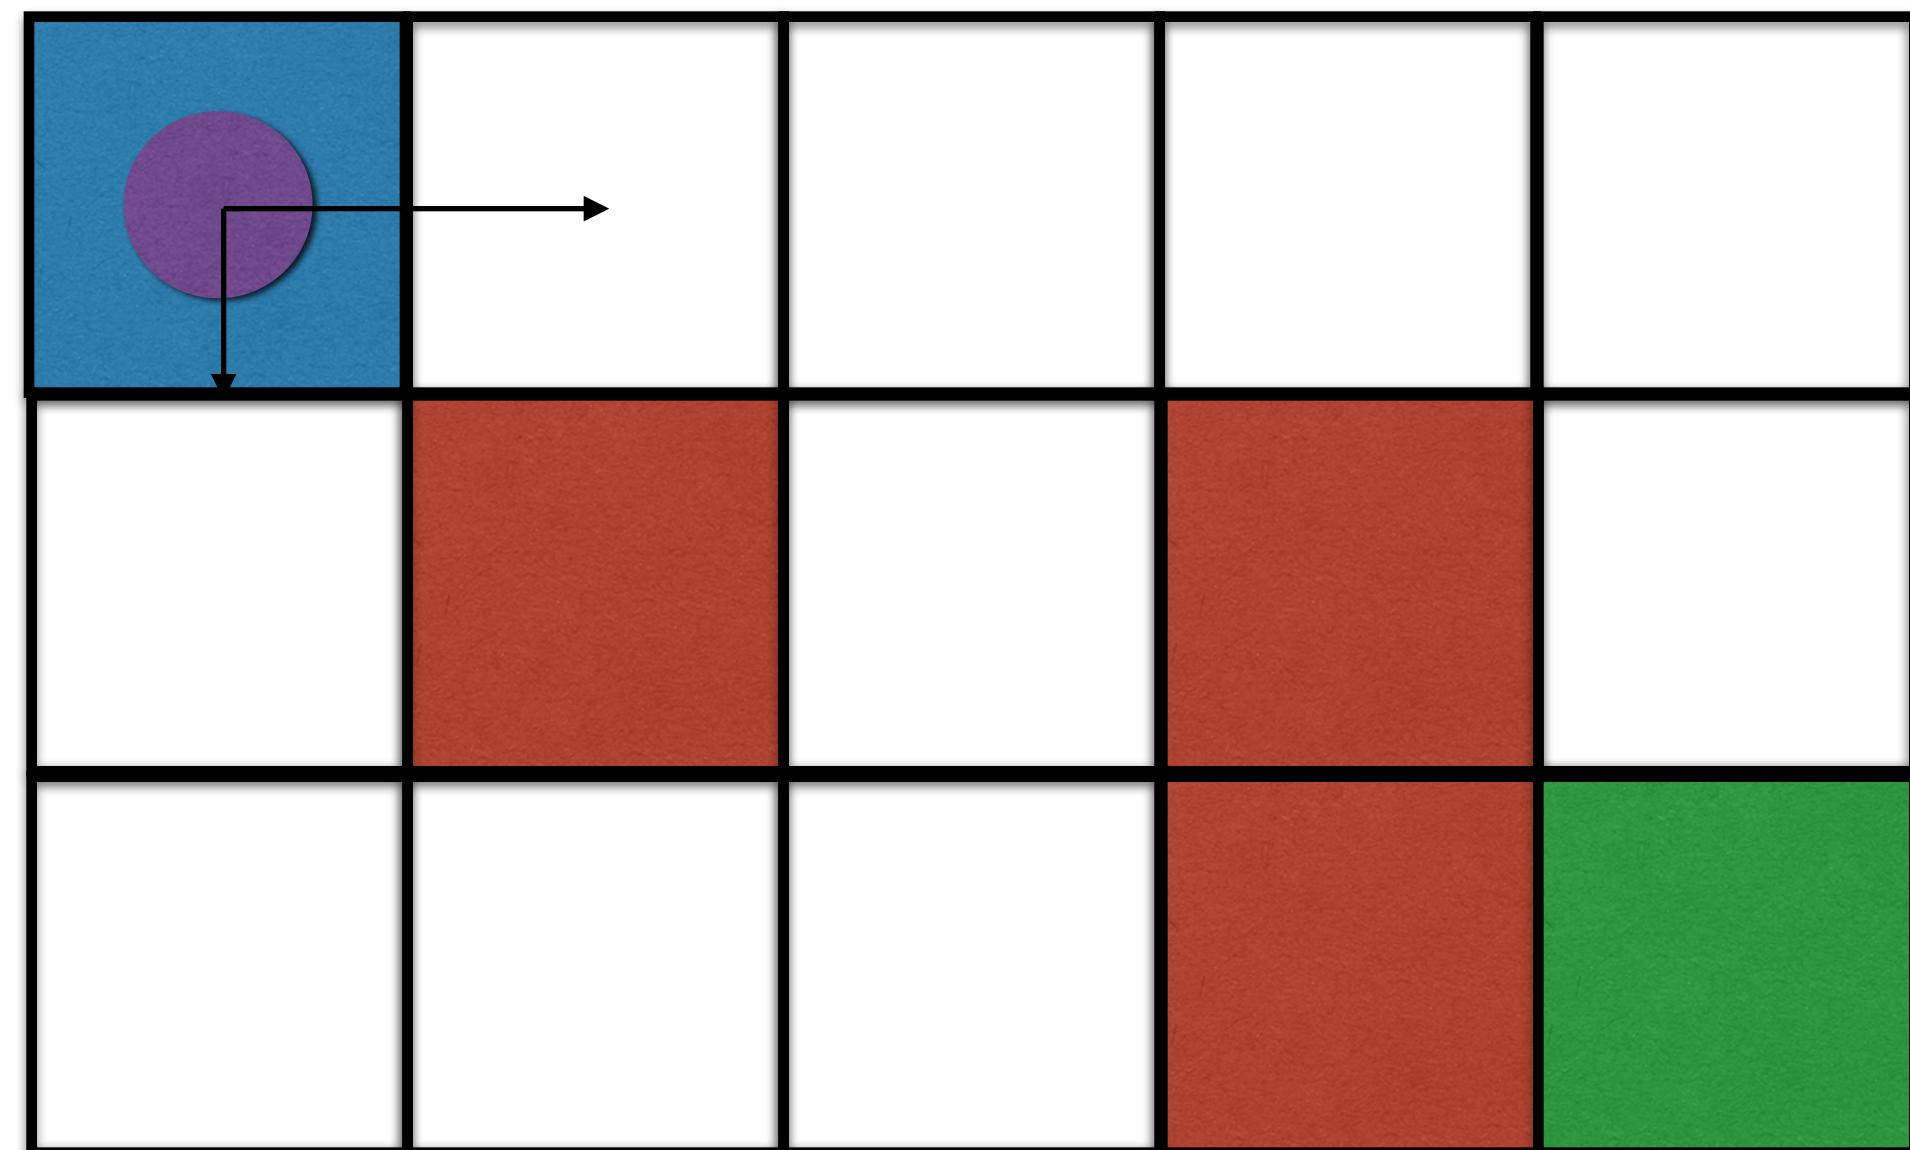
- In reinforcement learning terms, you are the *agent*.
- Everything else (the other person, the setting, etc.) is the *environment*.
- At the beginning of the date, you might not know how to act, so you try different things to see how the other person responds.
- As the date goes on, you slowly figure out how you should behave based on what you've tried so far, and how it went.
- If you're a good agent, you may even learn how to behave really, really well!

DATE:  
SUCCESS



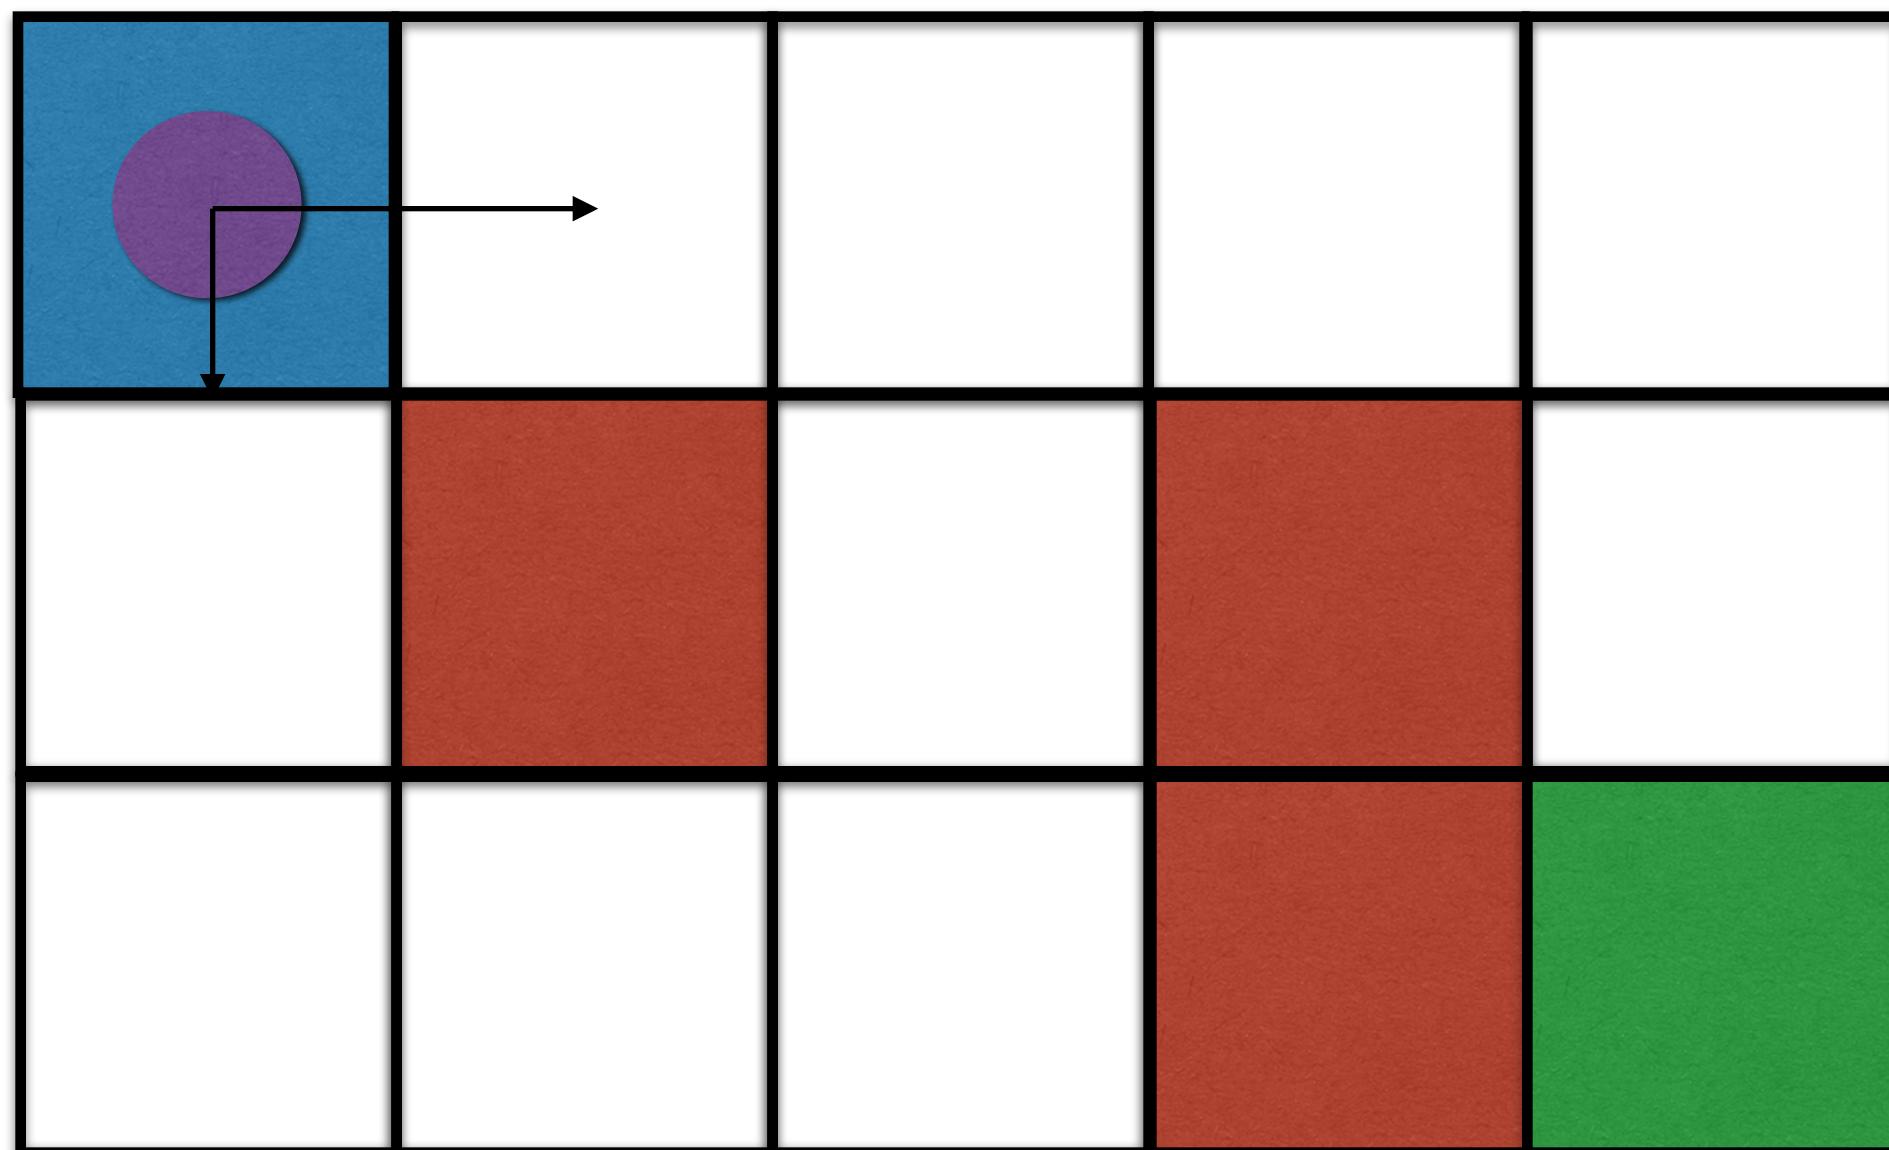
# RL Example: Gridworld

---



# RL Example: Gridworld

---

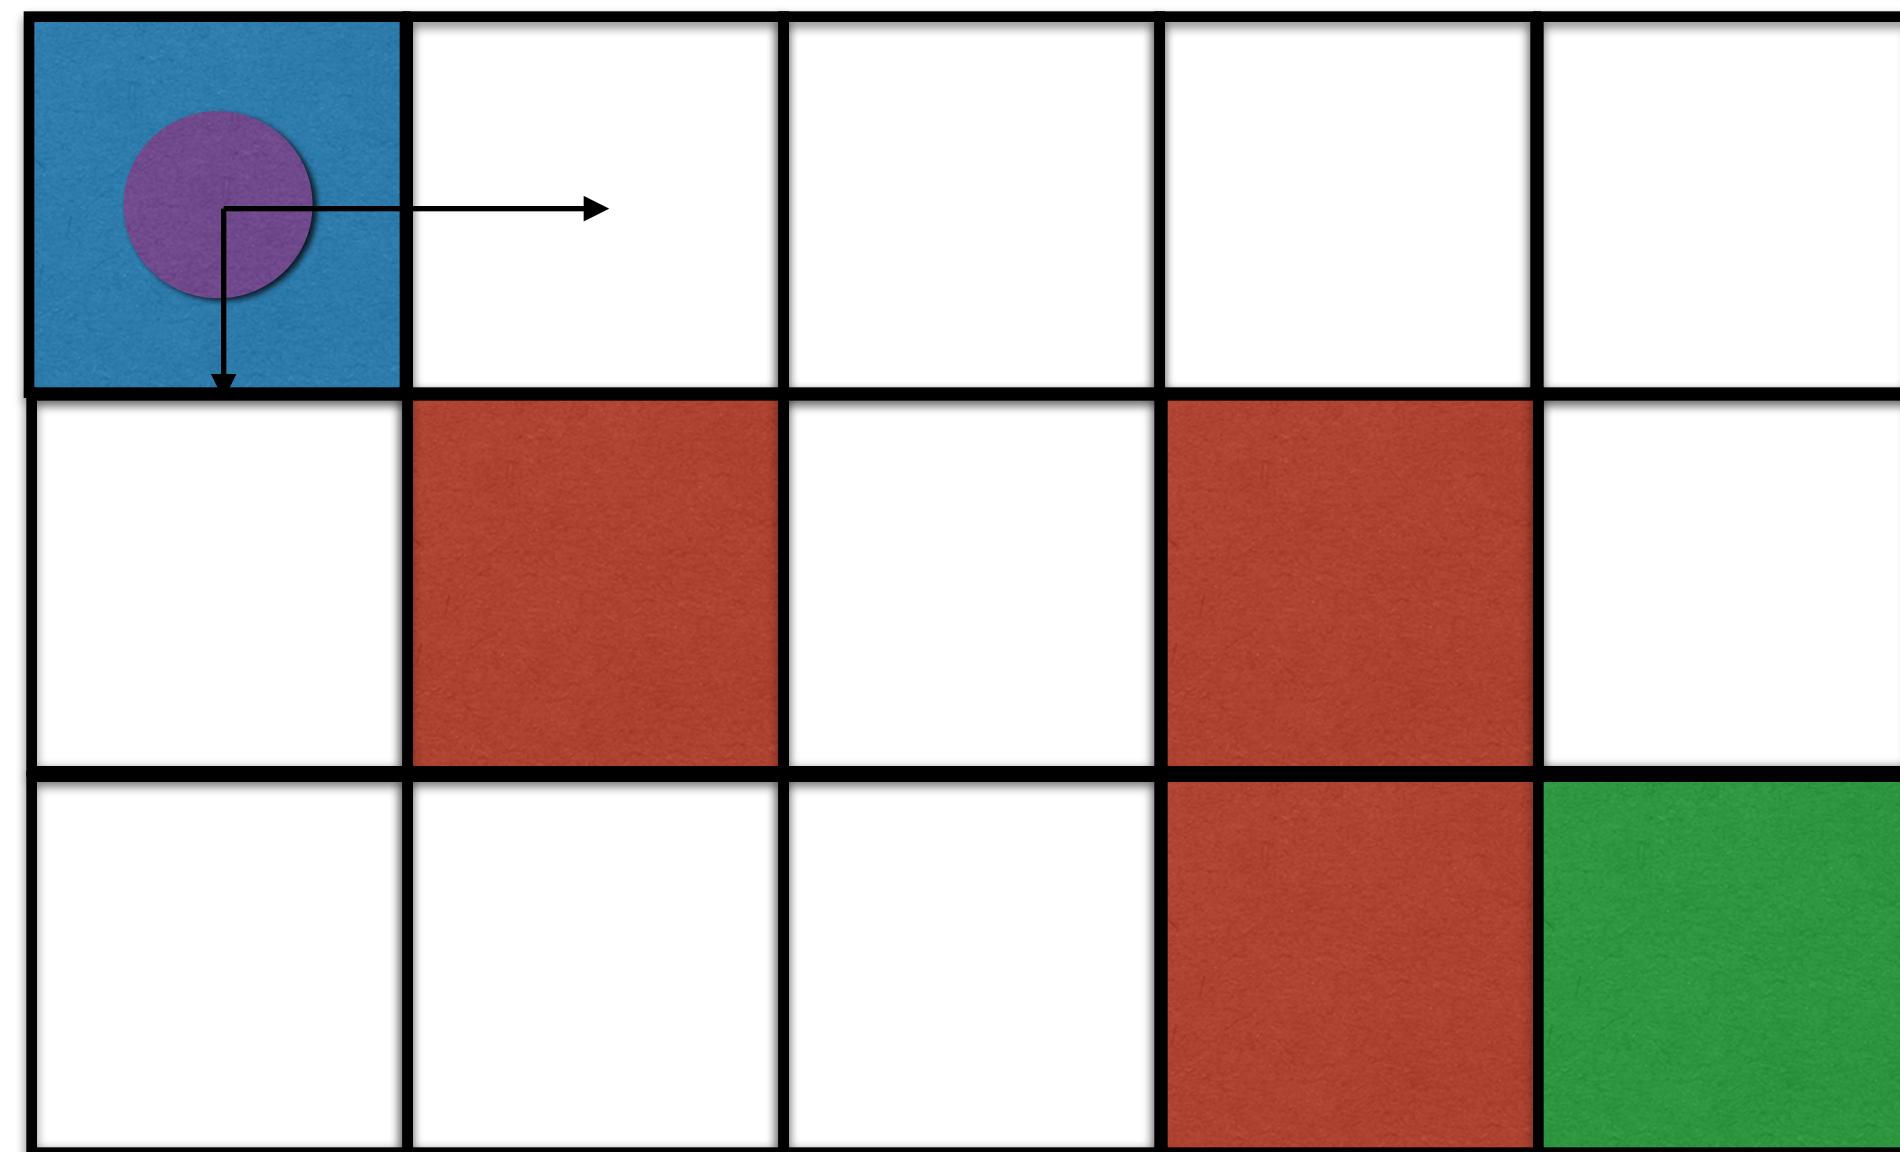


What is the environment?

What is the agent?

# RL Example: Gridworld

---



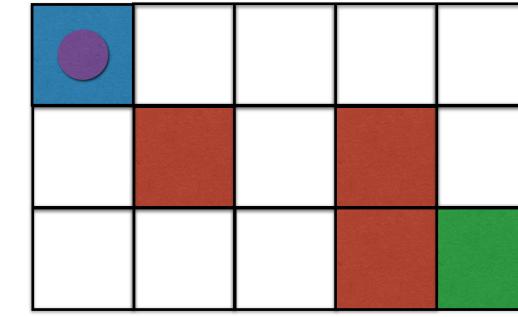
What is the environment?

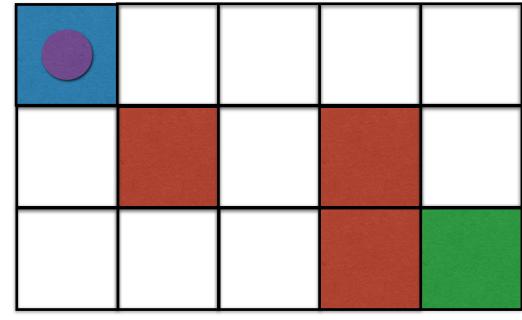
What is the agent?

The Problem: How do we get to the goal (green) from the start (blue) *as quickly as possible* while avoiding the obstacles (red)?

# The RL Setting

---

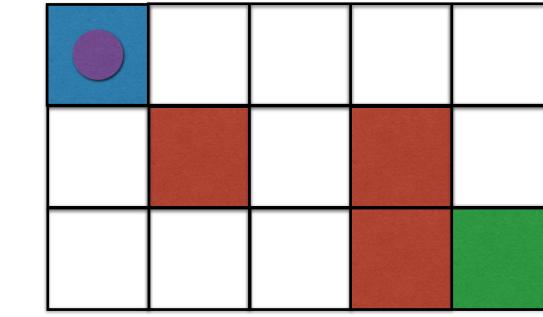




# The RL Setting

---

The environment:



# The RL Setting

---

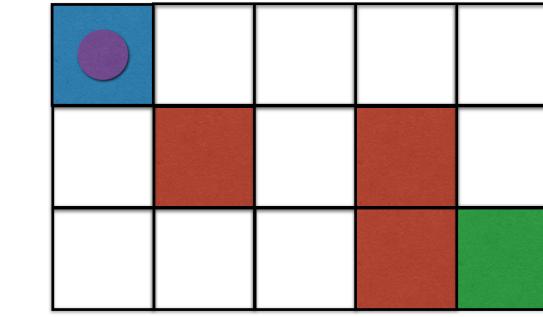
The environment:

- States ( $s$ ):

Configuration of the agent and environment.

# The RL Setting

---

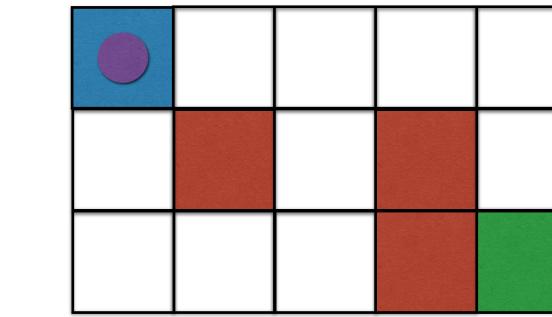


The environment:

- States ( $s$ ):  
Configuration of the agent and environment.
- Actions ( $a$ ):  
What can the agent do in a state?

# The RL Setting

---

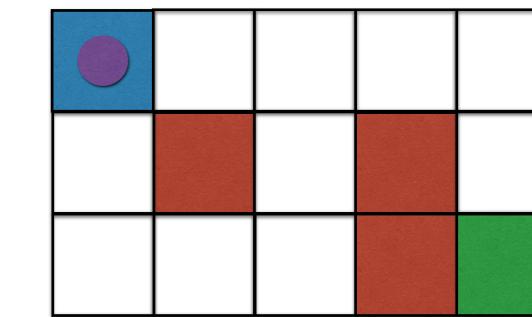


The environment:

- States ( $s$ ):  
Configuration of the agent and environment.
- Actions ( $a$ ):  
What can the agent do in a state?
- Reward function ( $R$ ):  
What reward does the agent get for each state?

# The RL Setting

---



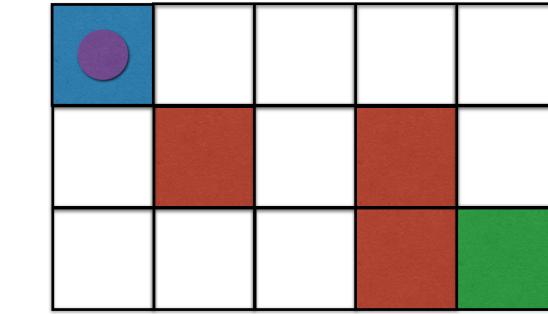
The environment:

- States ( $s$ ):  
Configuration of the agent and environment.
- Actions ( $a$ ):  
What can the agent do in a state?
- Reward function ( $R$ ):  
What reward does the agent get for each state?

The agent:

# The RL Setting

---



The environment:

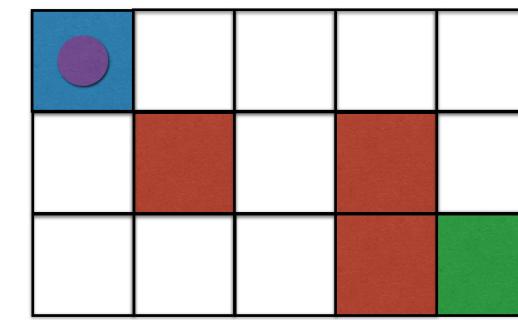
- States ( $s$ ):  
Configuration of the agent and environment.
- Actions ( $a$ ):  
What can the agent do in a state?
- Reward function ( $R$ ):  
What reward does the agent get for each state?

The agent:

- Policy ( $\pi$ ):  
Given a state, what action will the agent take?

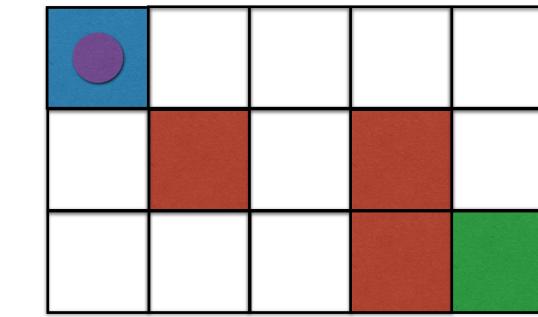
# Gridworld Revisited

---



# Gridworld Revisited

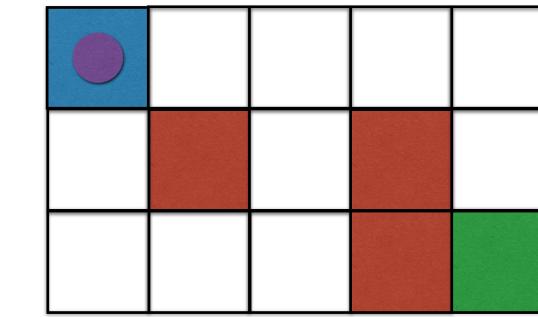
---



- The environment of Gridworld, in more detail:

# Gridworld Revisited

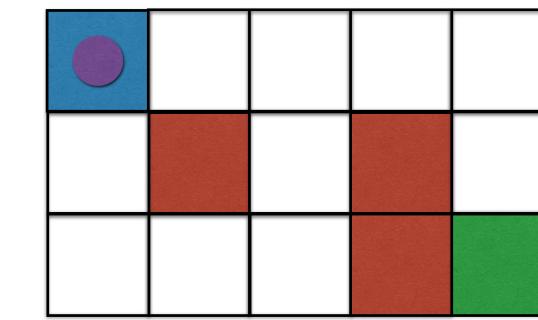
---



- The environment of Gridworld, in more detail:
  - States ( $s$ ):

# Gridworld Revisited

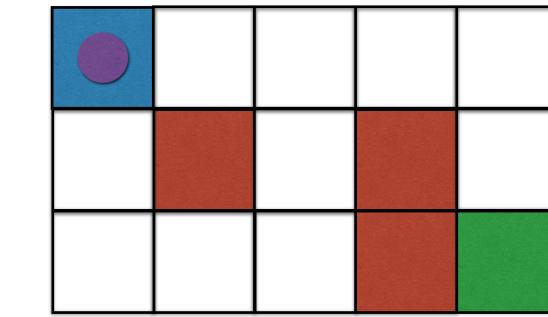
---



- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?

# Gridworld Revisited

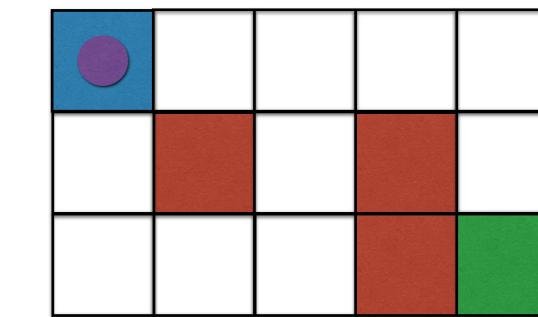
---



- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):

# Gridworld Revisited

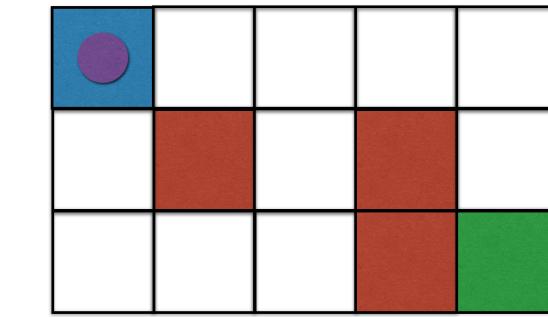
---



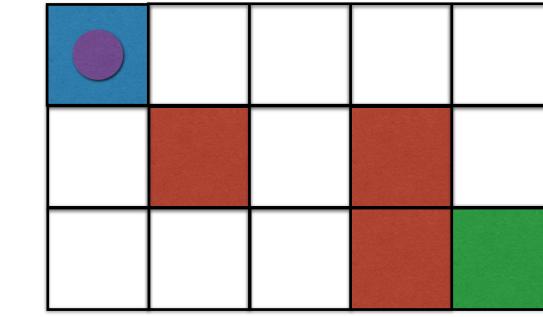
- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):  
Go to an adjacent square, or stay put.

# Gridworld Revisited

---



- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):  
Go to an adjacent square, or stay put.
  - Reward function: ???



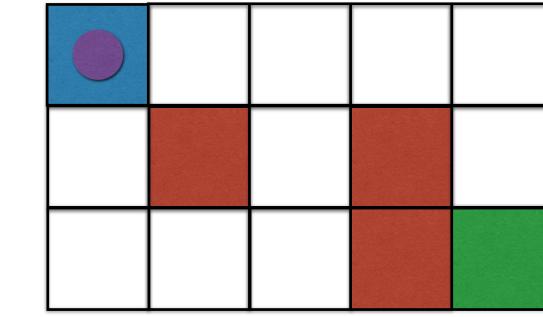
# Gridworld Revisited

---

- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):  
Go to an adjacent square, or stay put.
  - Reward function: ???

The Problem:

How do we get to the goal from the start *as quickly as possible* while avoiding the obstacles?



# Gridworld Revisited

---

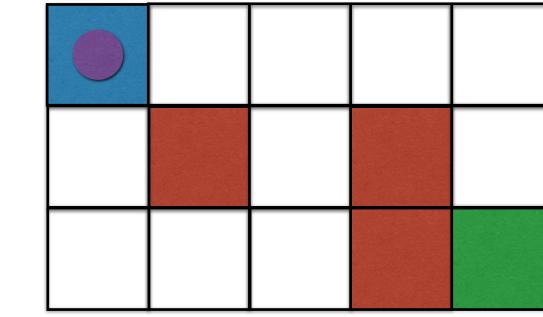
- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):  
Go to an adjacent square, or stay put.
  - Reward function: ???

The Problem:

How do we get to the goal from the start *as quickly as possible* while avoiding the obstacles?

In RL terminology:

What is the optimal policy  $\pi^*$  that *maximizes* my expected reward over time?



# Gridworld Revisited

---

- The environment of Gridworld, in more detail:
  - States ( $s$ ):  
What square is the agent in?
  - Actions ( $a$ ):  
Go to an adjacent square, or stay put.
  - Reward function: ???

The Problem:

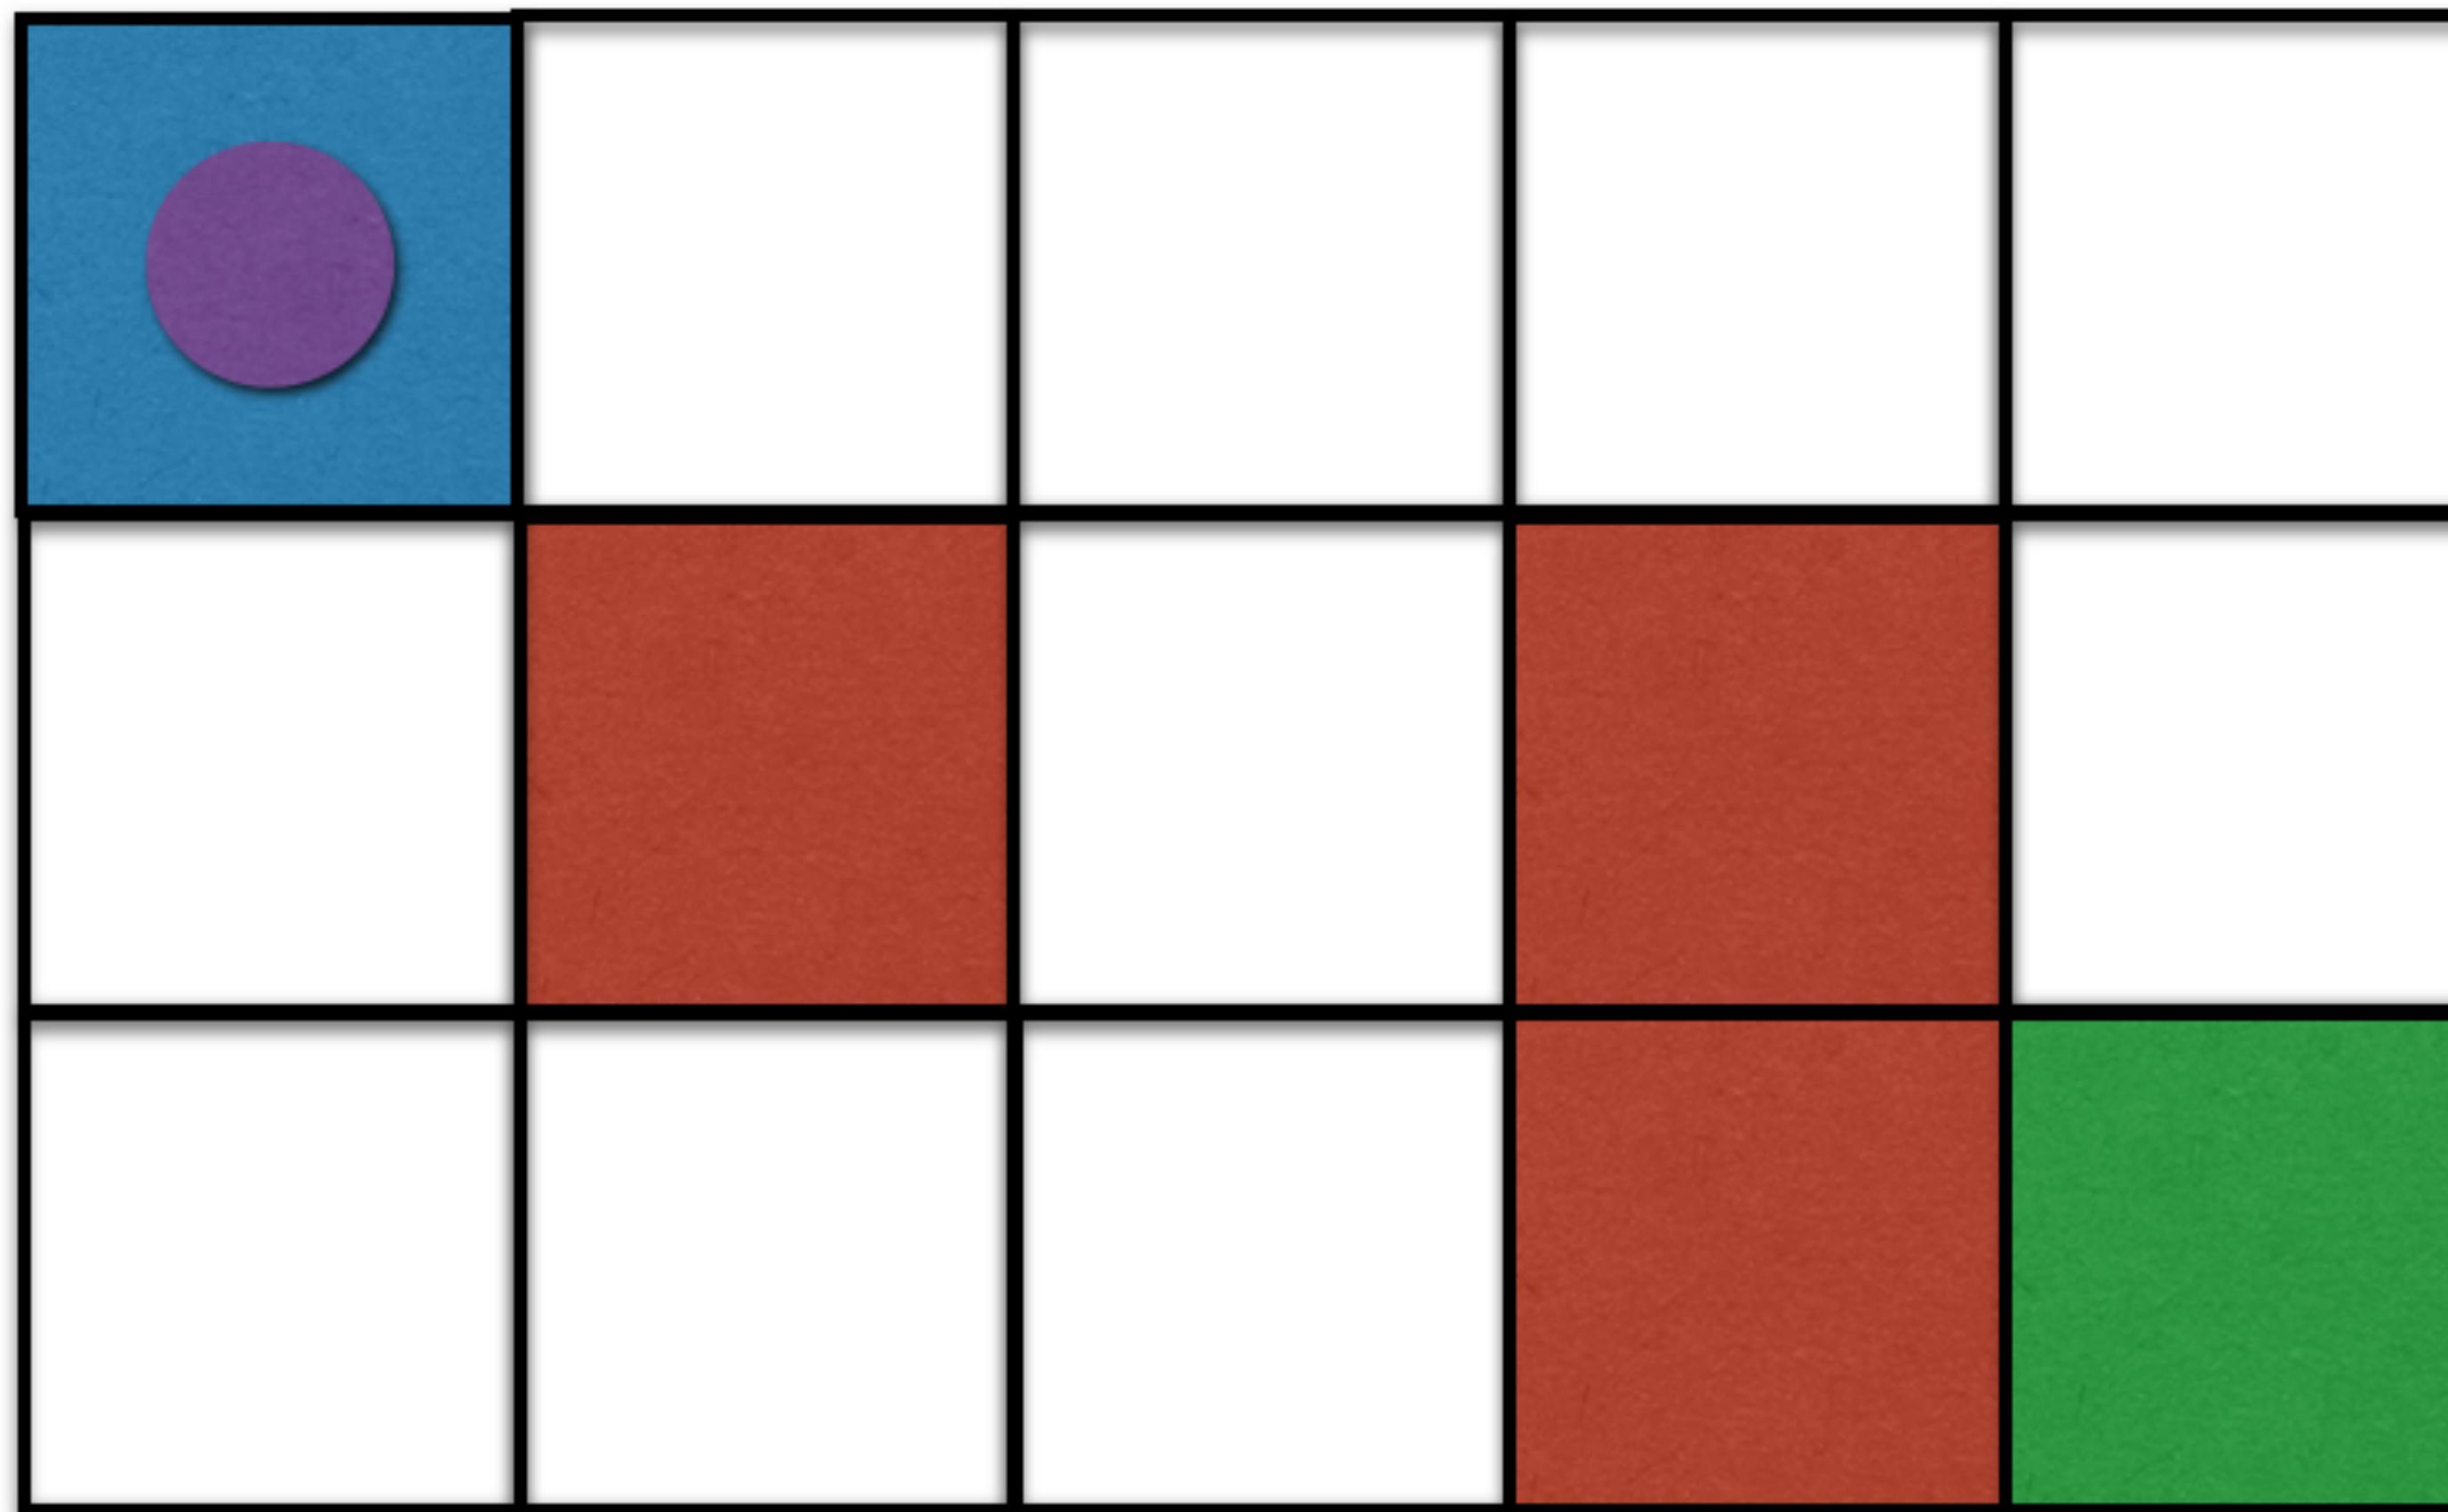
How do we get to the goal from the start *as quickly as possible* while avoiding the obstacles?

In RL terminology:

What is the optimal policy  $\pi^*$  that *maximizes* my expected reward over time?

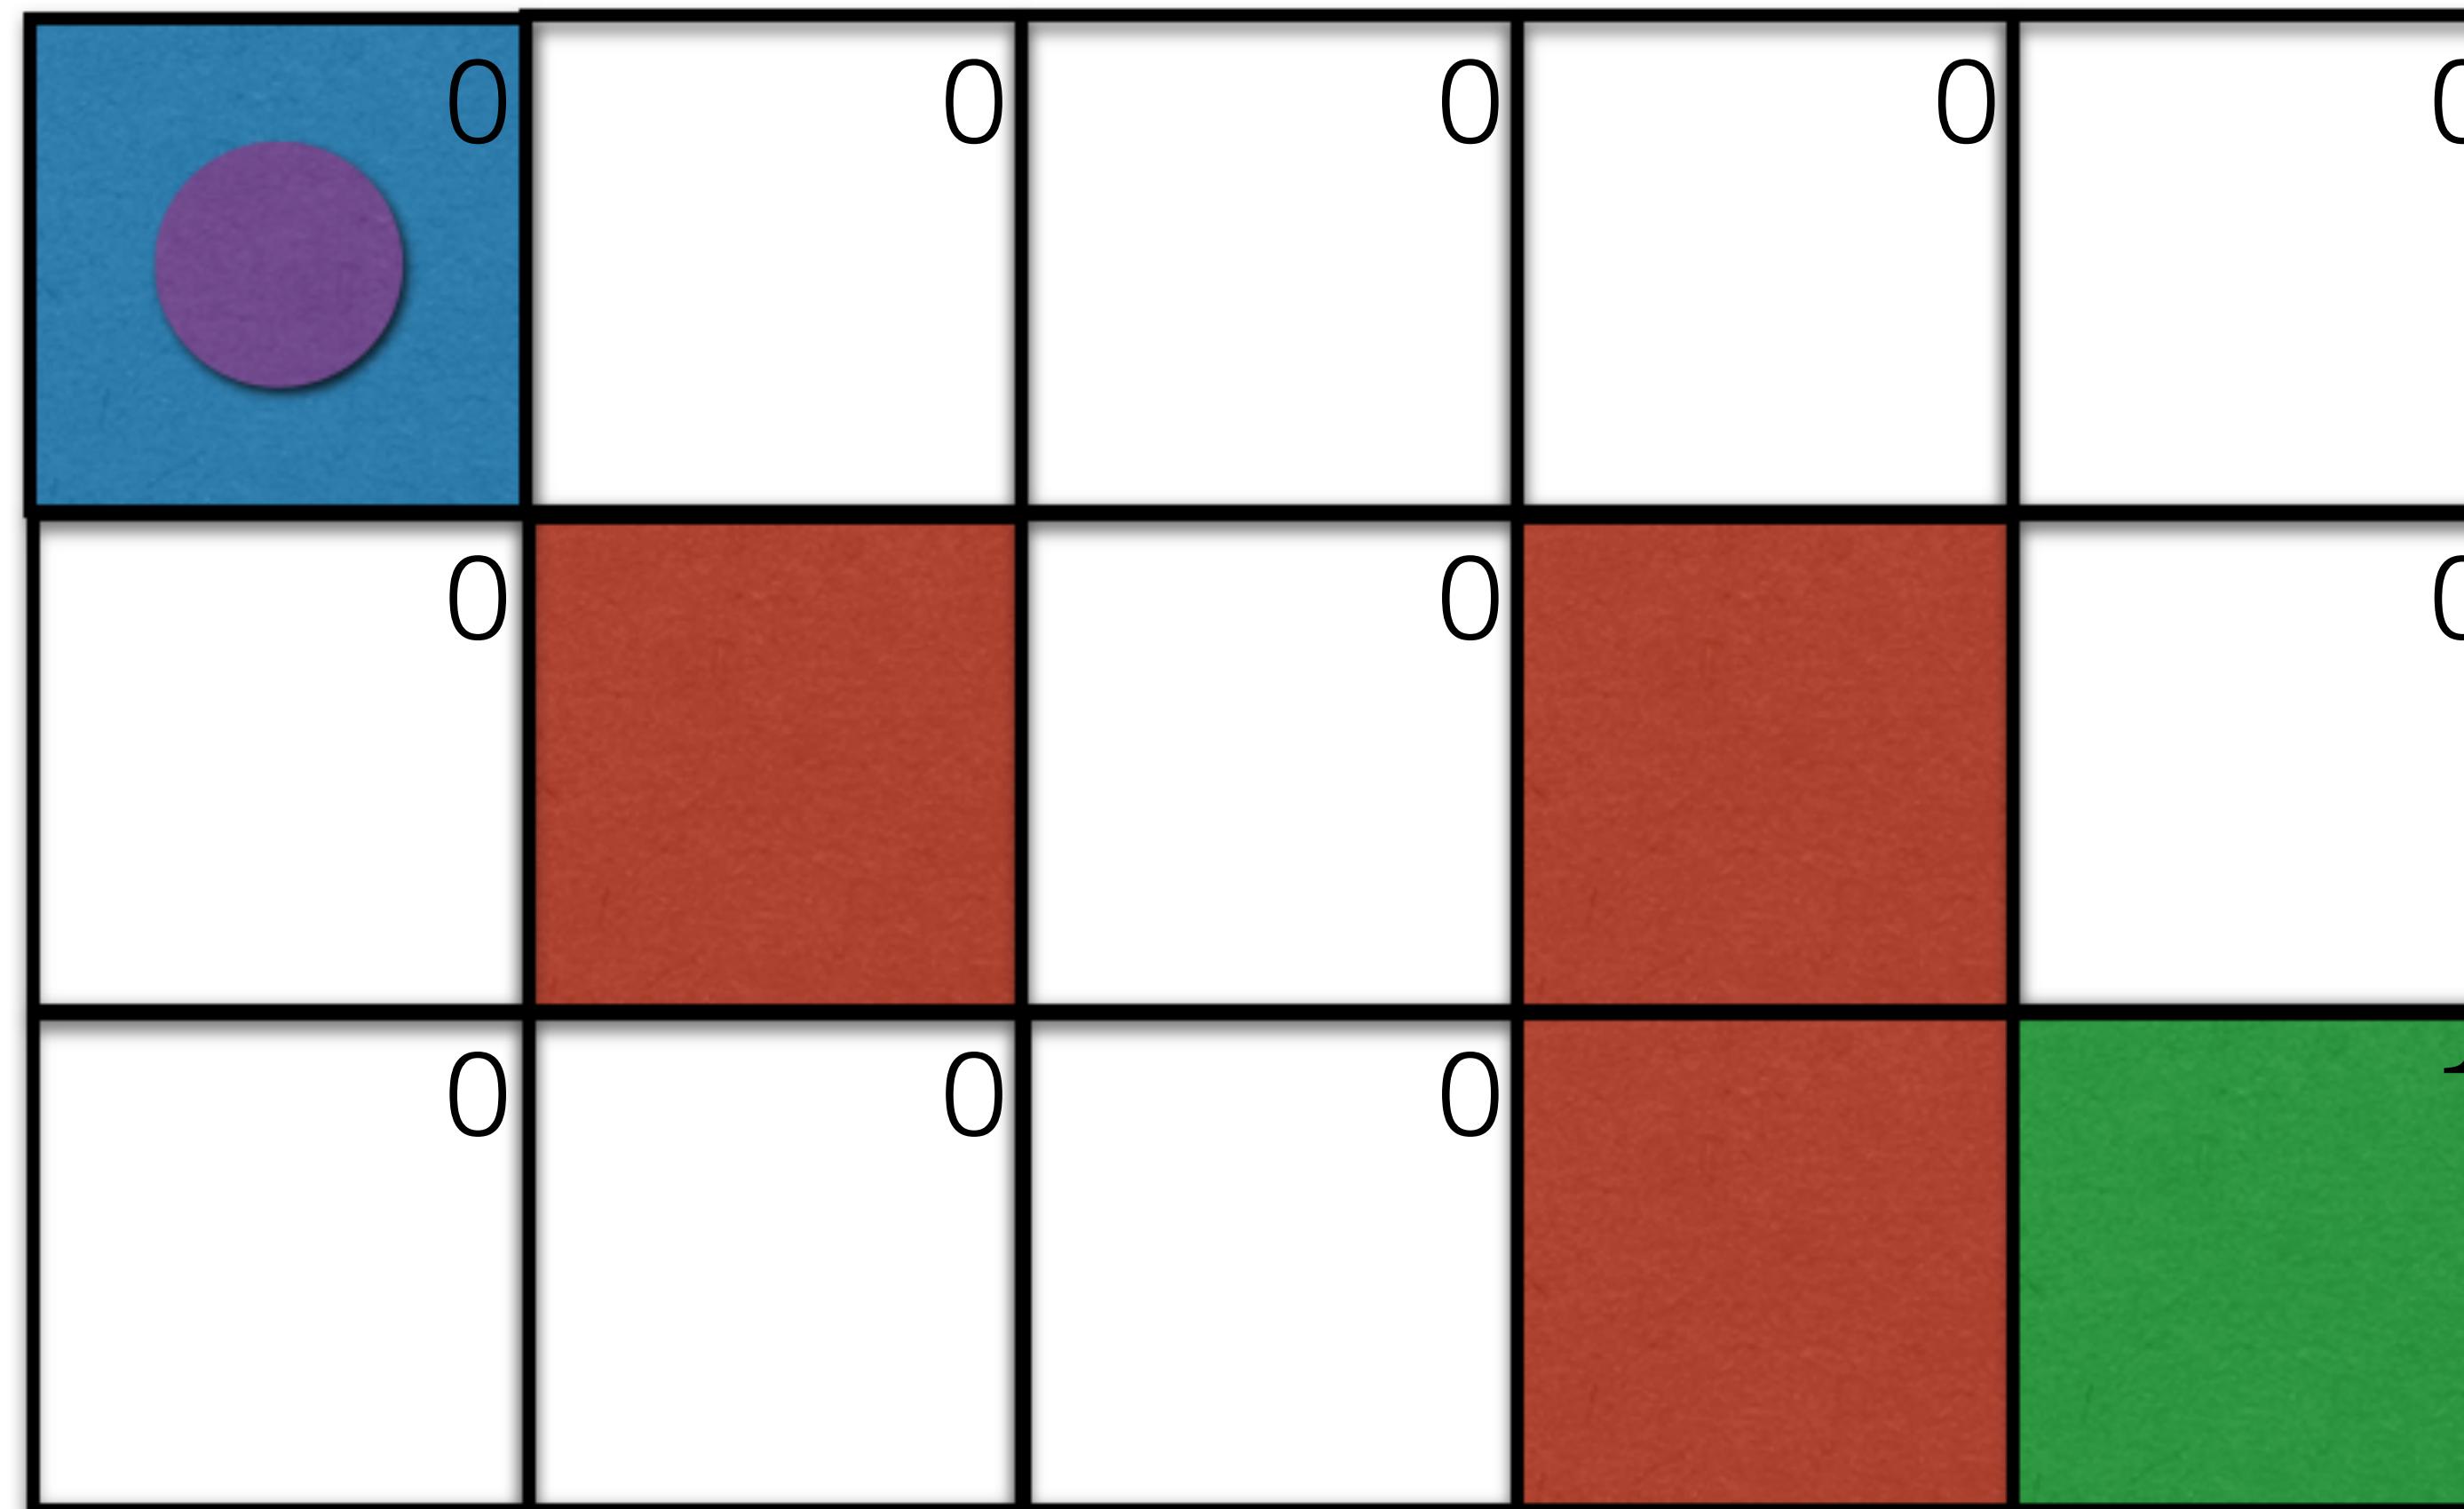
# Gridworld Reward Function

---



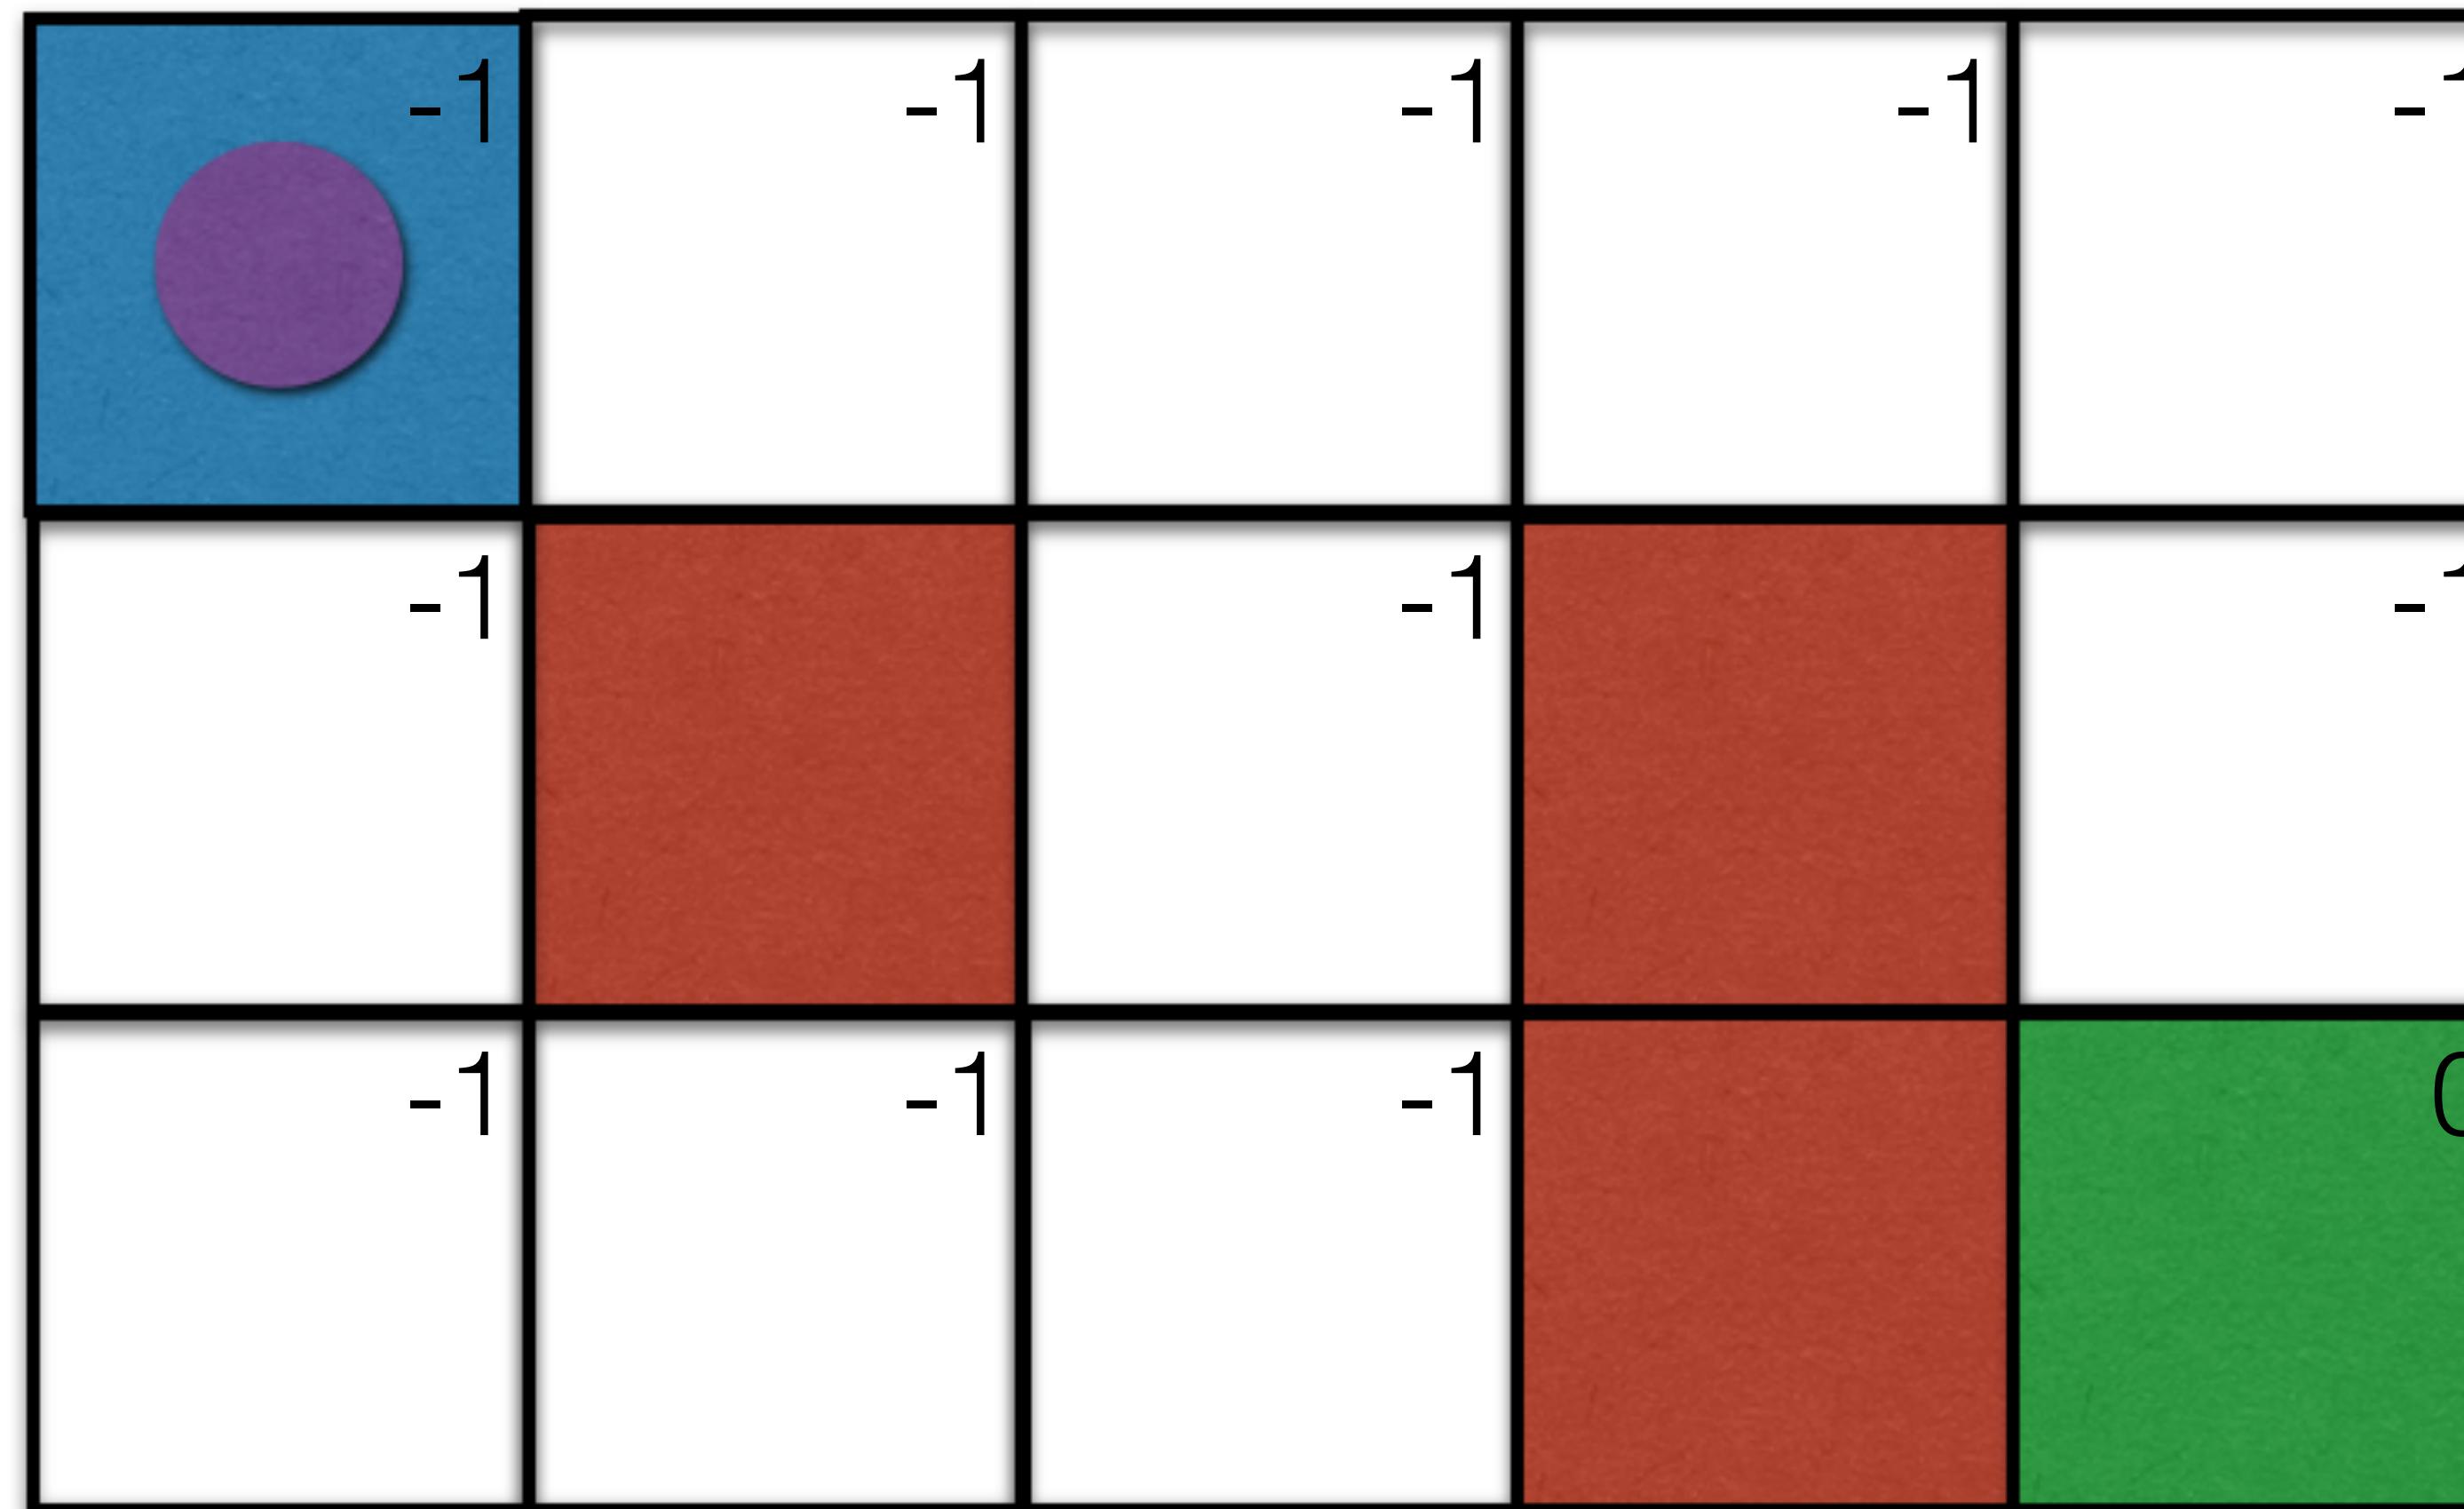
# Gridworld Reward Function

---



# Gridworld Reward Function

---



# Value Function

---

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?
  - We use our policy  $\pi$  to determine which actions to take.

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?
  - We use our policy  $\pi$  to determine which actions to take.
  - So, the value function *also depends on our policy*.

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?
  - We use our policy  $\pi$  to determine which actions to take.
  - So, the value function *also depends on our policy*.
- How do we determine the value of a state?

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?
  - We use our policy  $\pi$  to determine which actions to take.
  - So, the value function *also depends on our policy*.
- How do we determine the value of a state?
  - The value of a state is the *reward* of the state plus the *value* of the state we end up in next.

# Value Function

---

- Reward function:  $R(s)$  = reward of being in state  $s$
- Value function:  $V(s)$  = value of being in state  $s$
- The value of  $s$  is the *long-term expected reward* starting from  $s$ .
- How do we determine where to go after  $s$ ?
  - We use our policy  $\pi$  to determine which actions to take.
  - So, the value function *also depends on our policy*.
- How do we determine the value of a state?
  - The value of a state is the *reward* of the state plus the *value* of the state we end up in next.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

# Value Function

---

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

# Value Function

---

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- How do we solve this equation? Use recursion!

# Value Function

---

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- How do we solve this equation? Use recursion!
- What's our base case?

# Value Function

---

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- How do we solve this equation? Use recursion!
- What's our base case?
  - If we're at our *goal*, then there is no next state, so the value is just the reward.

# Value Function

---

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- How do we solve this equation? Use recursion!
- What's our base case?
  - If we're at our *goal*, then there is no next state, so the value is just the reward.

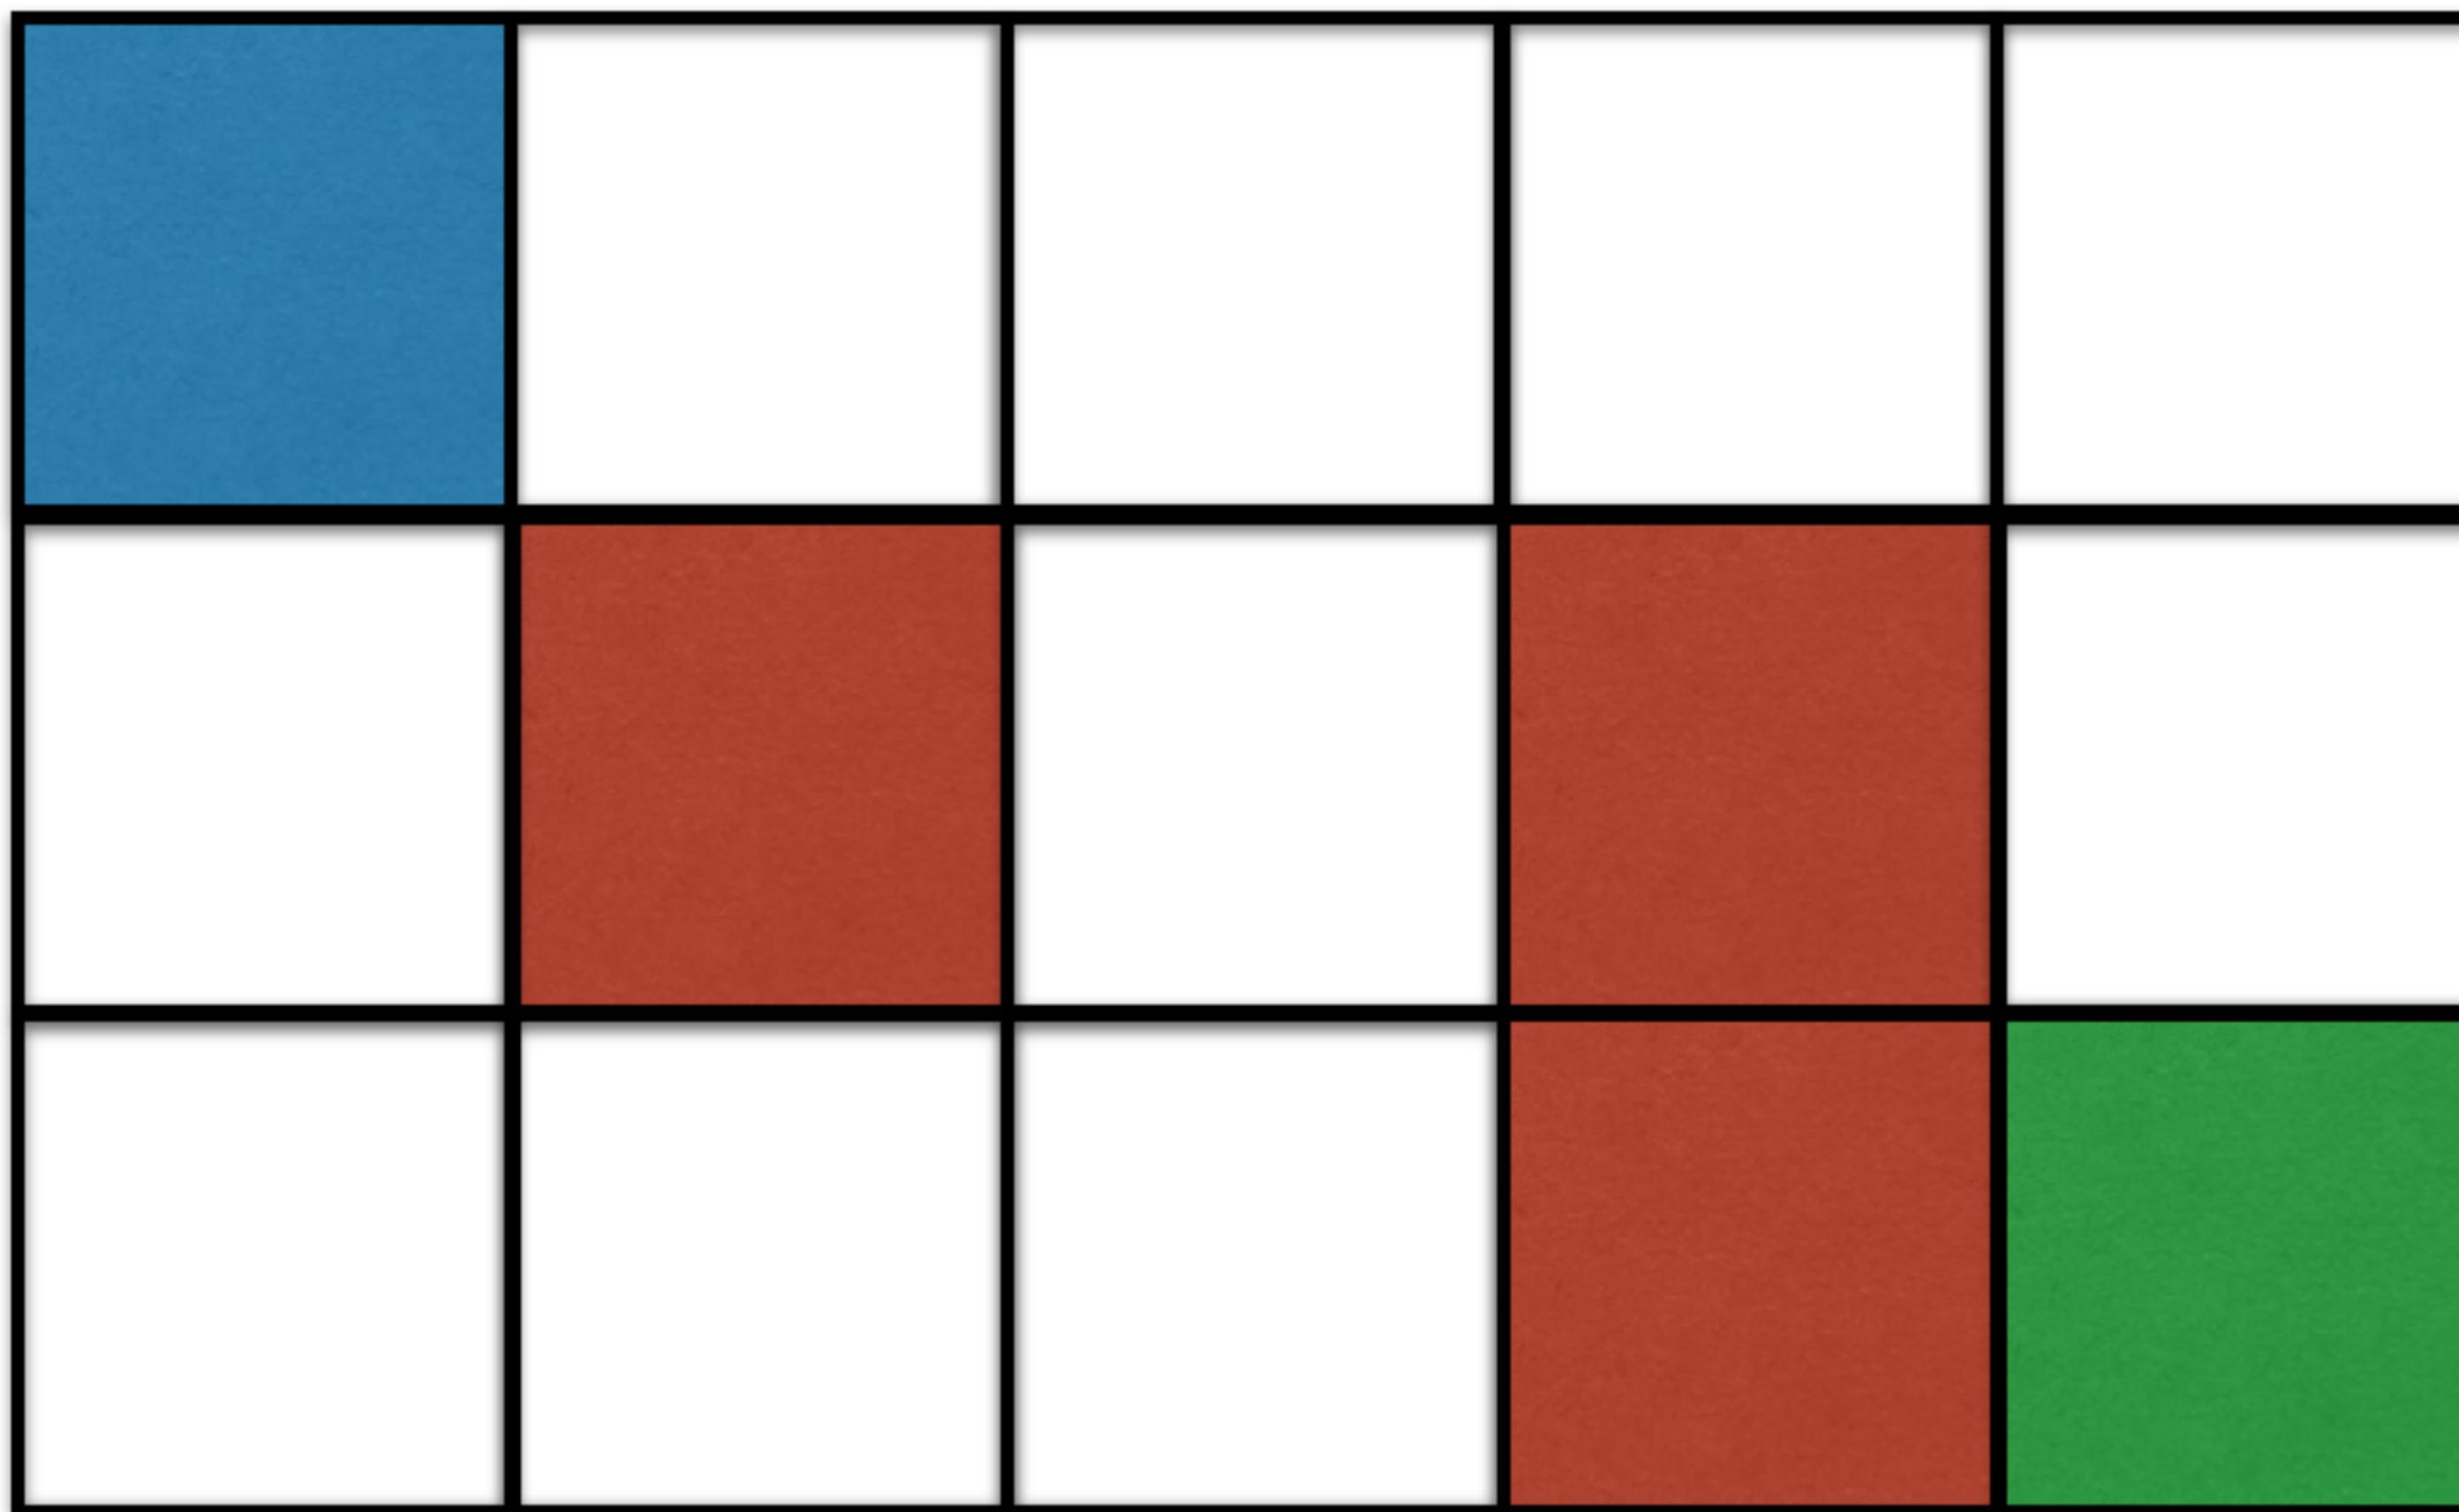
```
def V(s):  
    reward = R(s)  
    if is_goal(s):  
        return reward  
    return reward +  
        sum([P(s, pi(s), n_s) * V(n_s) for n_s in states])
```

# Gridworld Value Function Example

---

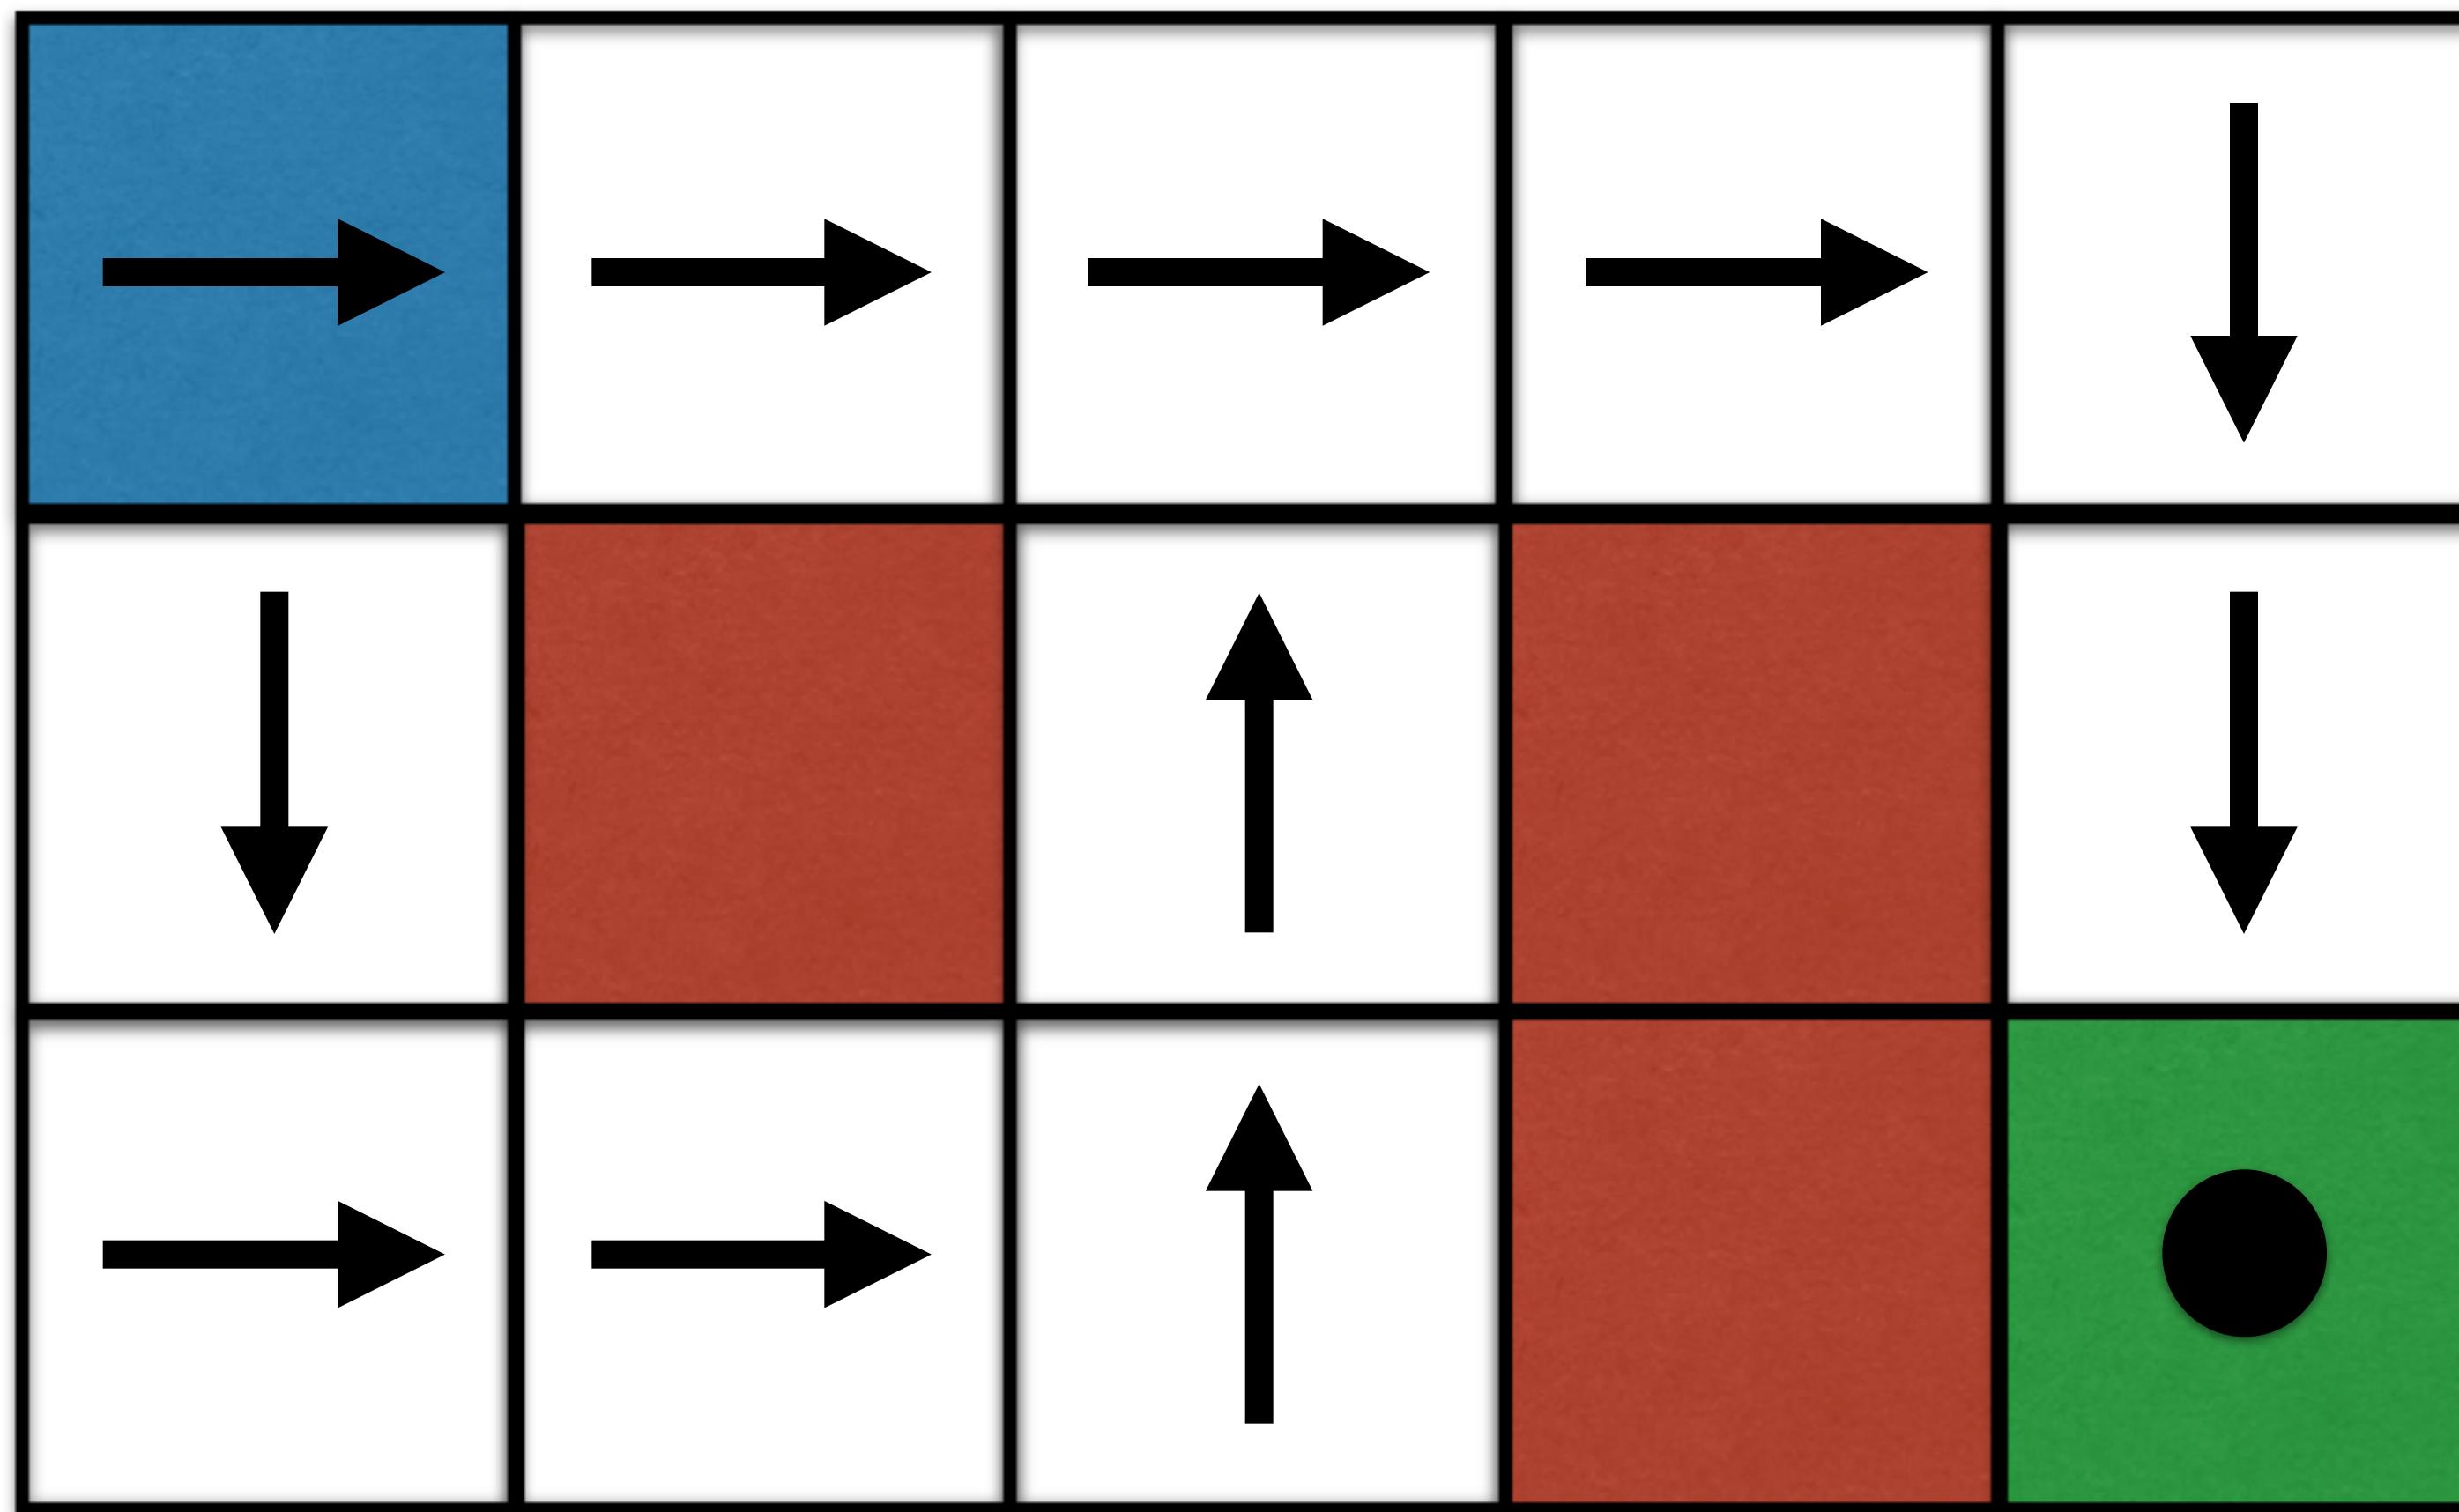
# Gridworld Value Function Example

---



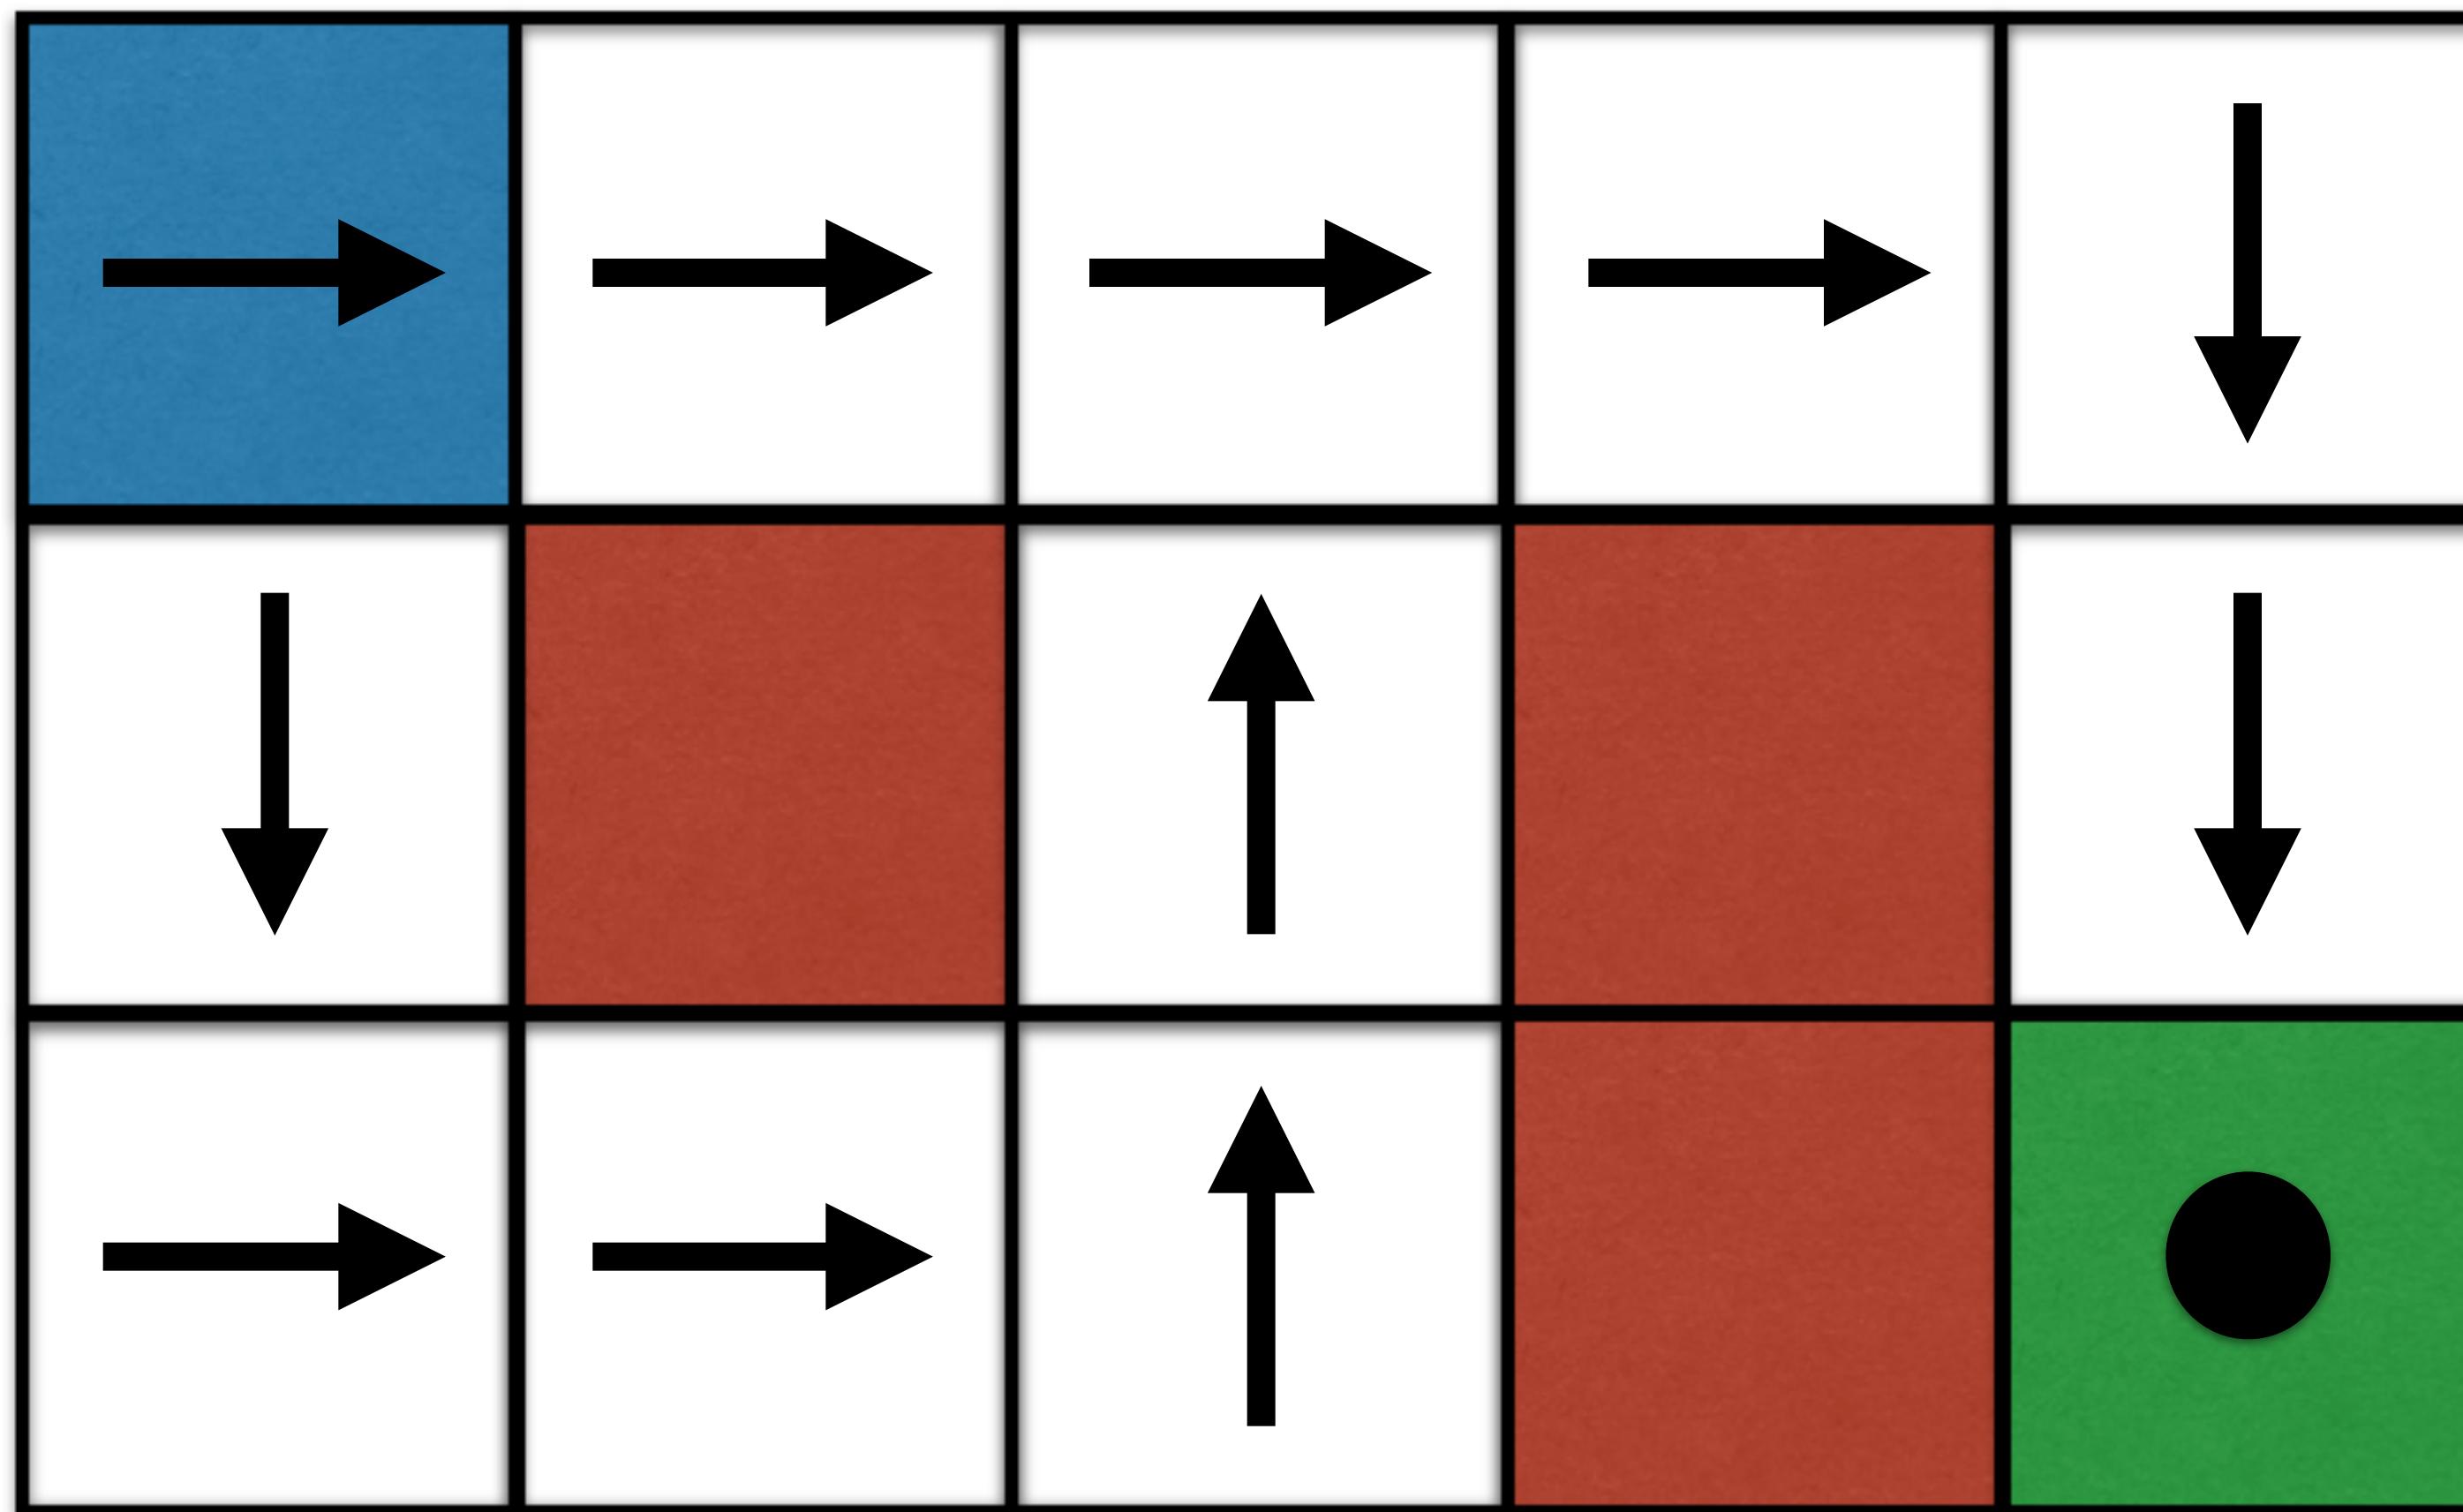
# Gridworld Value Function Example

---



# Gridworld Value Function Example

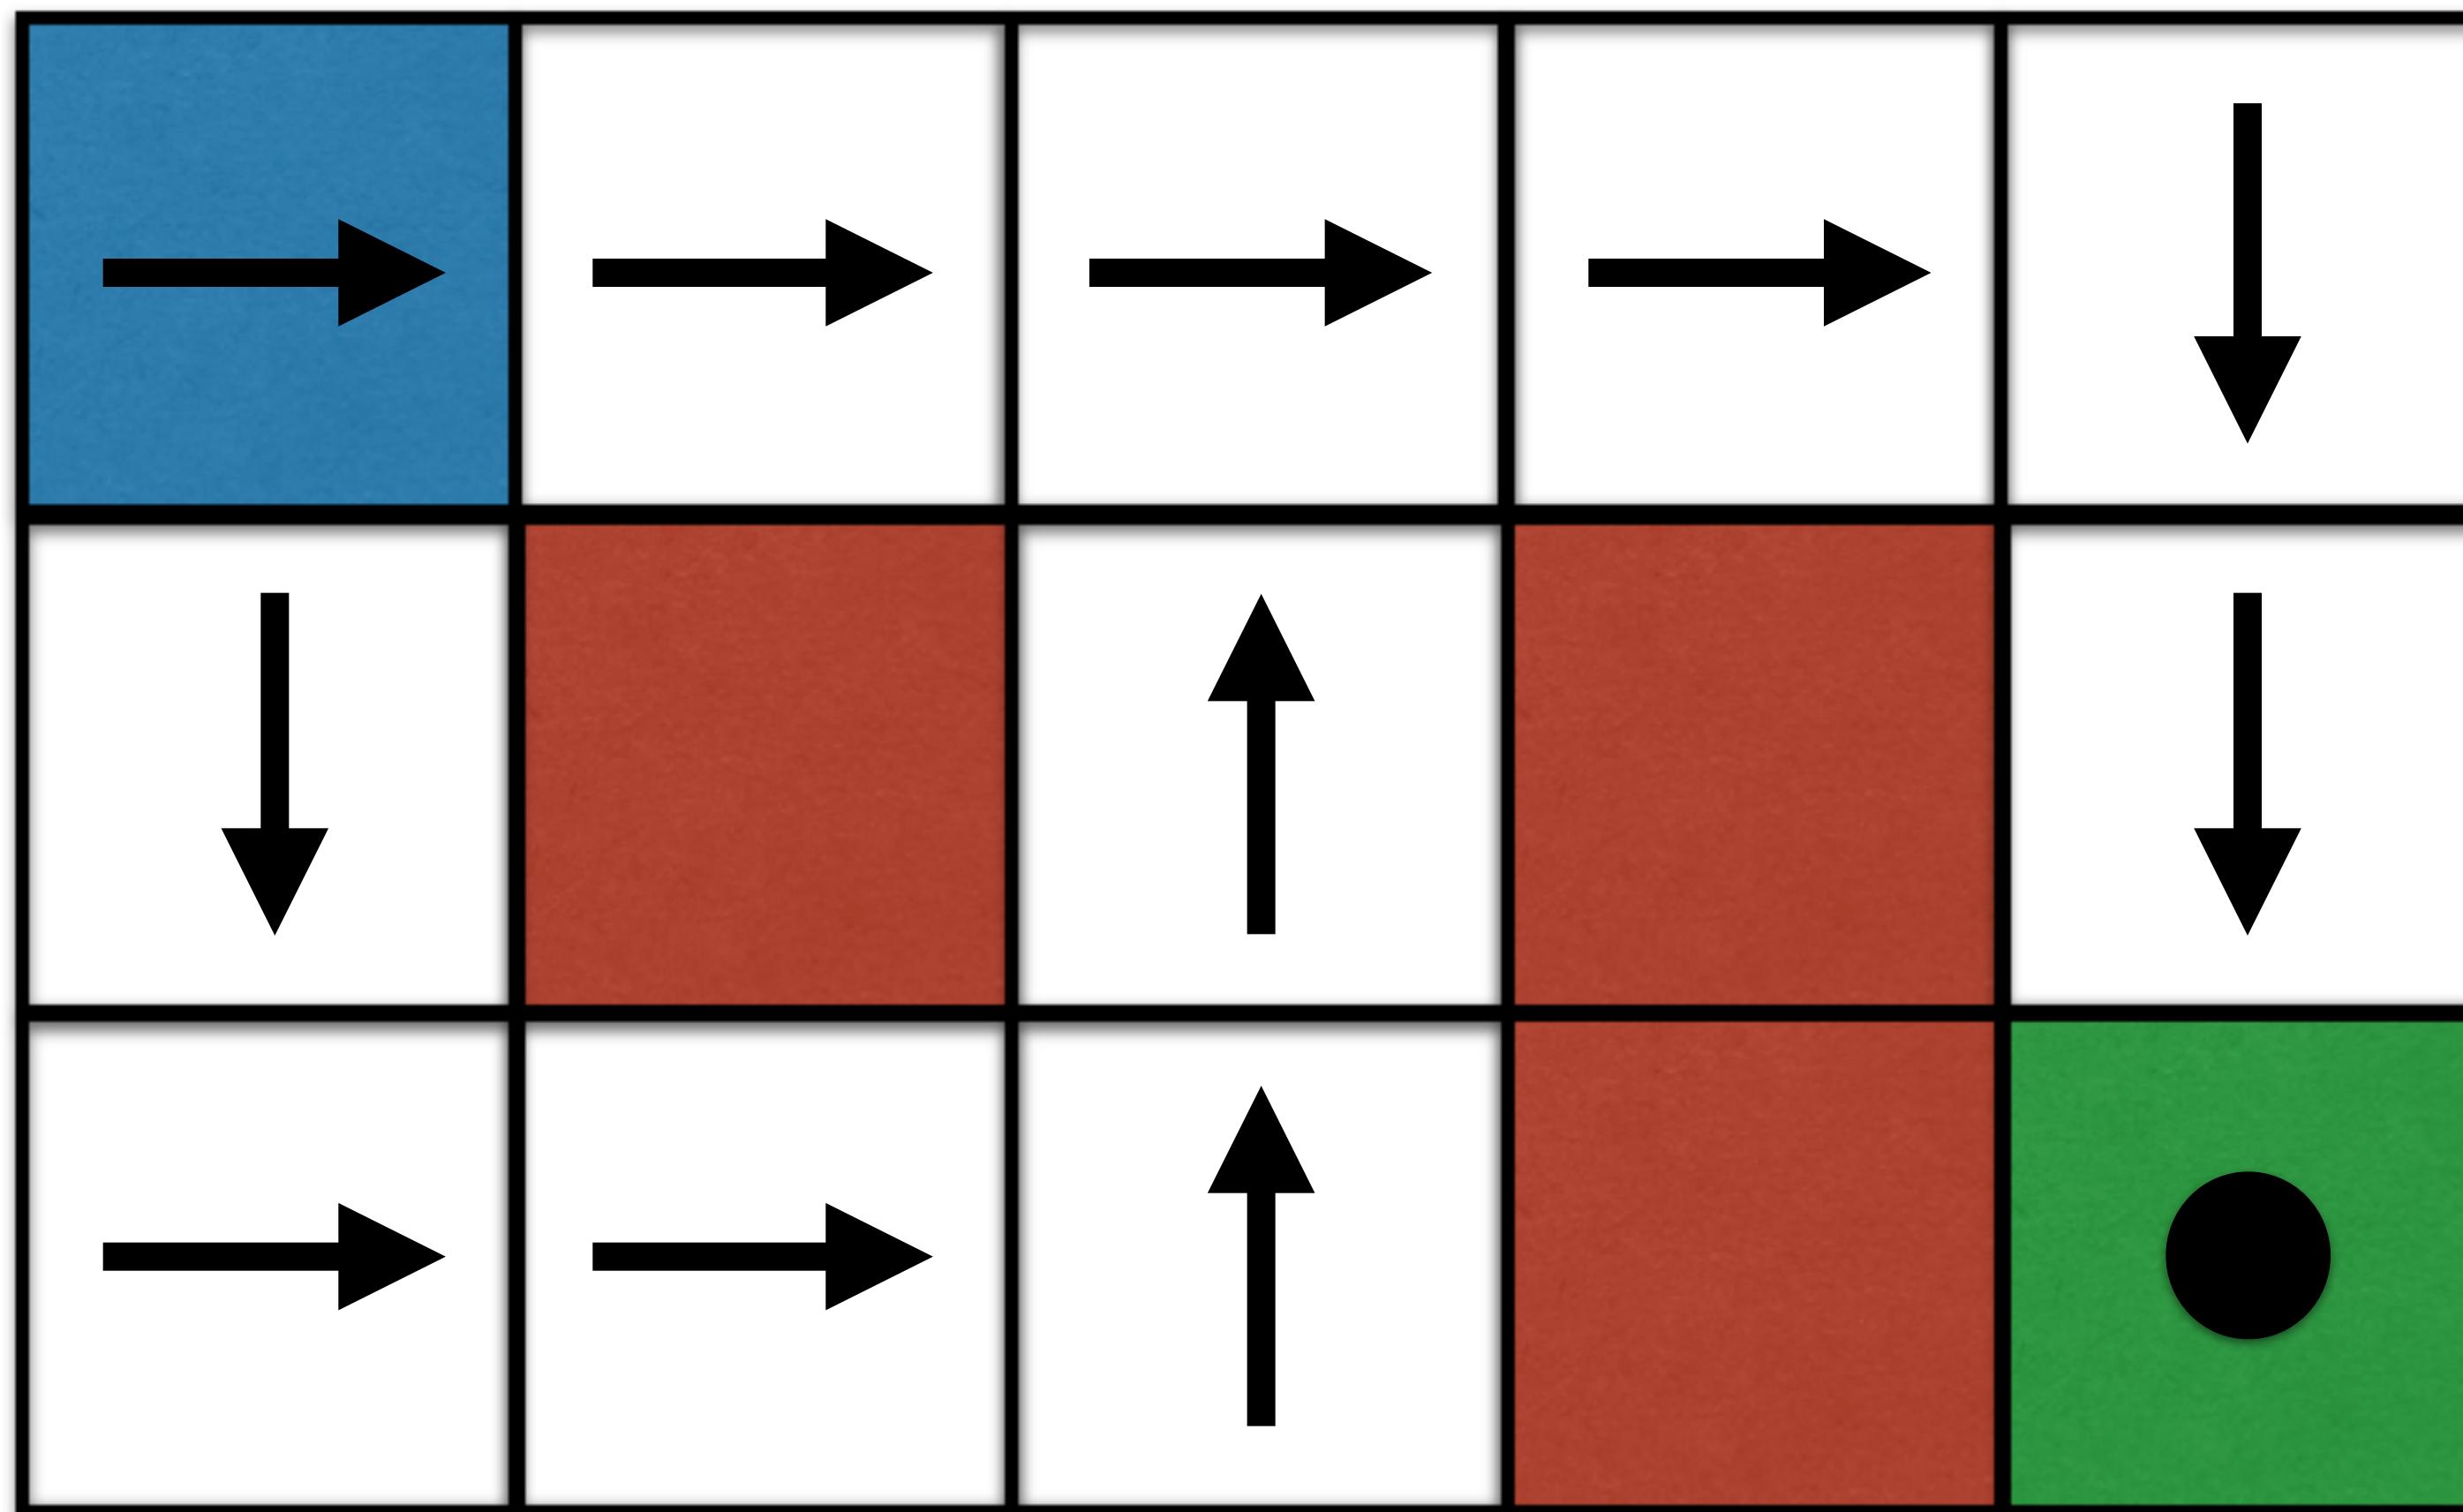
---



- Arrows denote the policy  $\pi$ .

# Gridworld Value Function Example

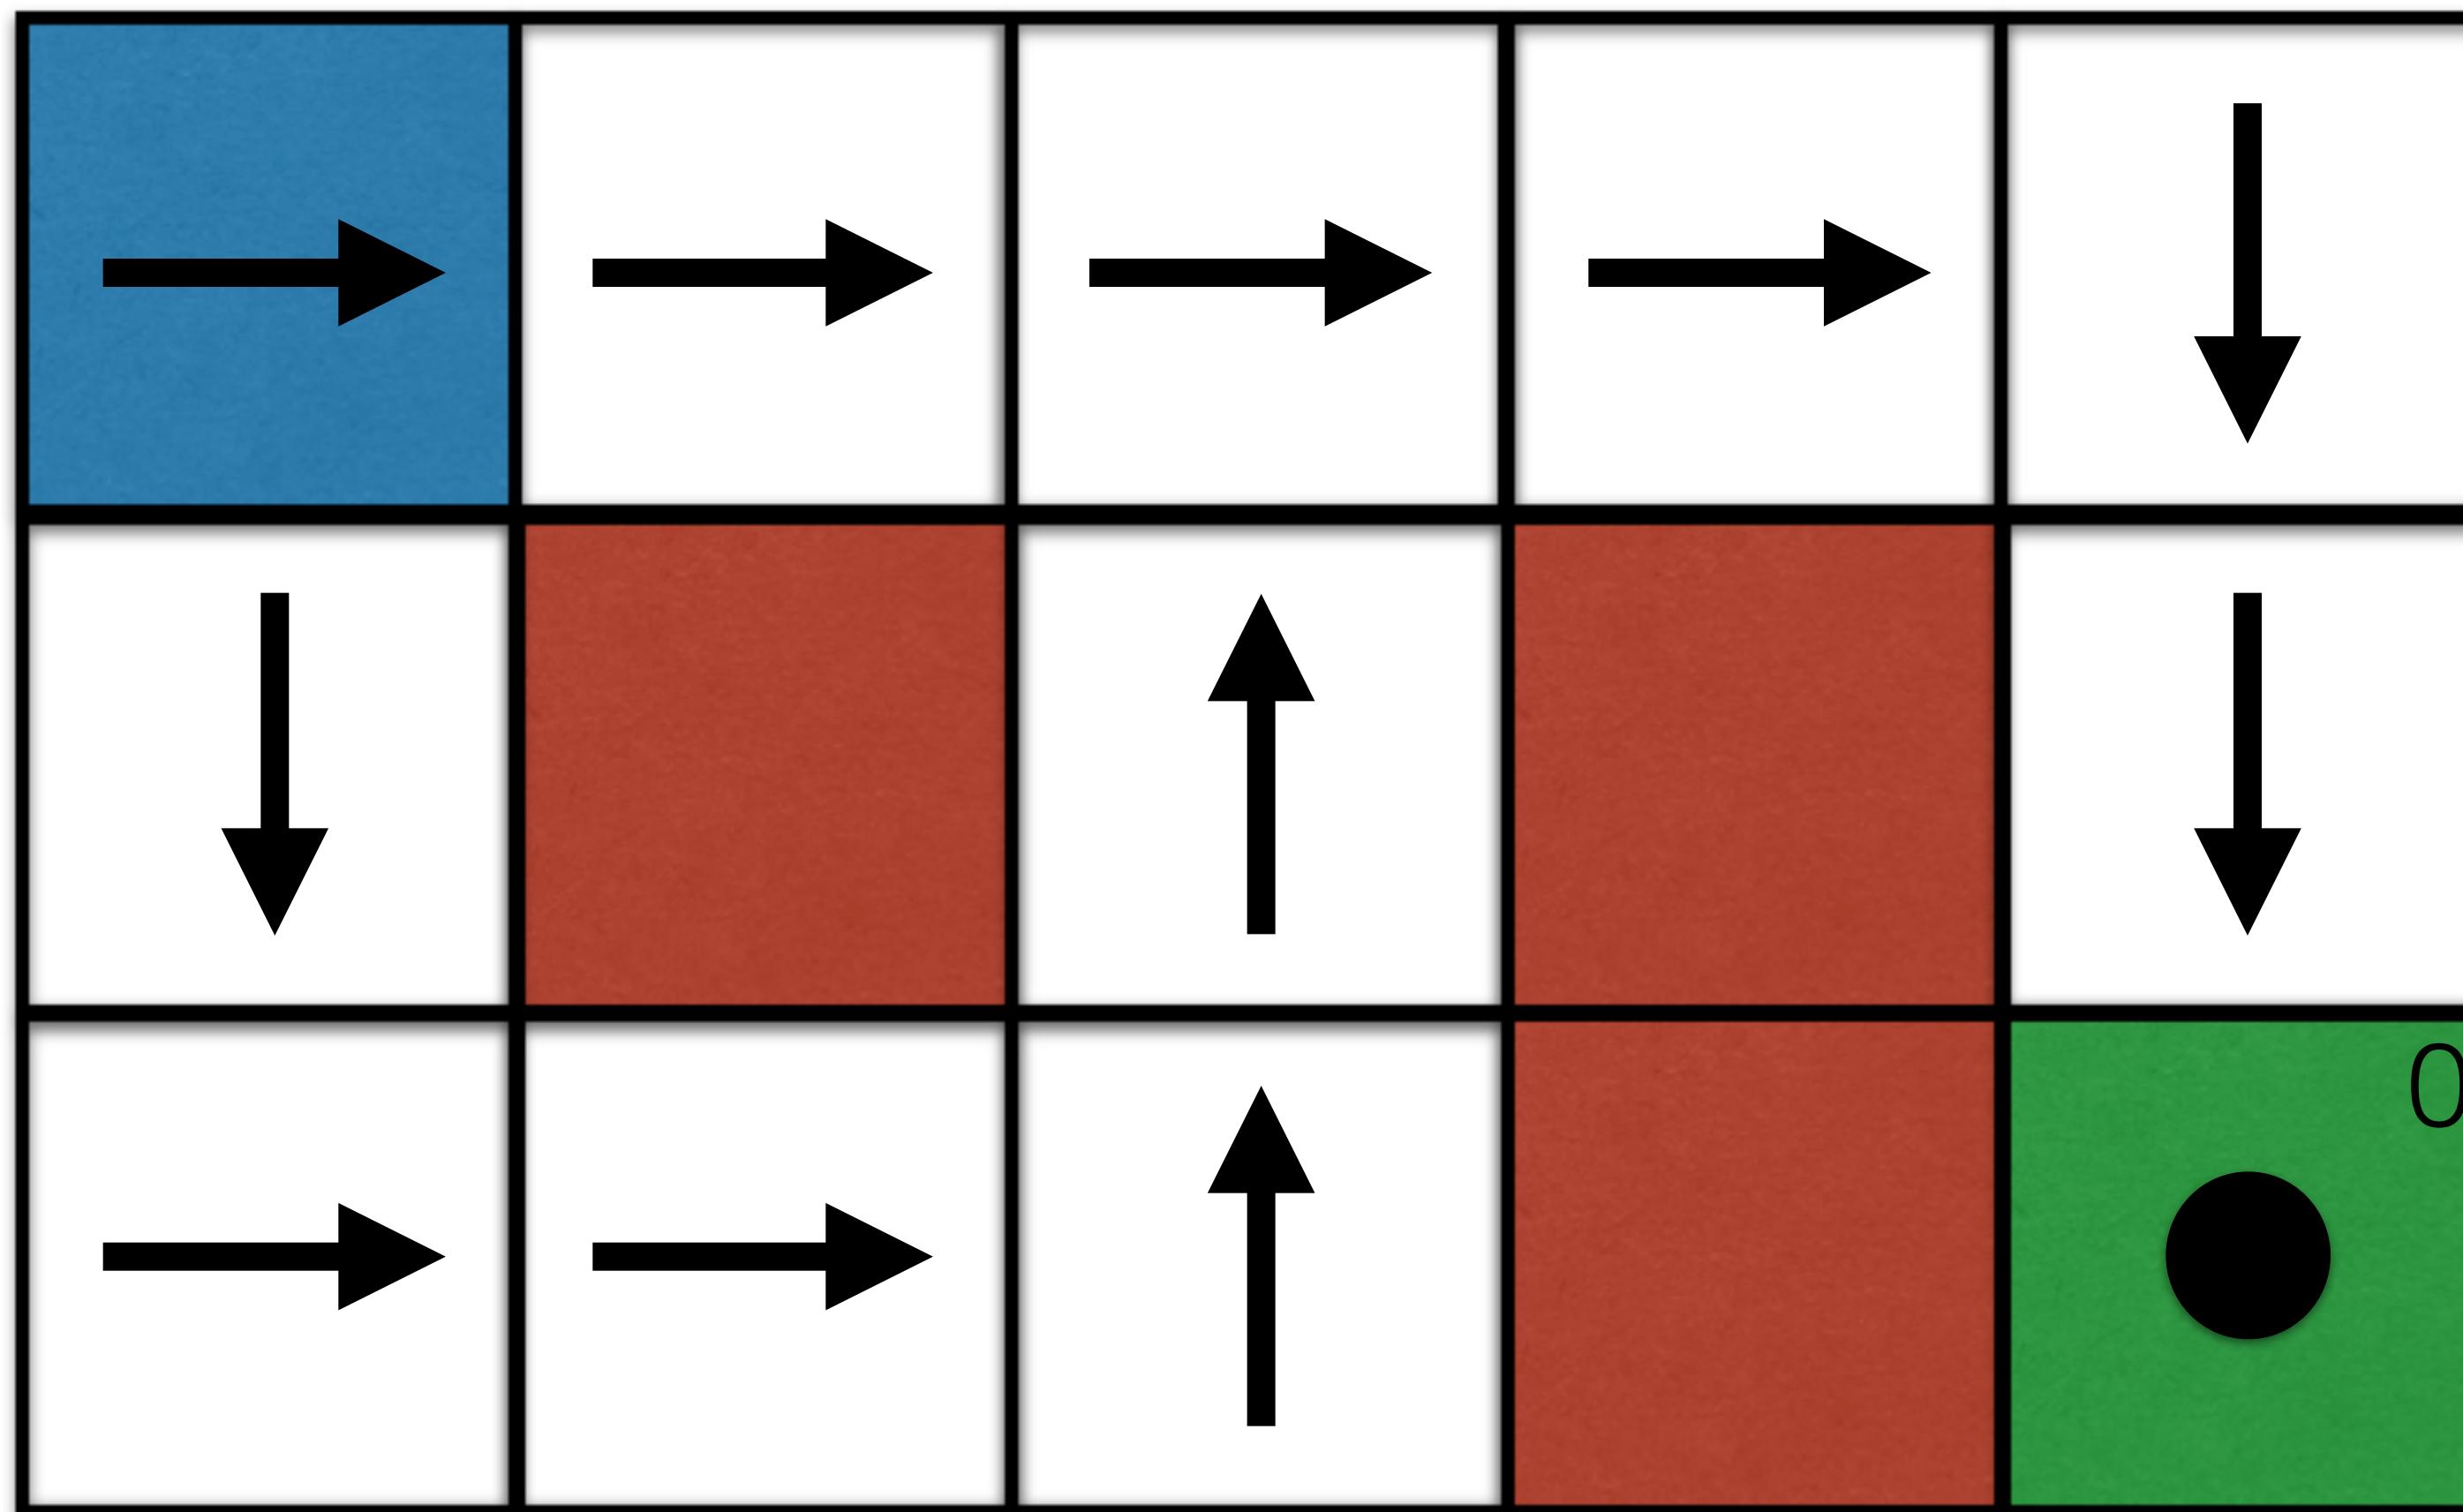
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

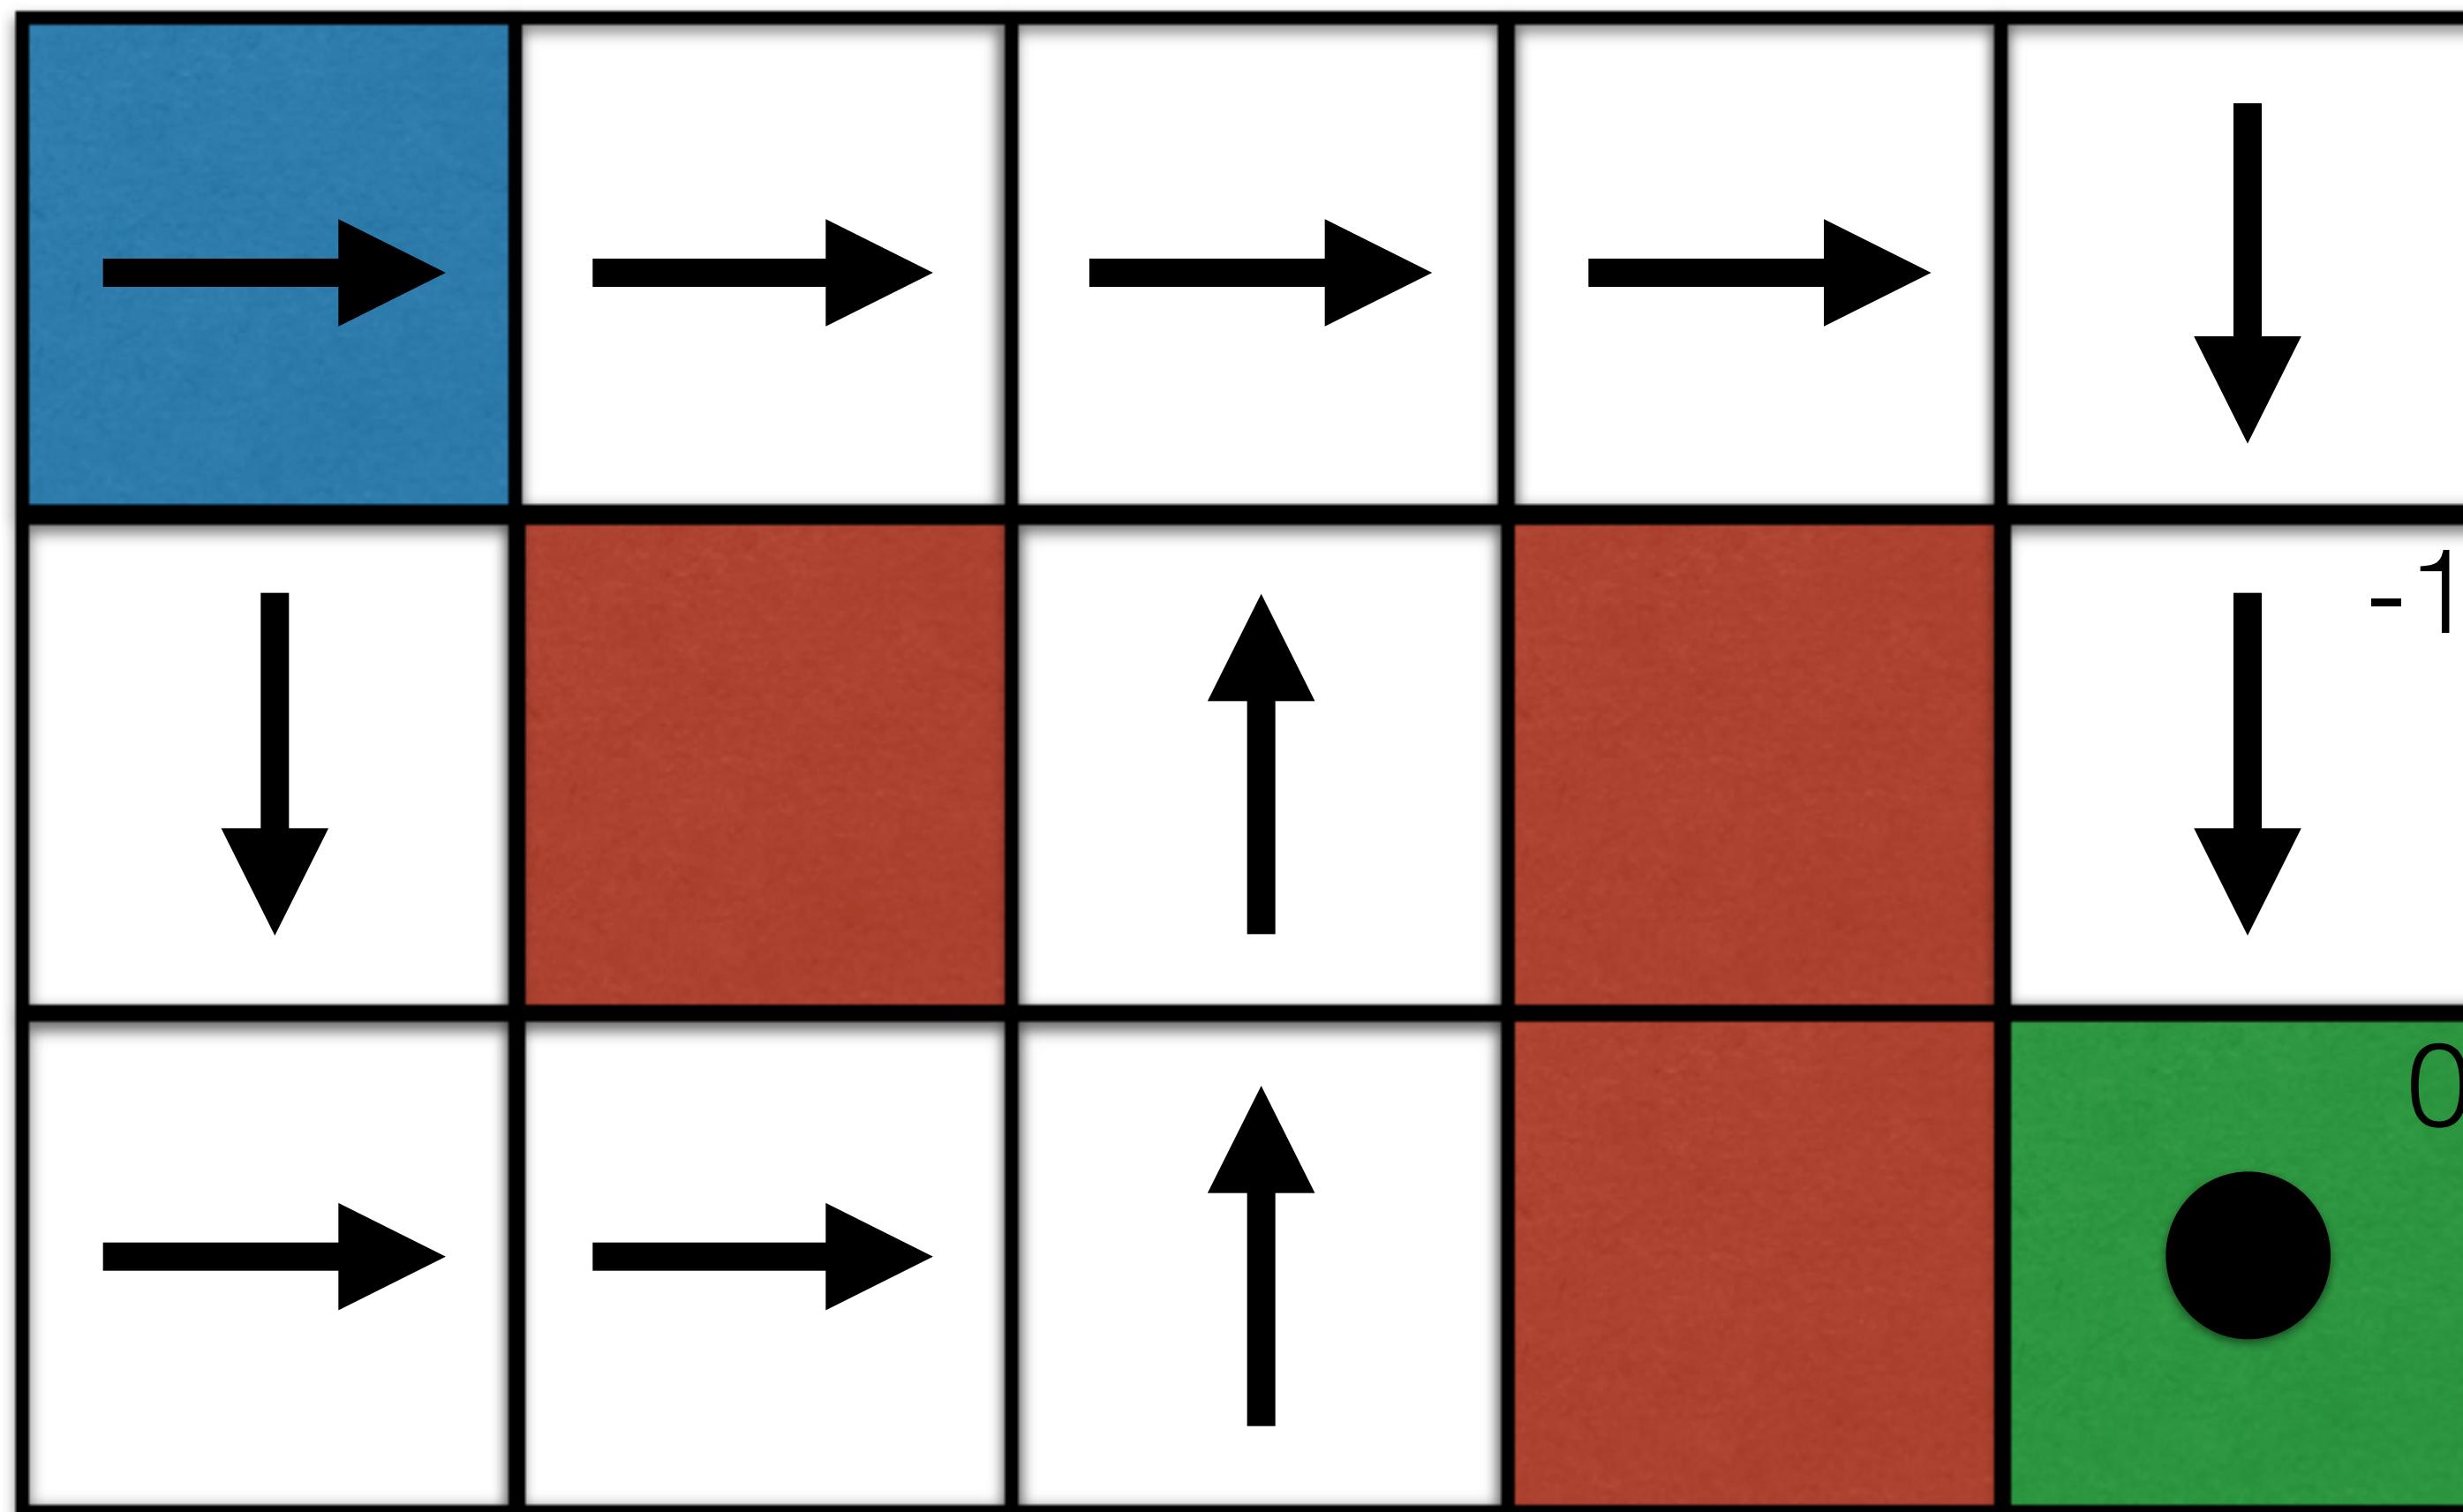
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

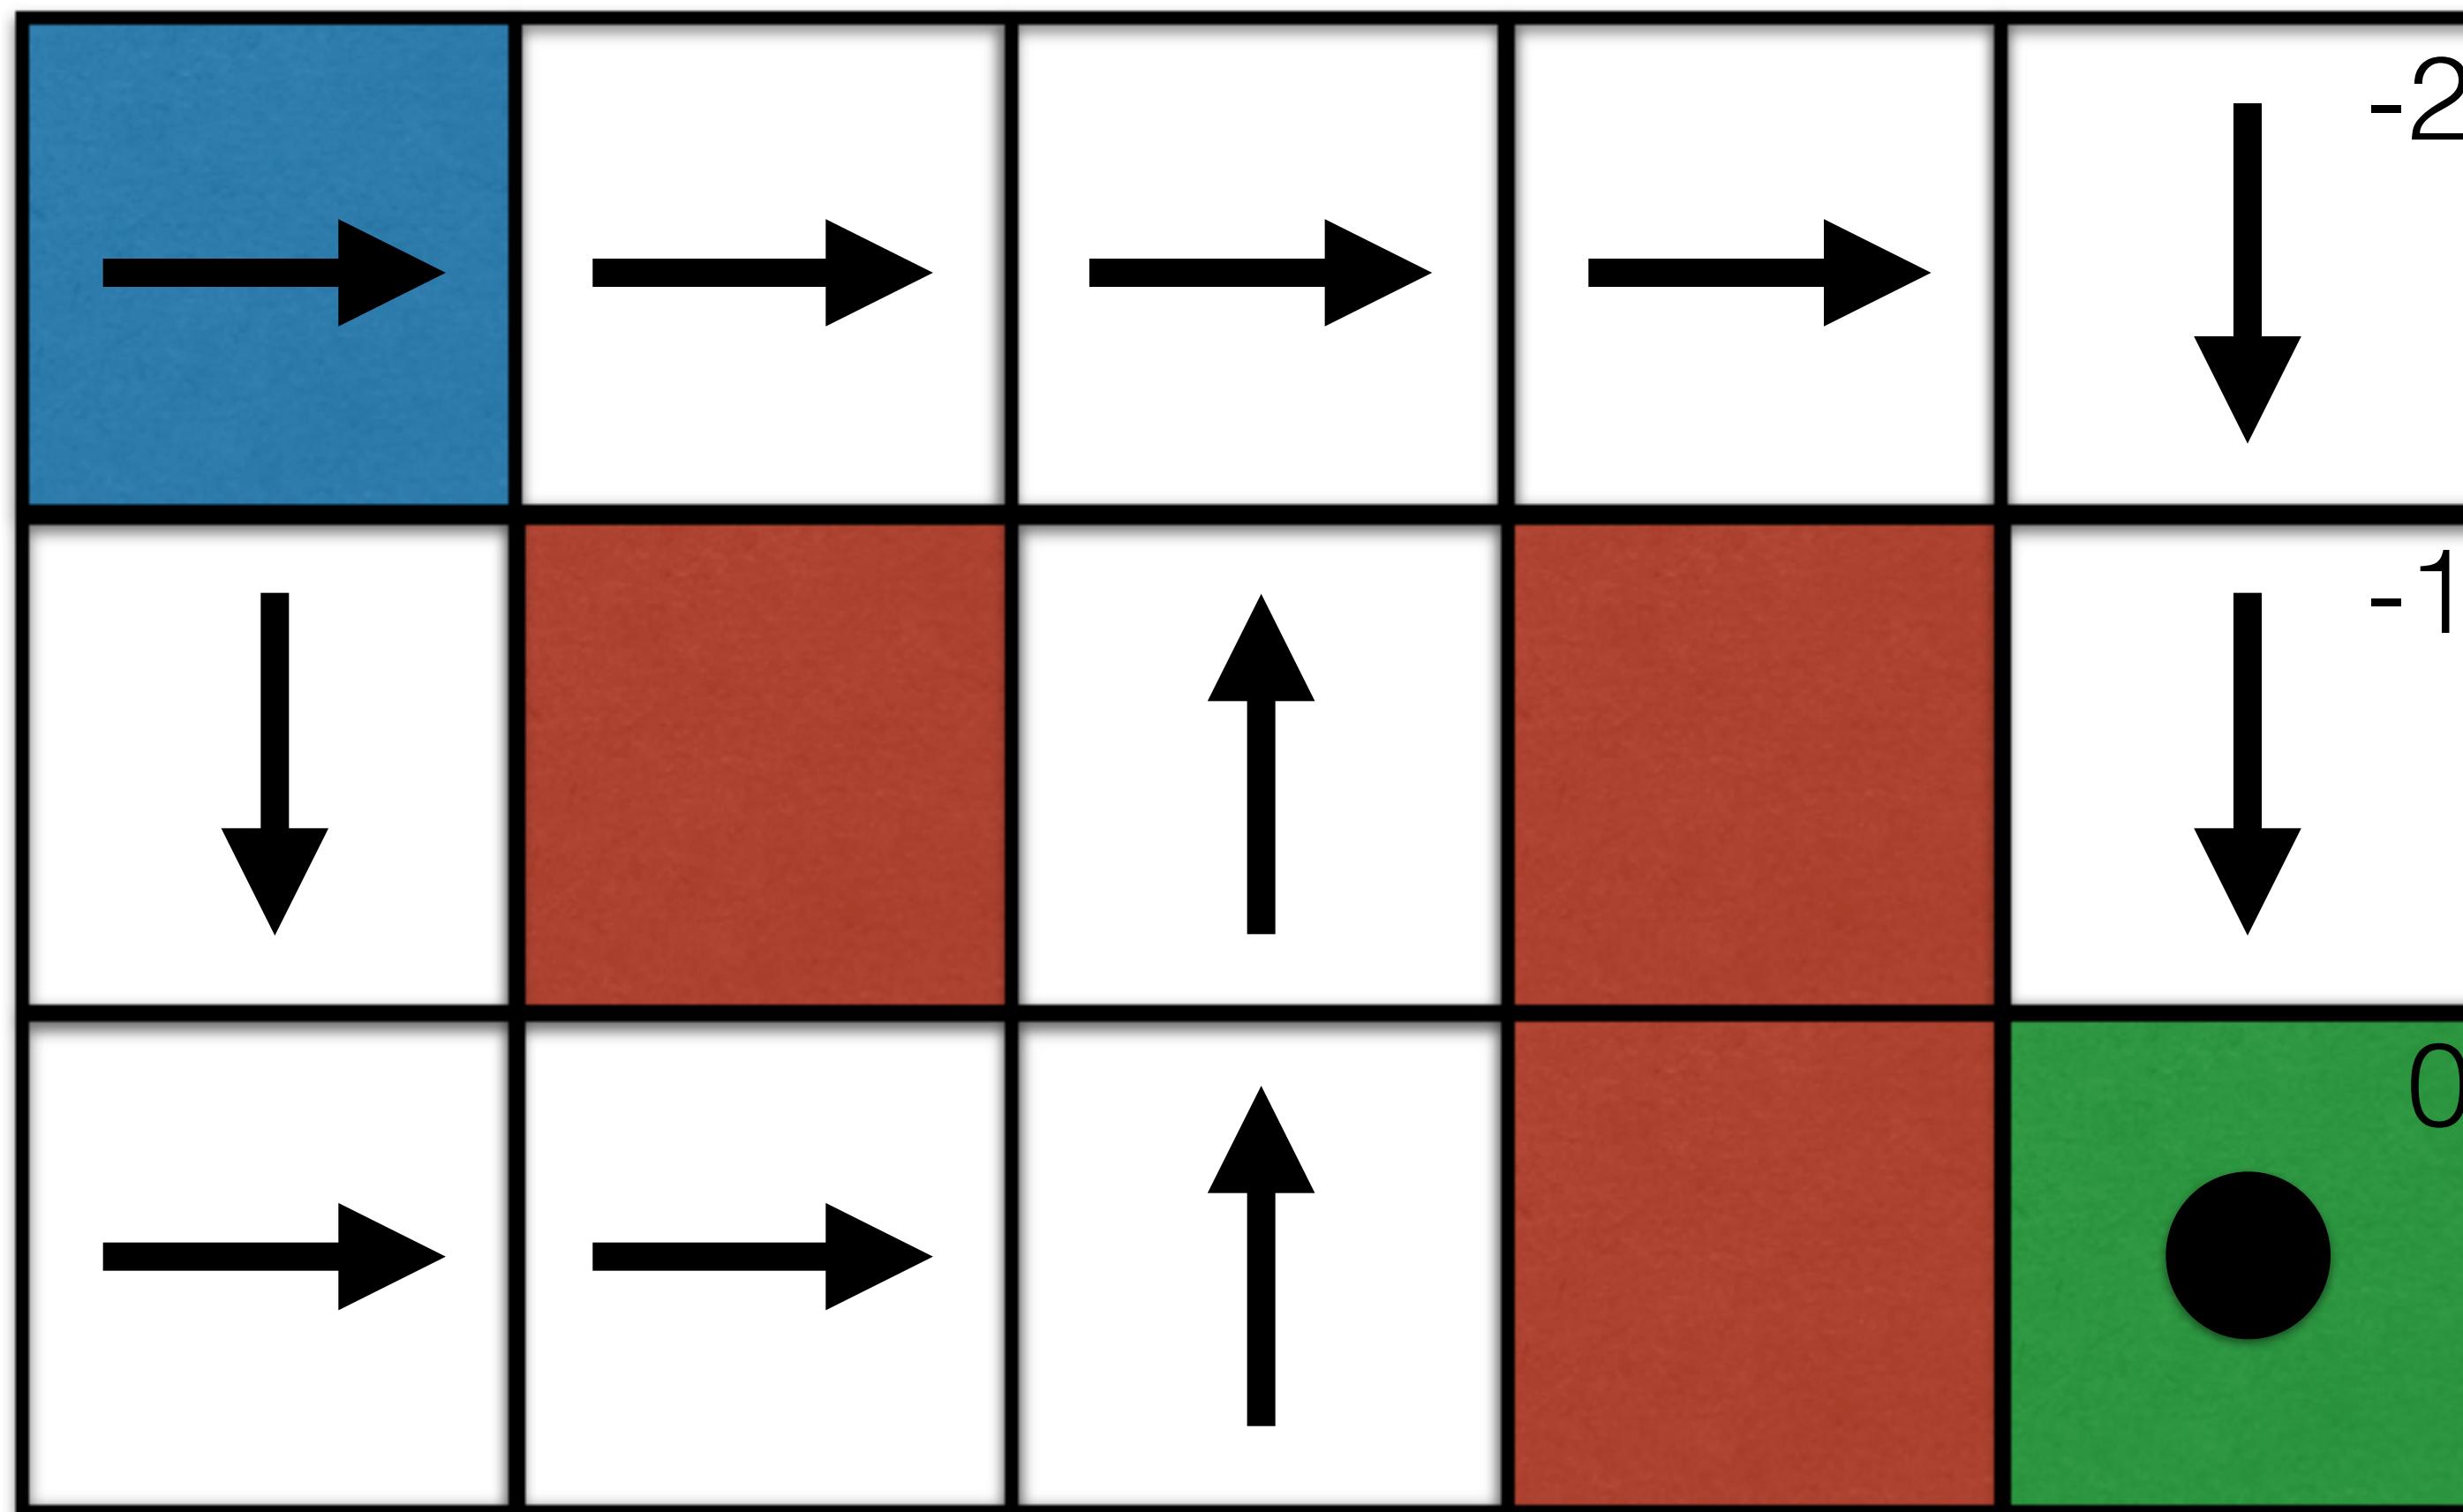
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

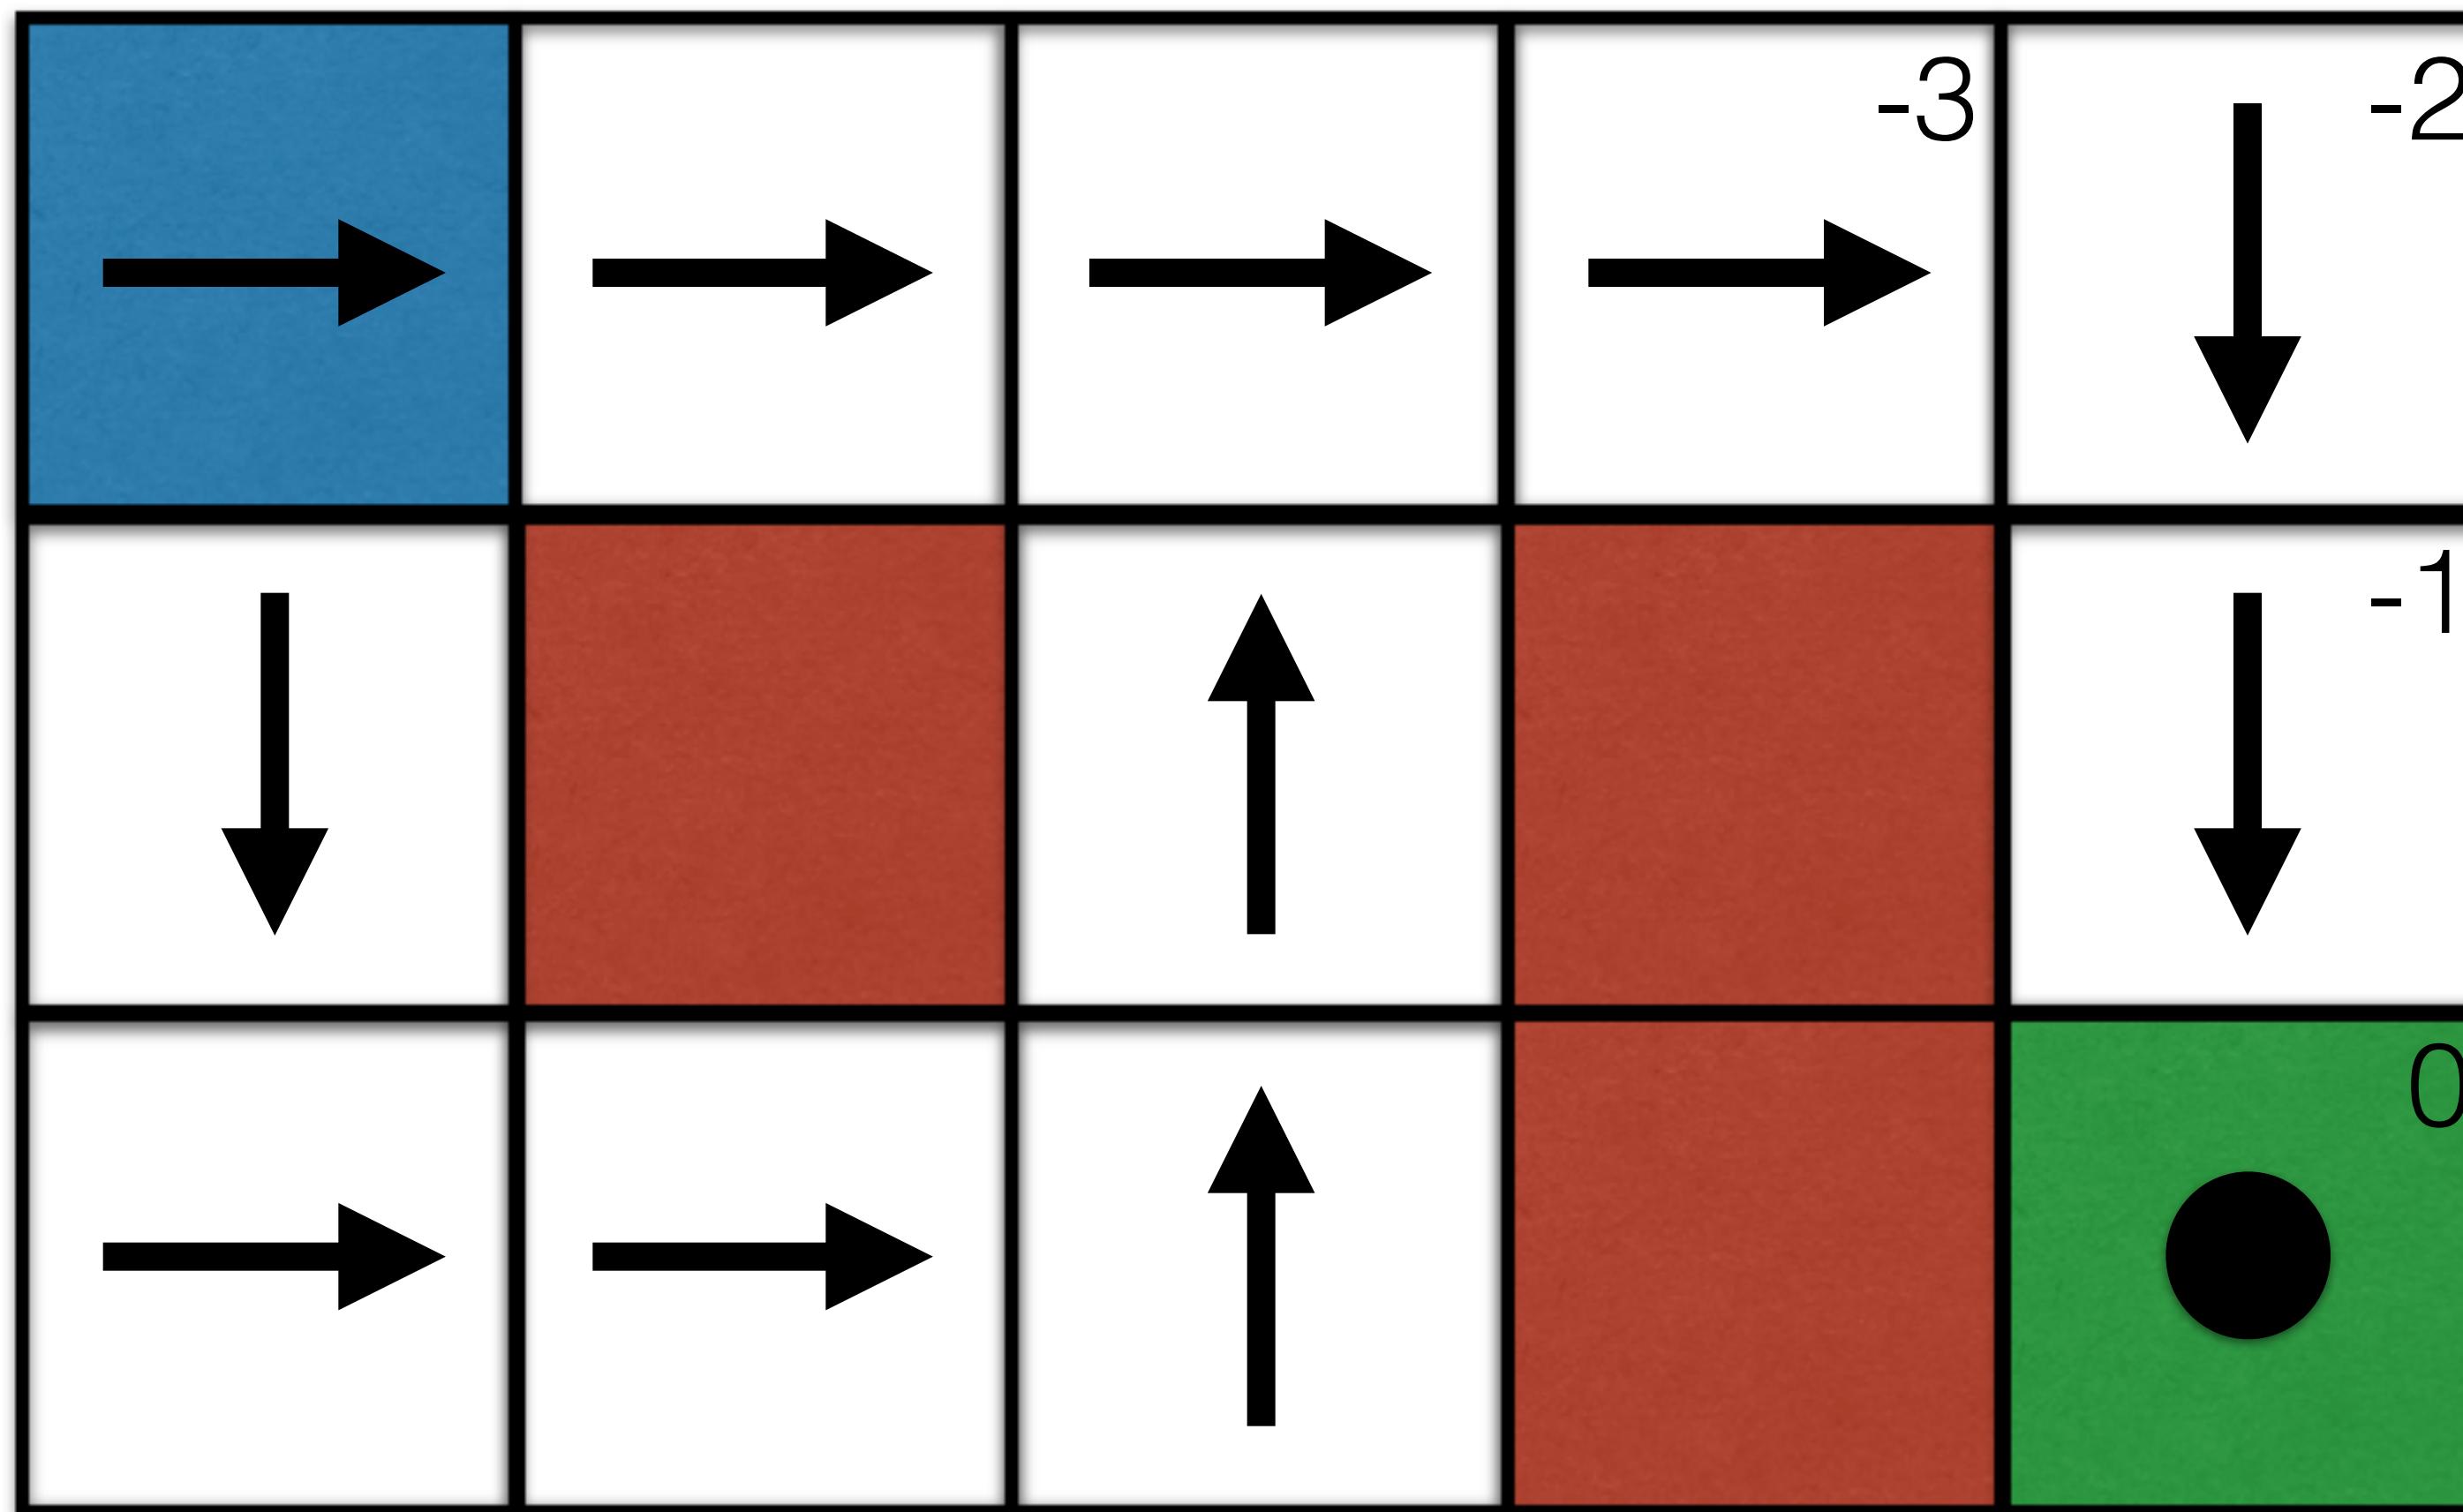
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

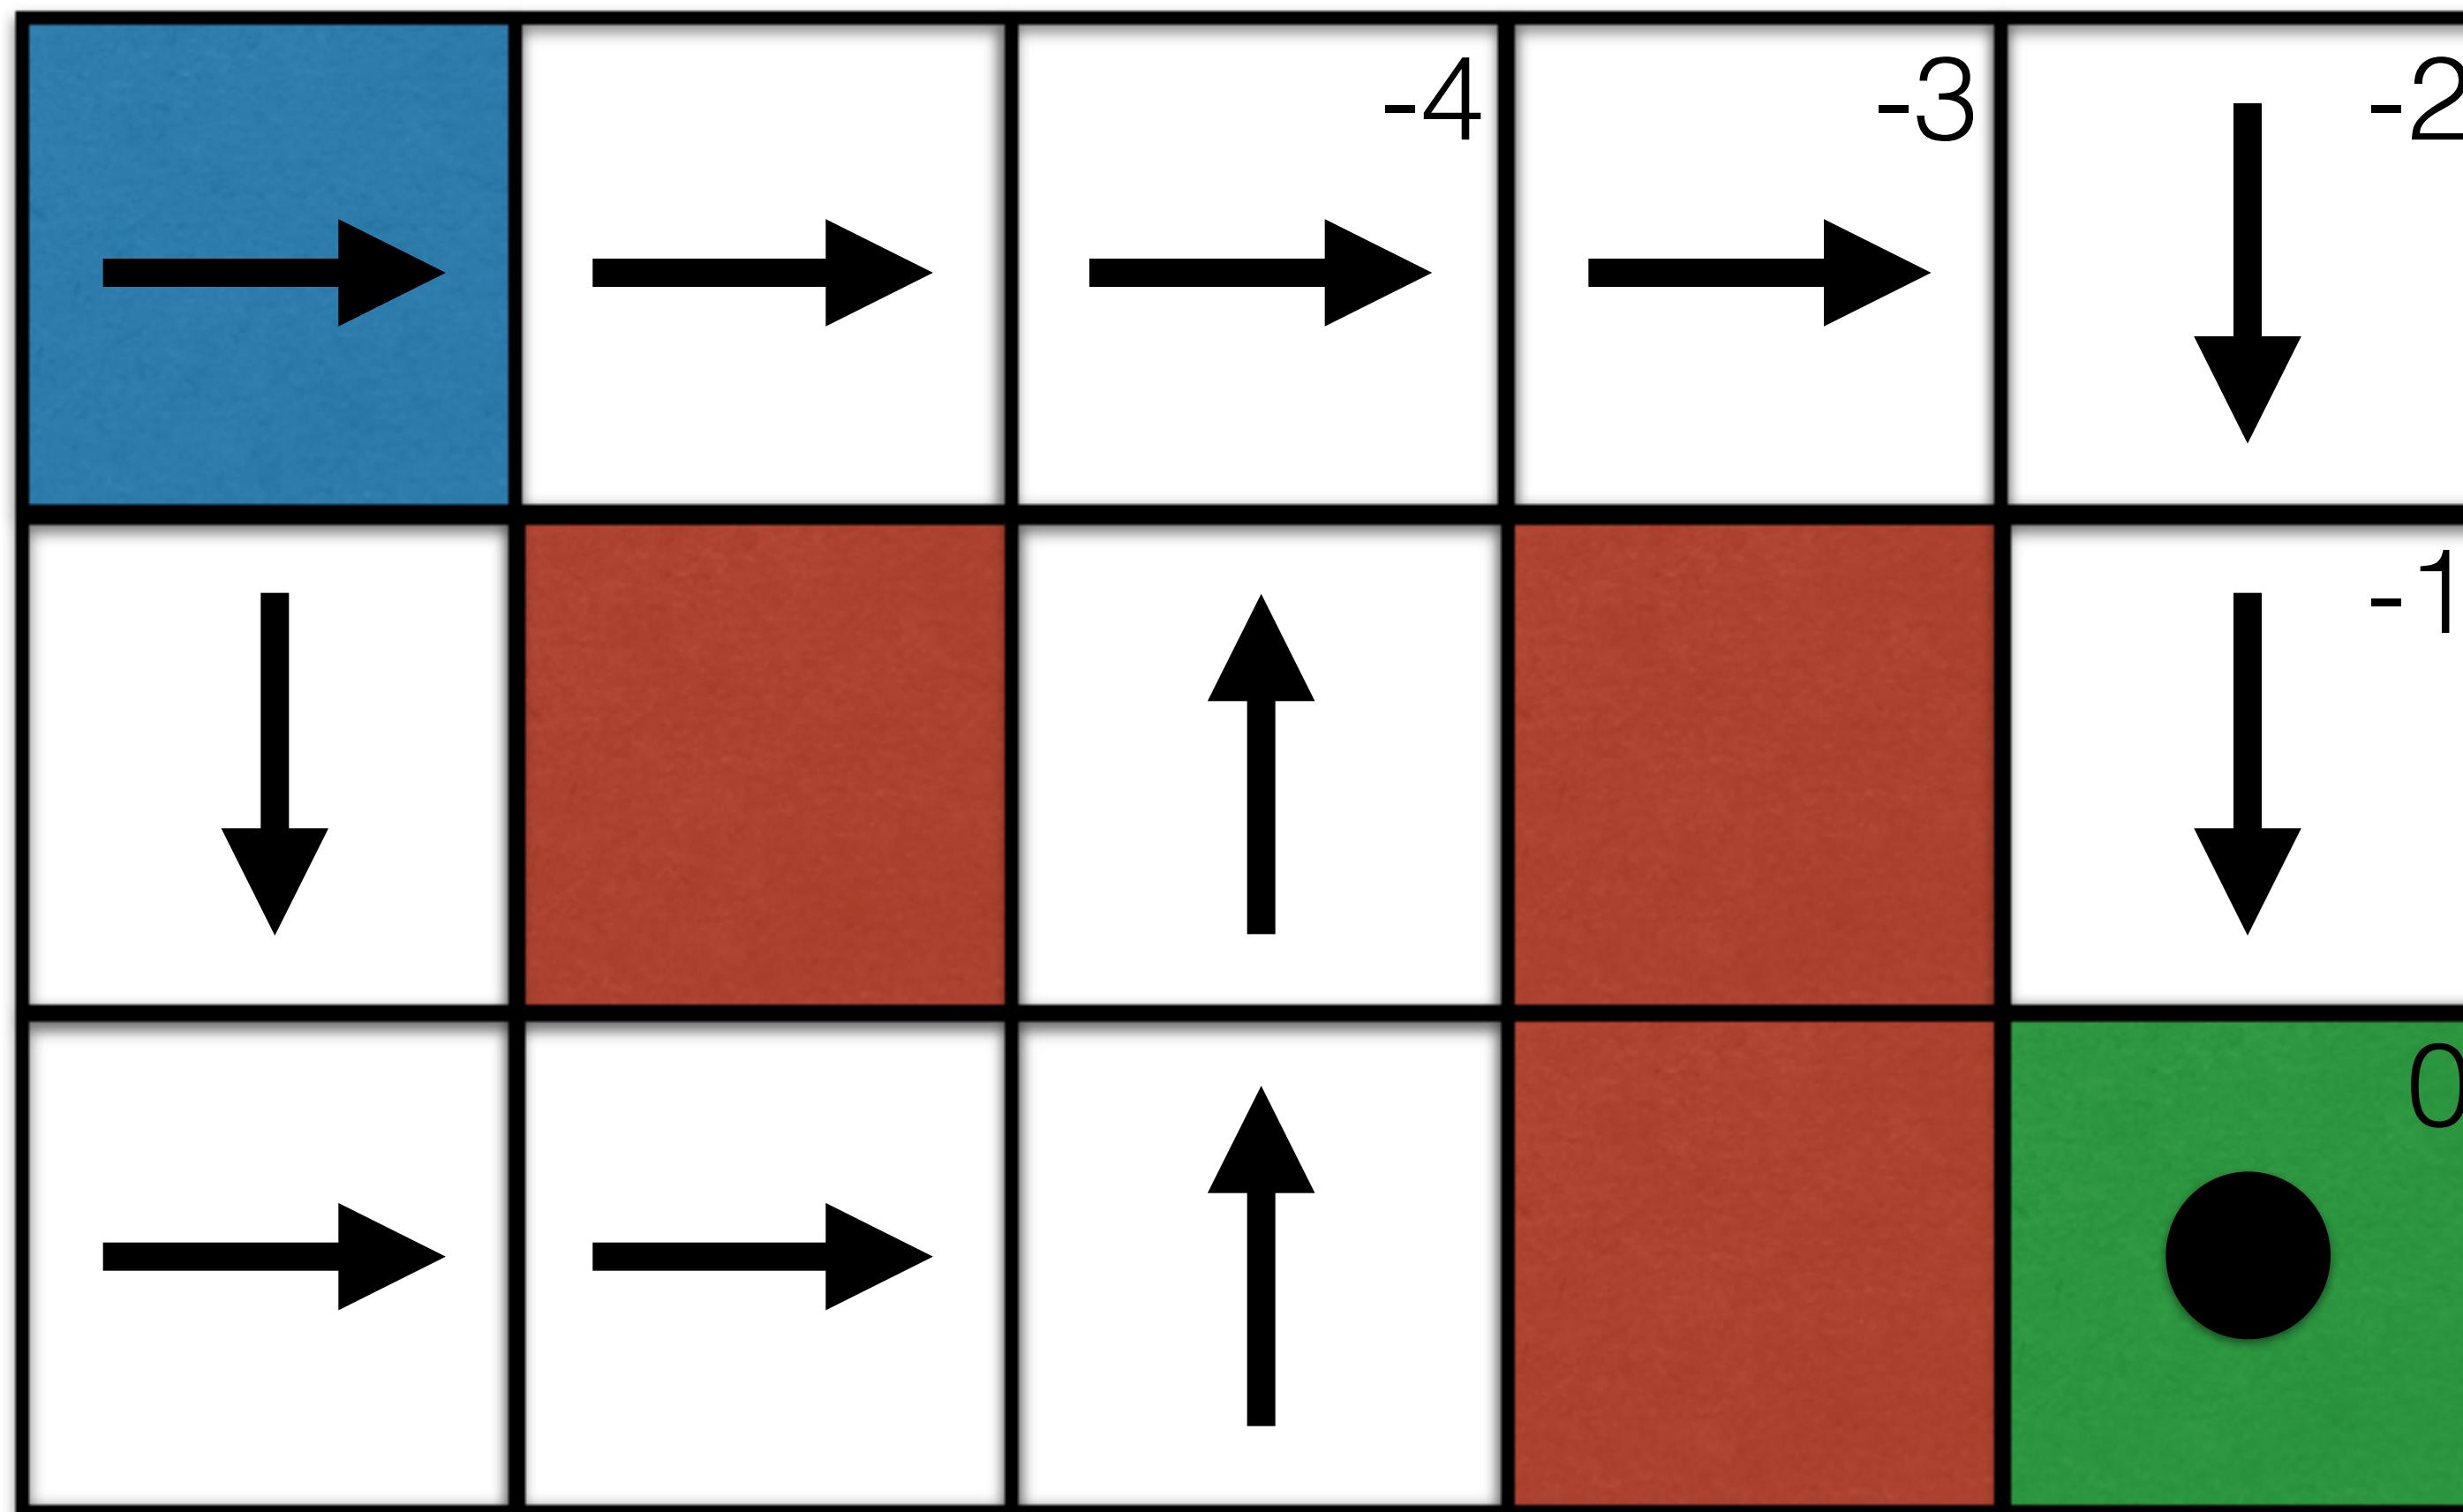
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

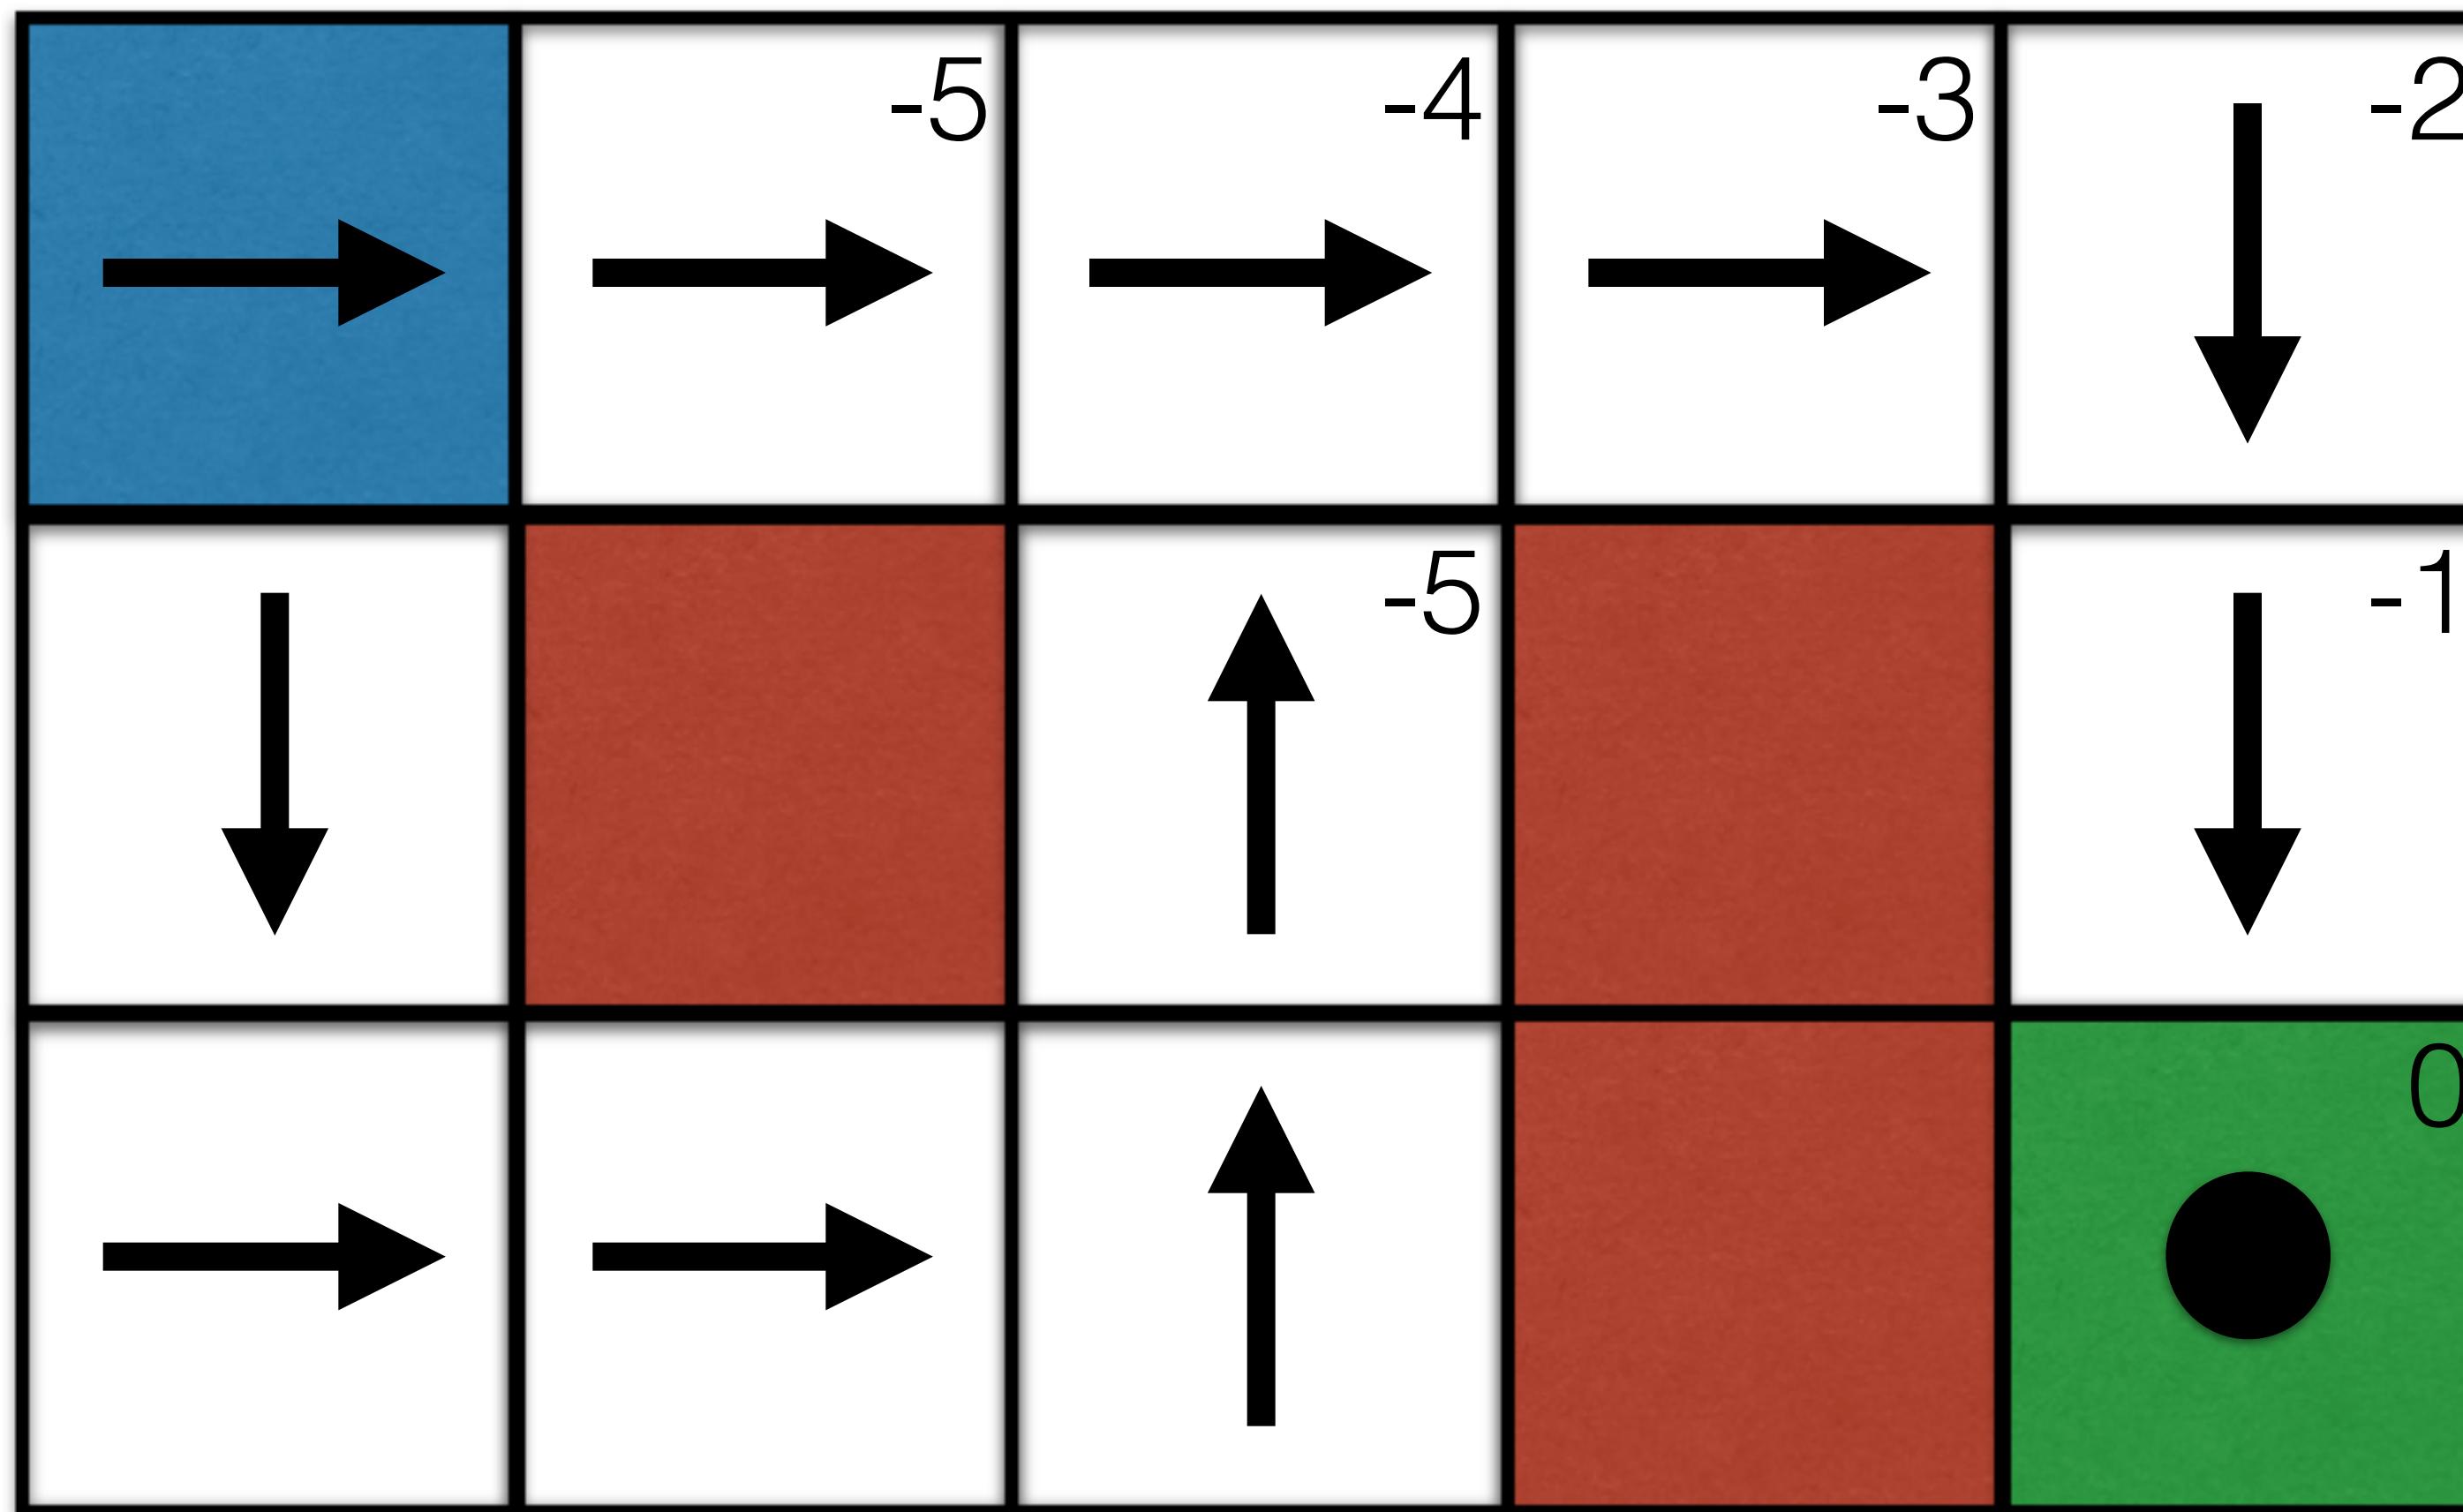
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

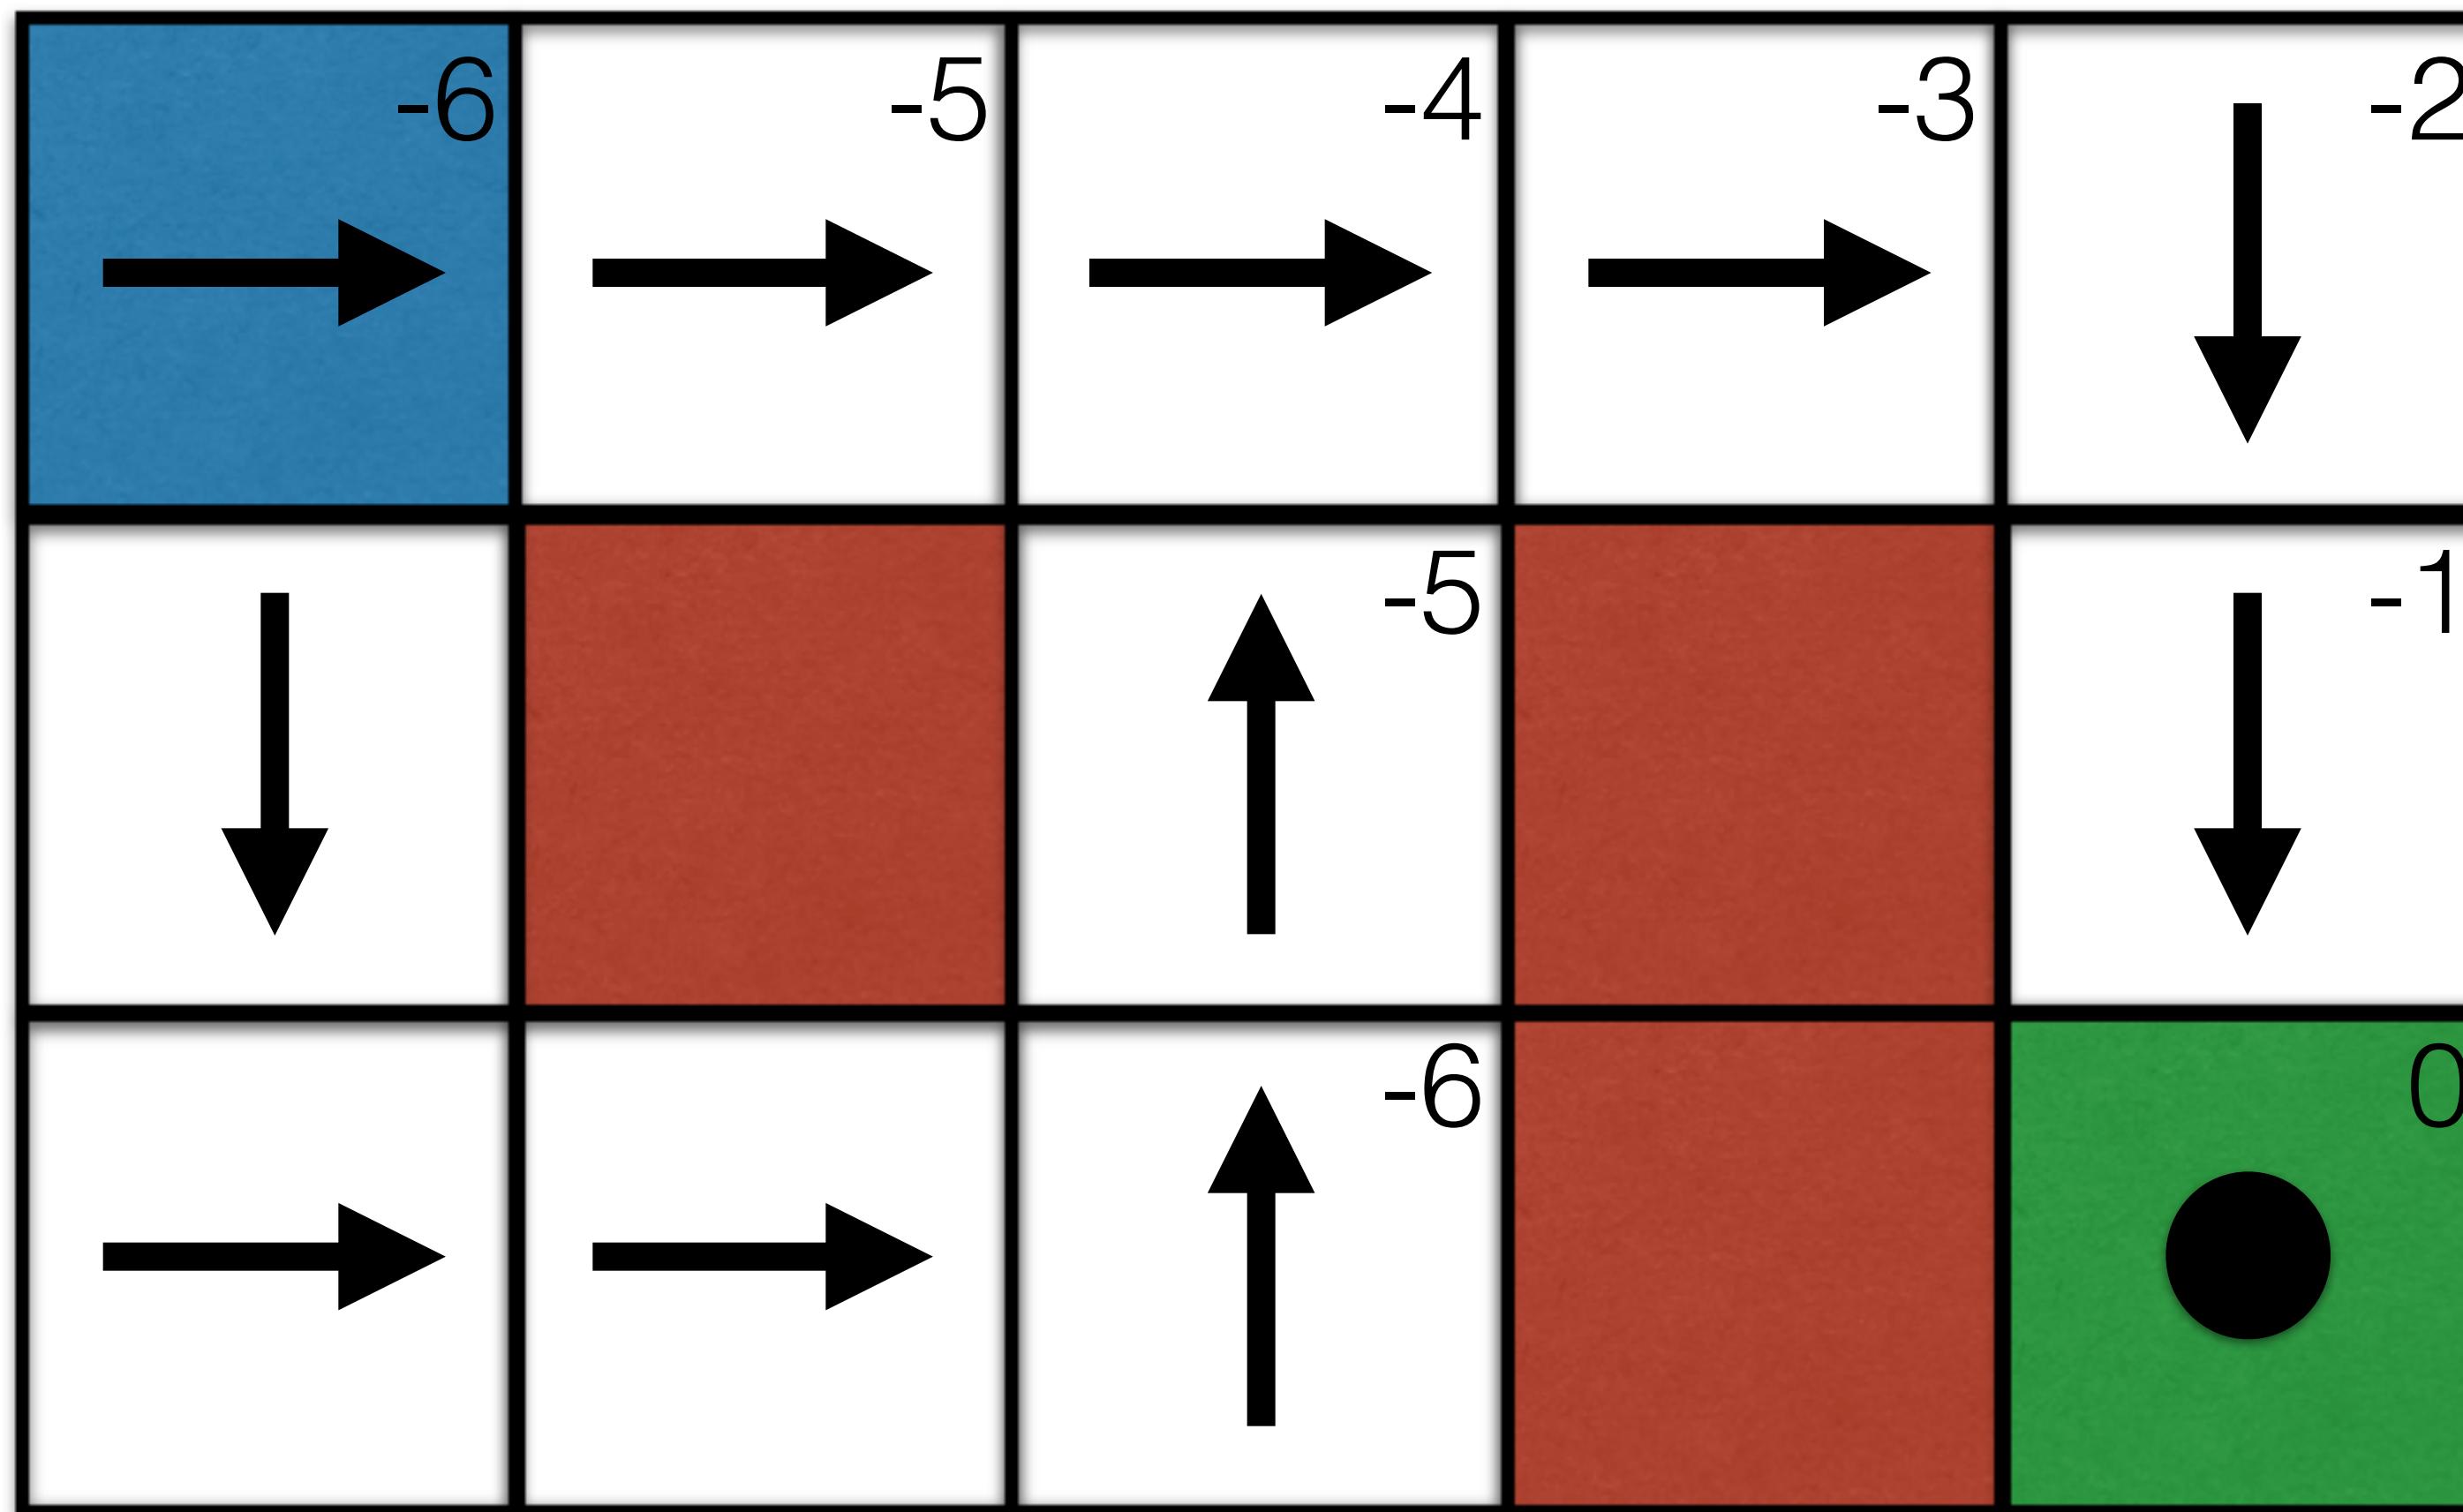
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

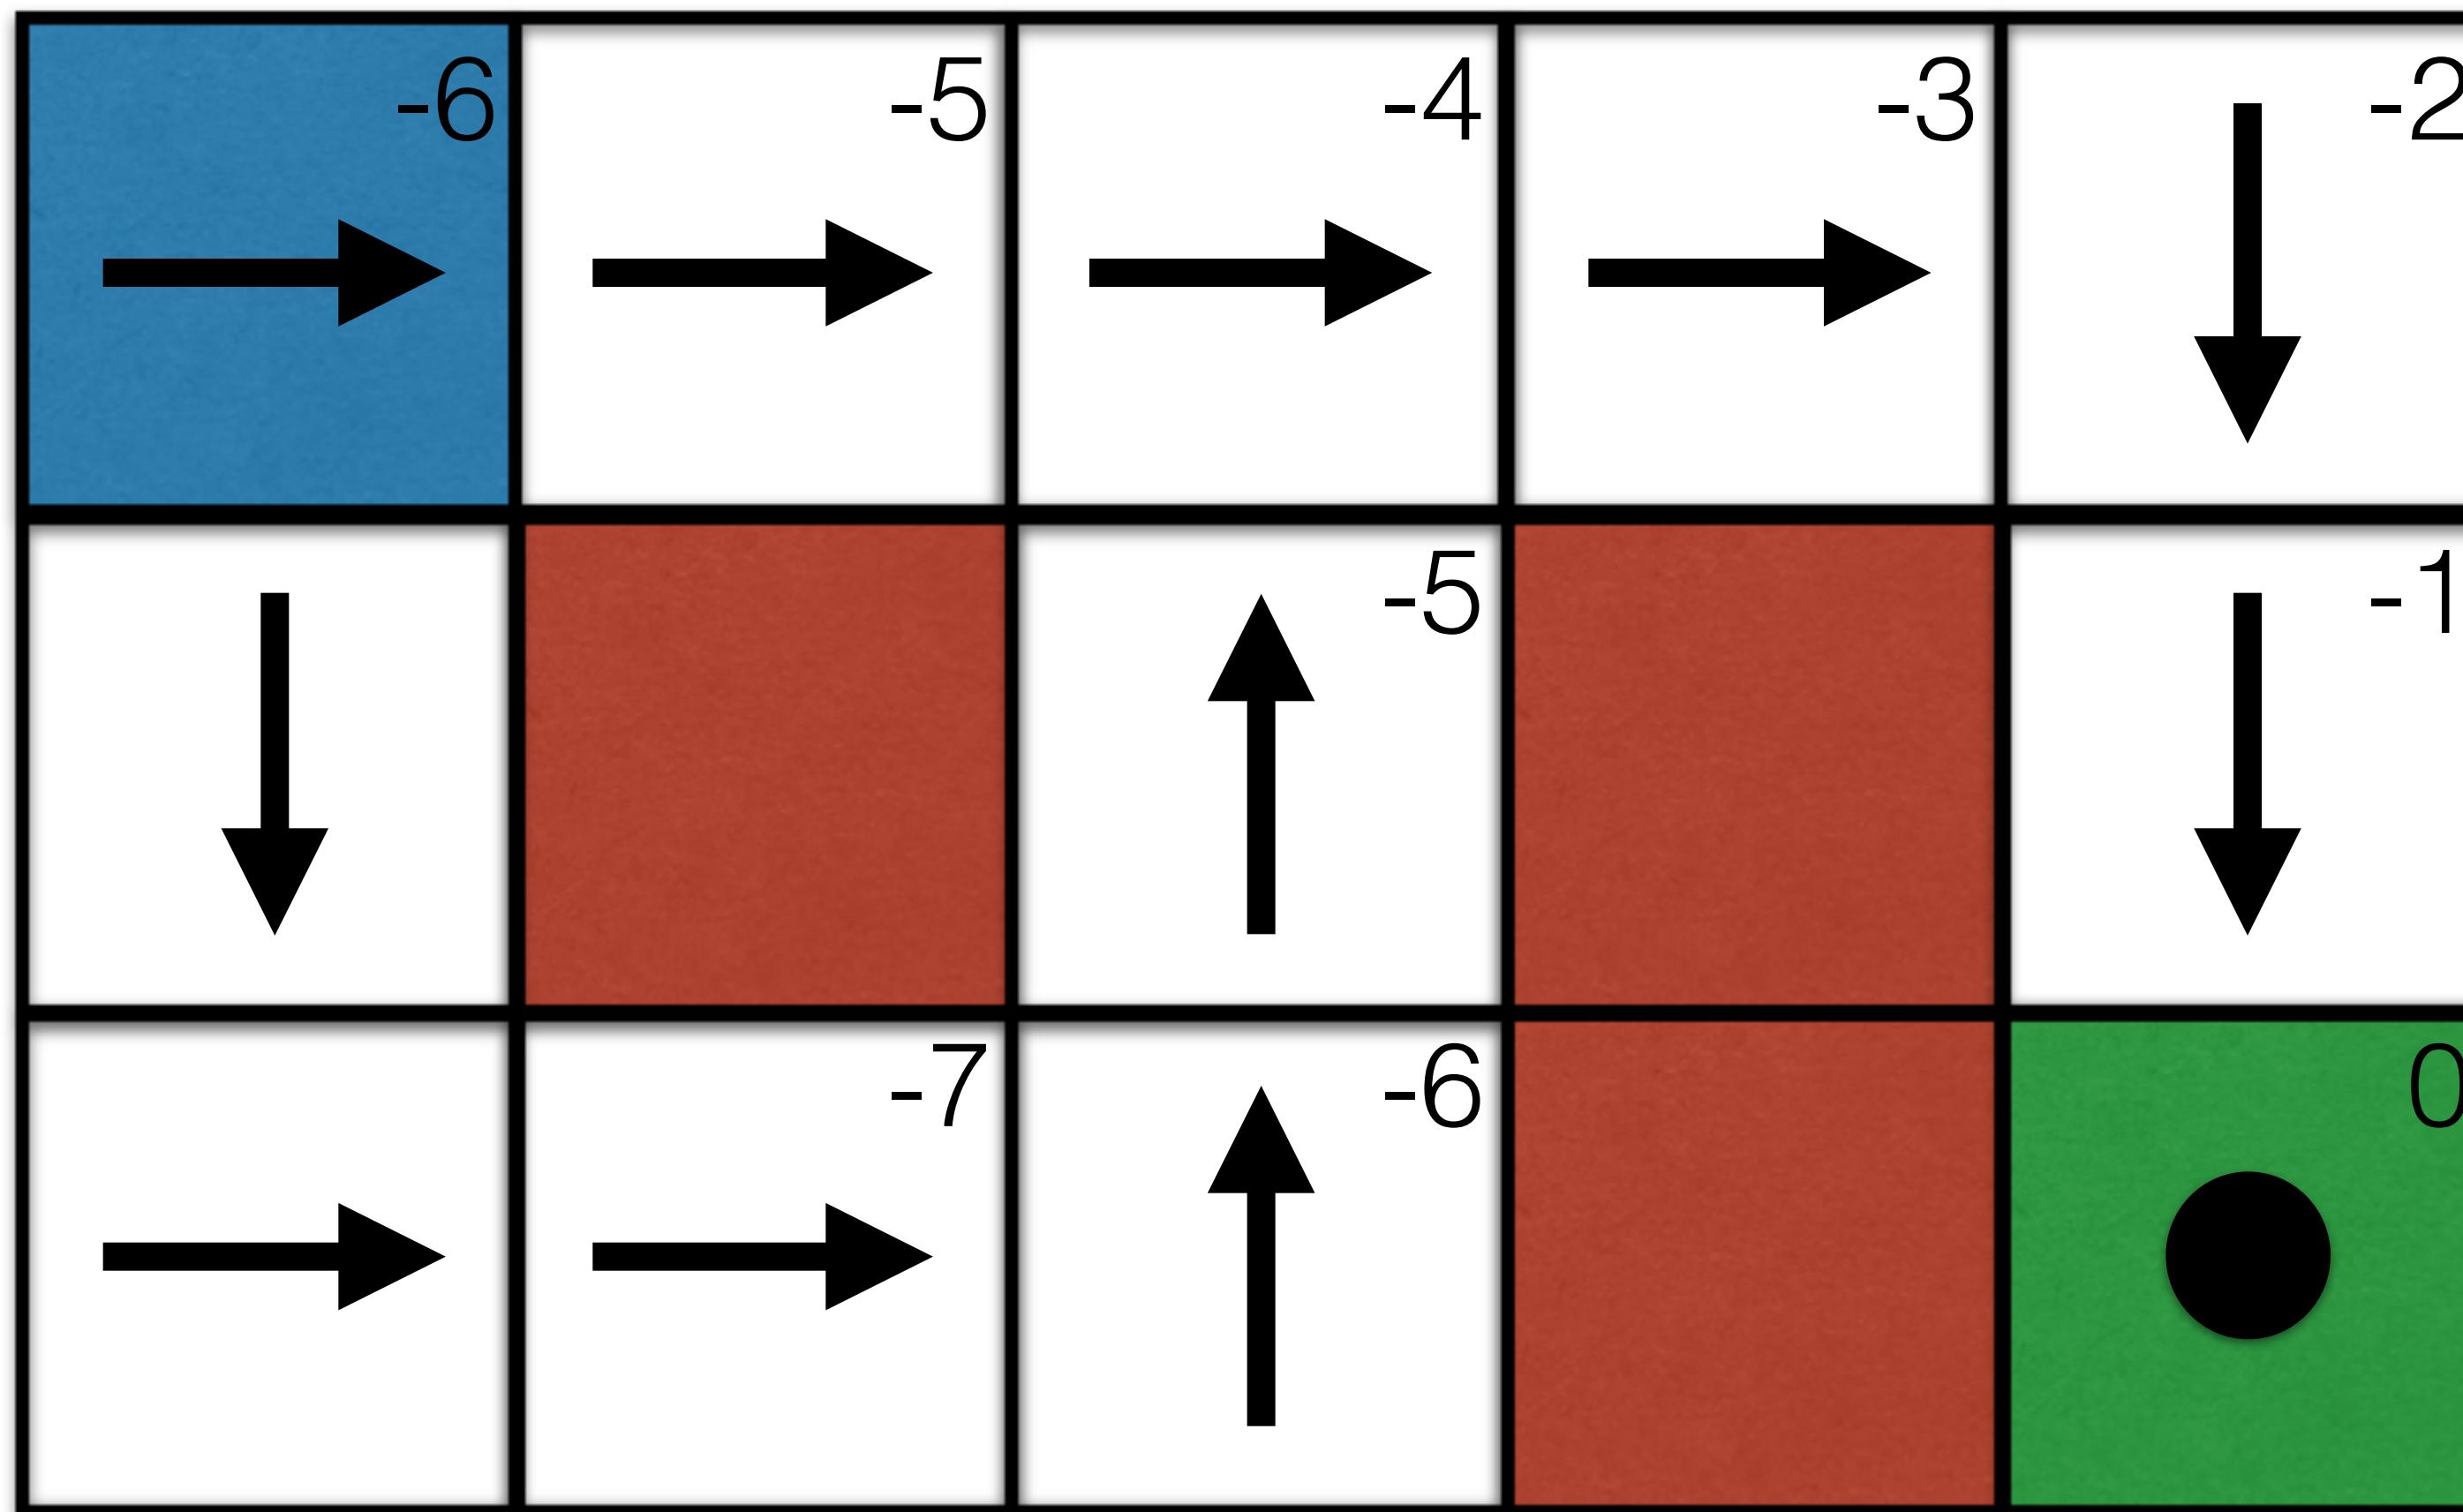
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

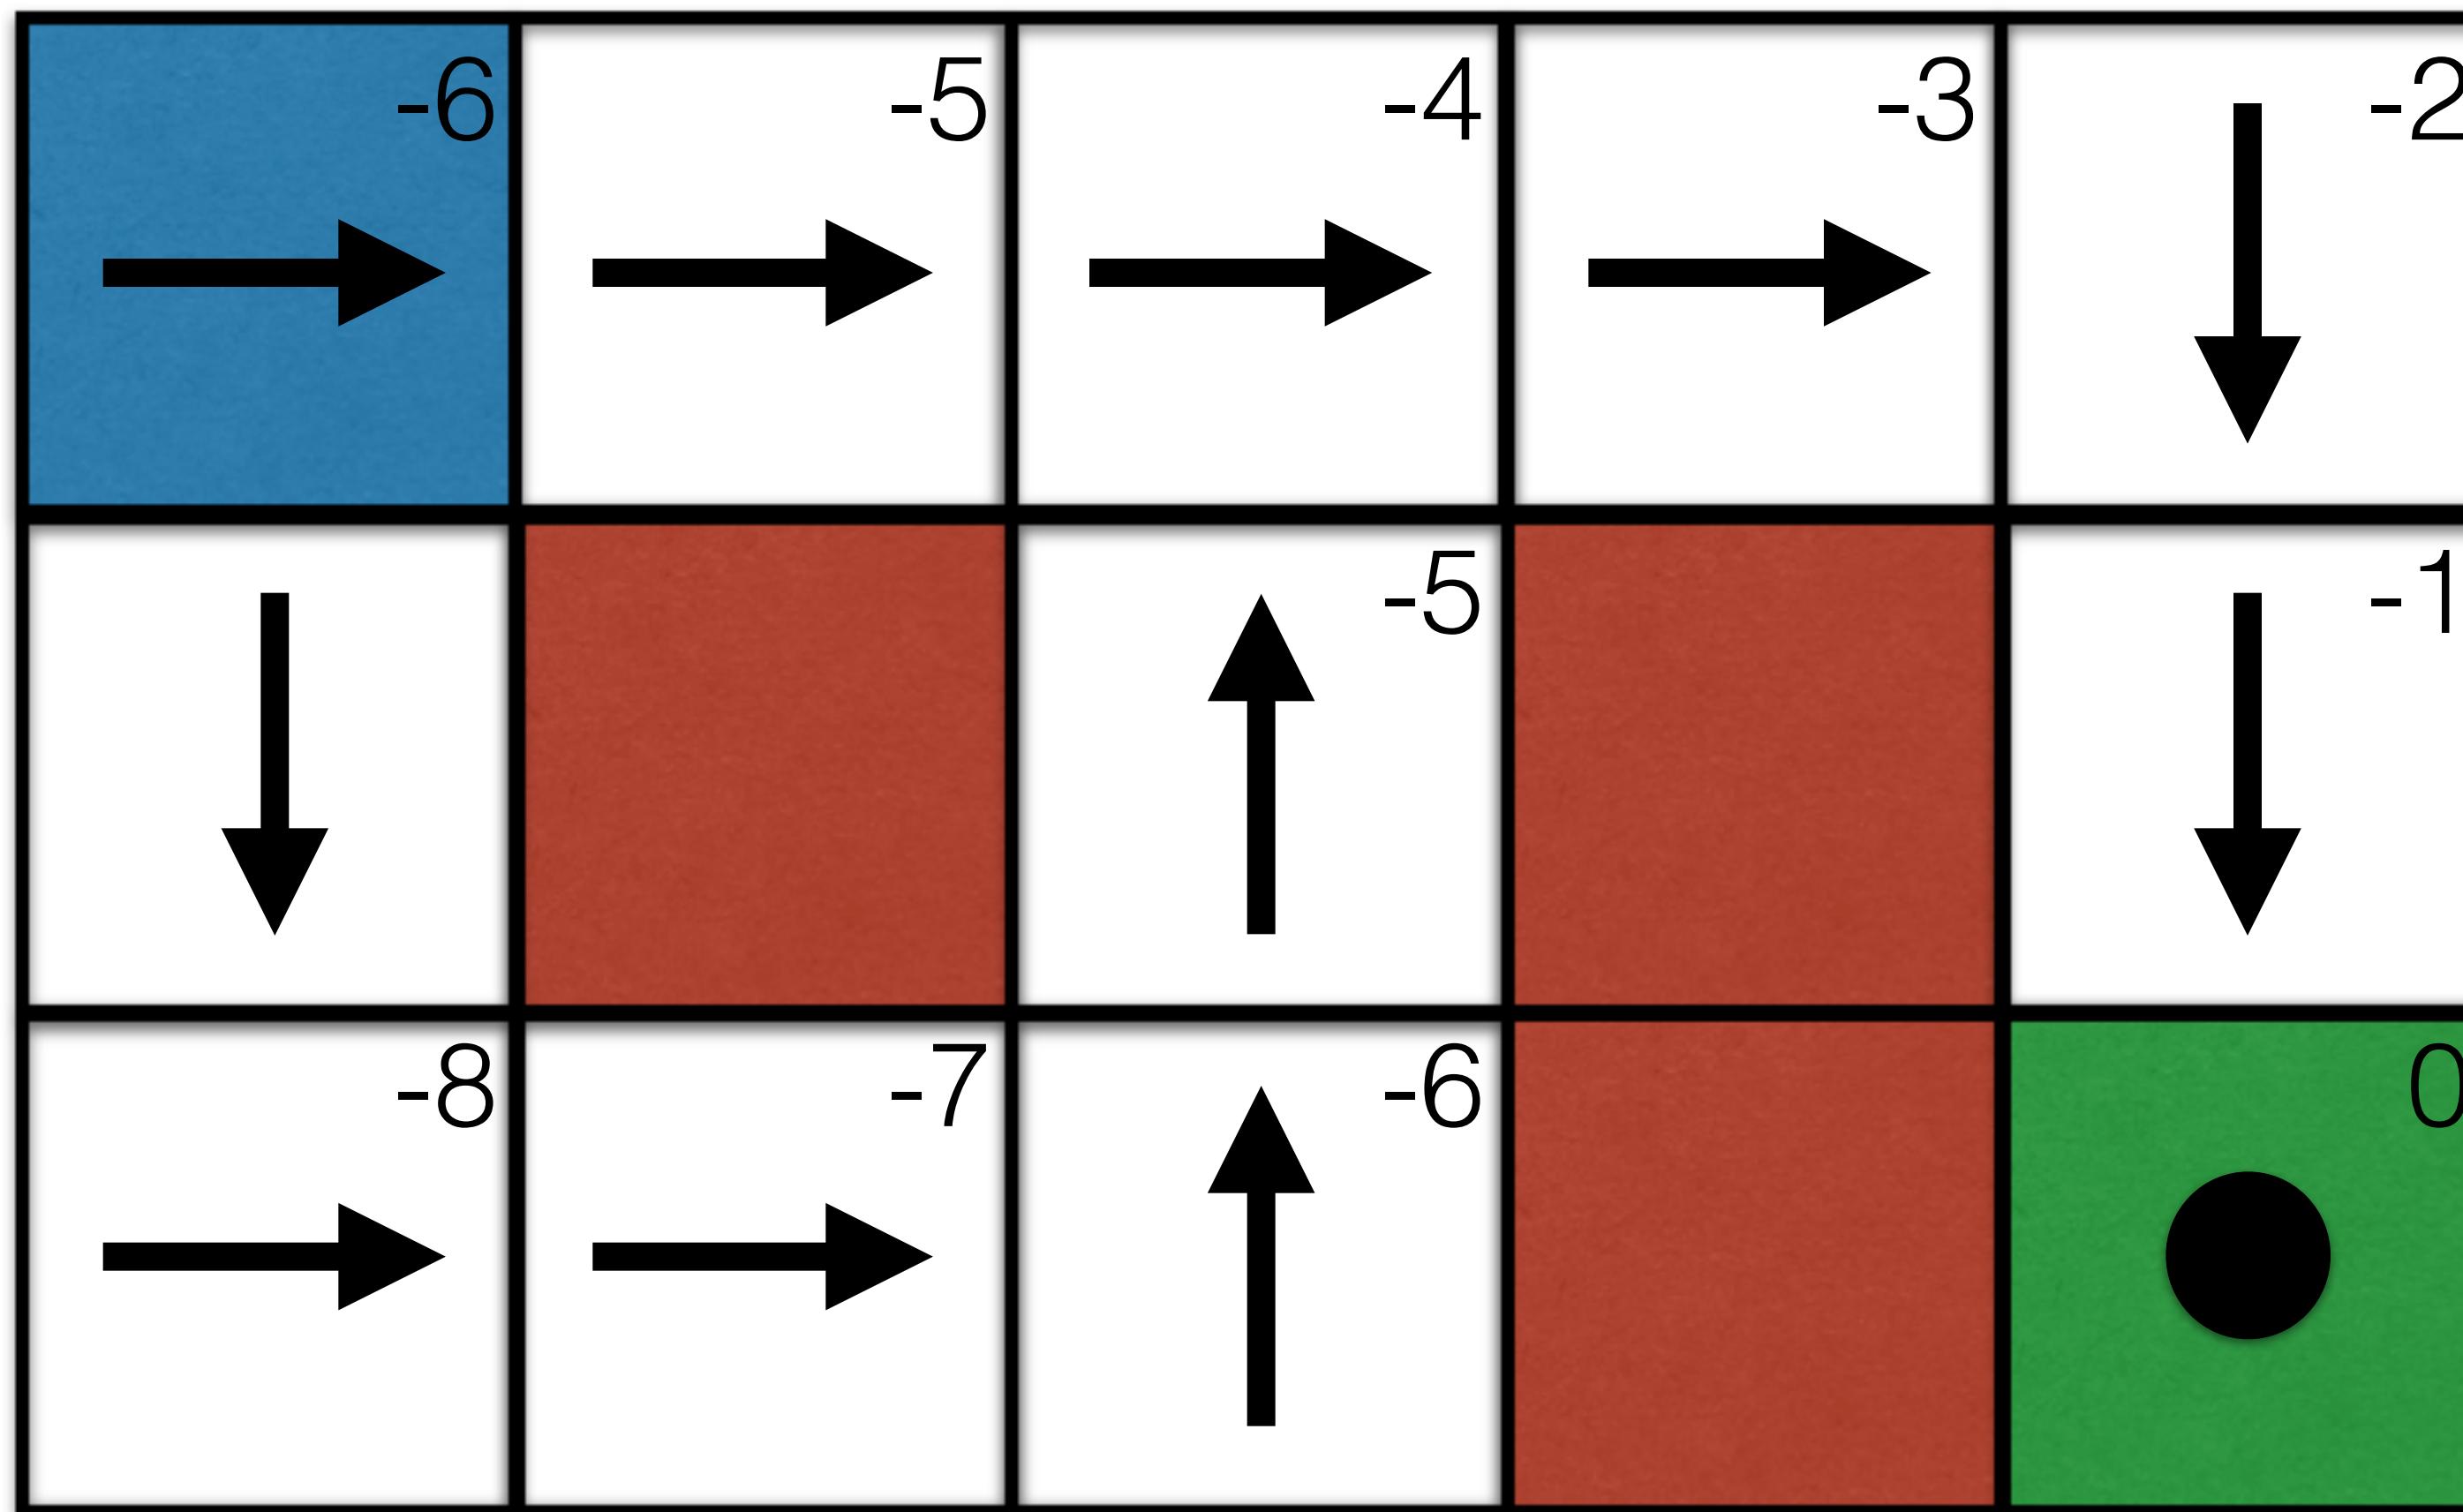
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

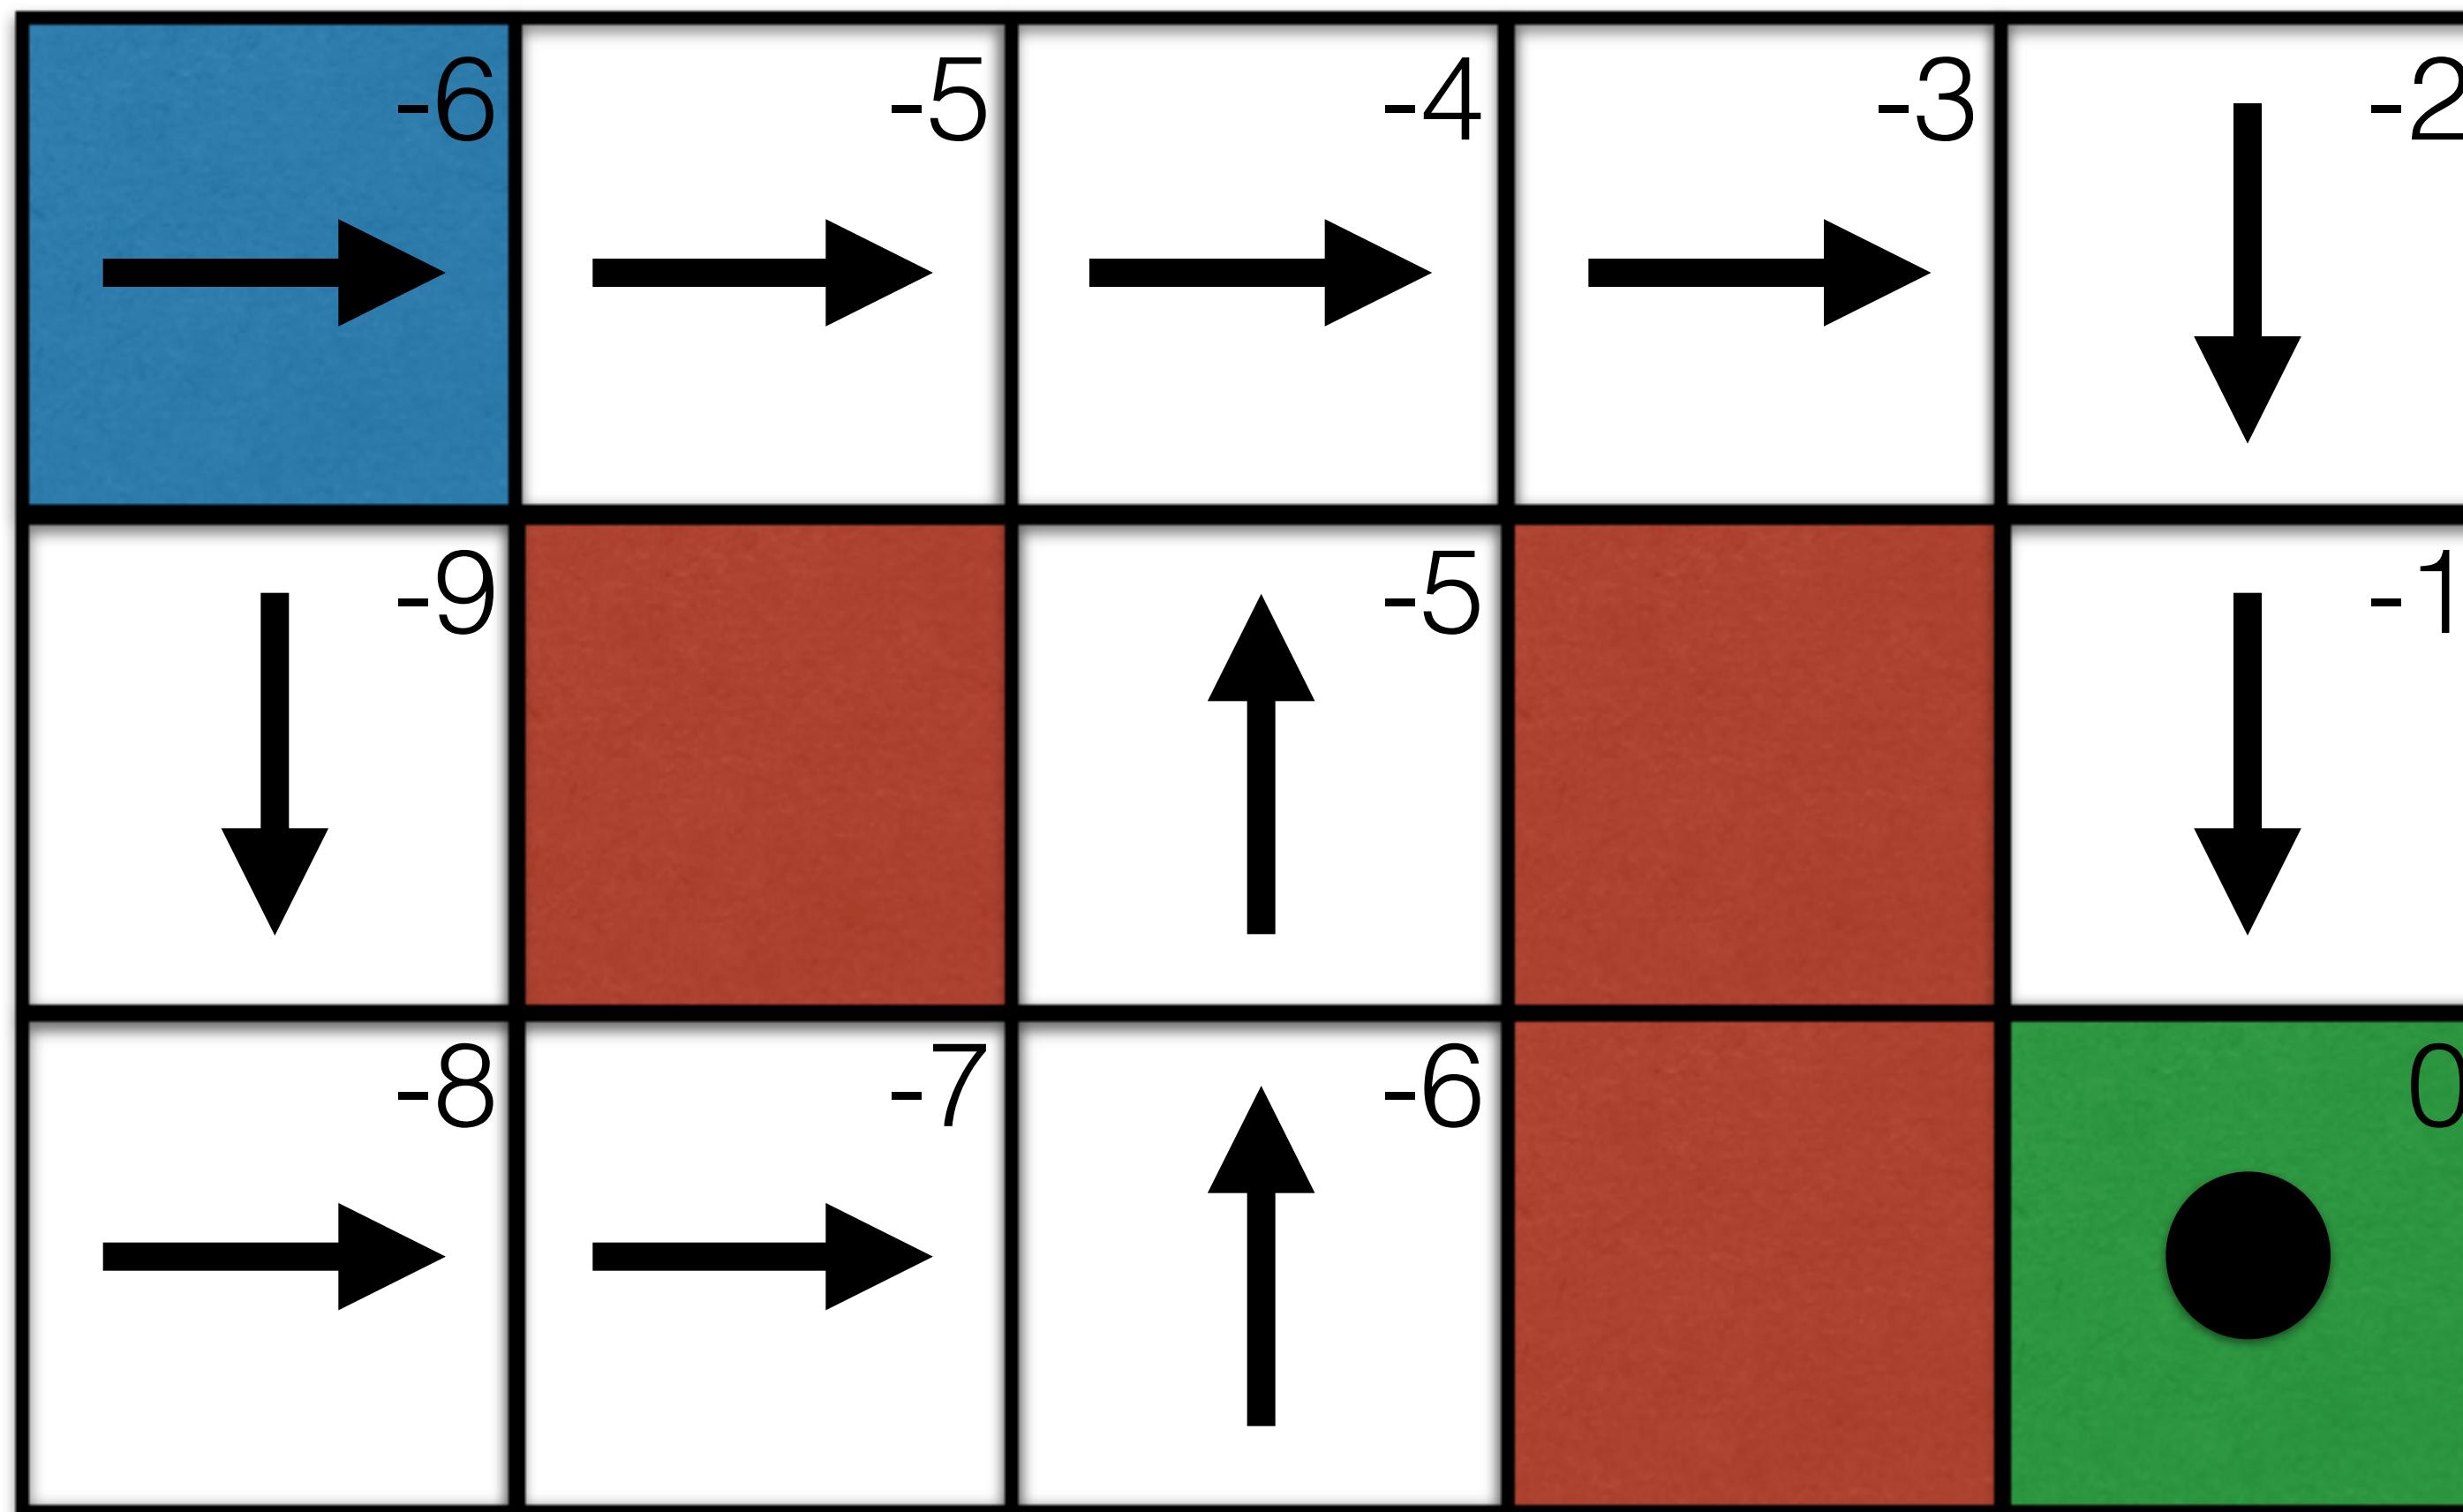
---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Gridworld Value Function Example

---



- Arrows denote the policy  $\pi$ .
- What are the values of the states, assuming no random movements?

# Policy Evaluation

---

# Policy Evaluation

---

- What we just did was *policy evaluation*, determining the value of states given our policy  $\pi$ . Simply put, we figured out how good our policy is.

# Policy Evaluation

---

- What we just did was *policy evaluation*, determining the value of states given our policy  $\pi$ . Simply put, we figured out how good our policy is.
- But remember, what we are really interested in is the optimal policy  $\pi^*$ ! How do we find this?

# Policy Evaluation

---

- What we just did was *policy evaluation*, determining the value of states given our policy  $\pi$ . Simply put, we figured out how good our policy is.
- But remember, what we are really interested in is the optimal policy  $\pi^*$ ! How do we find this?
- We need one more step - *policy iteration*.

# Policy Evaluation

---

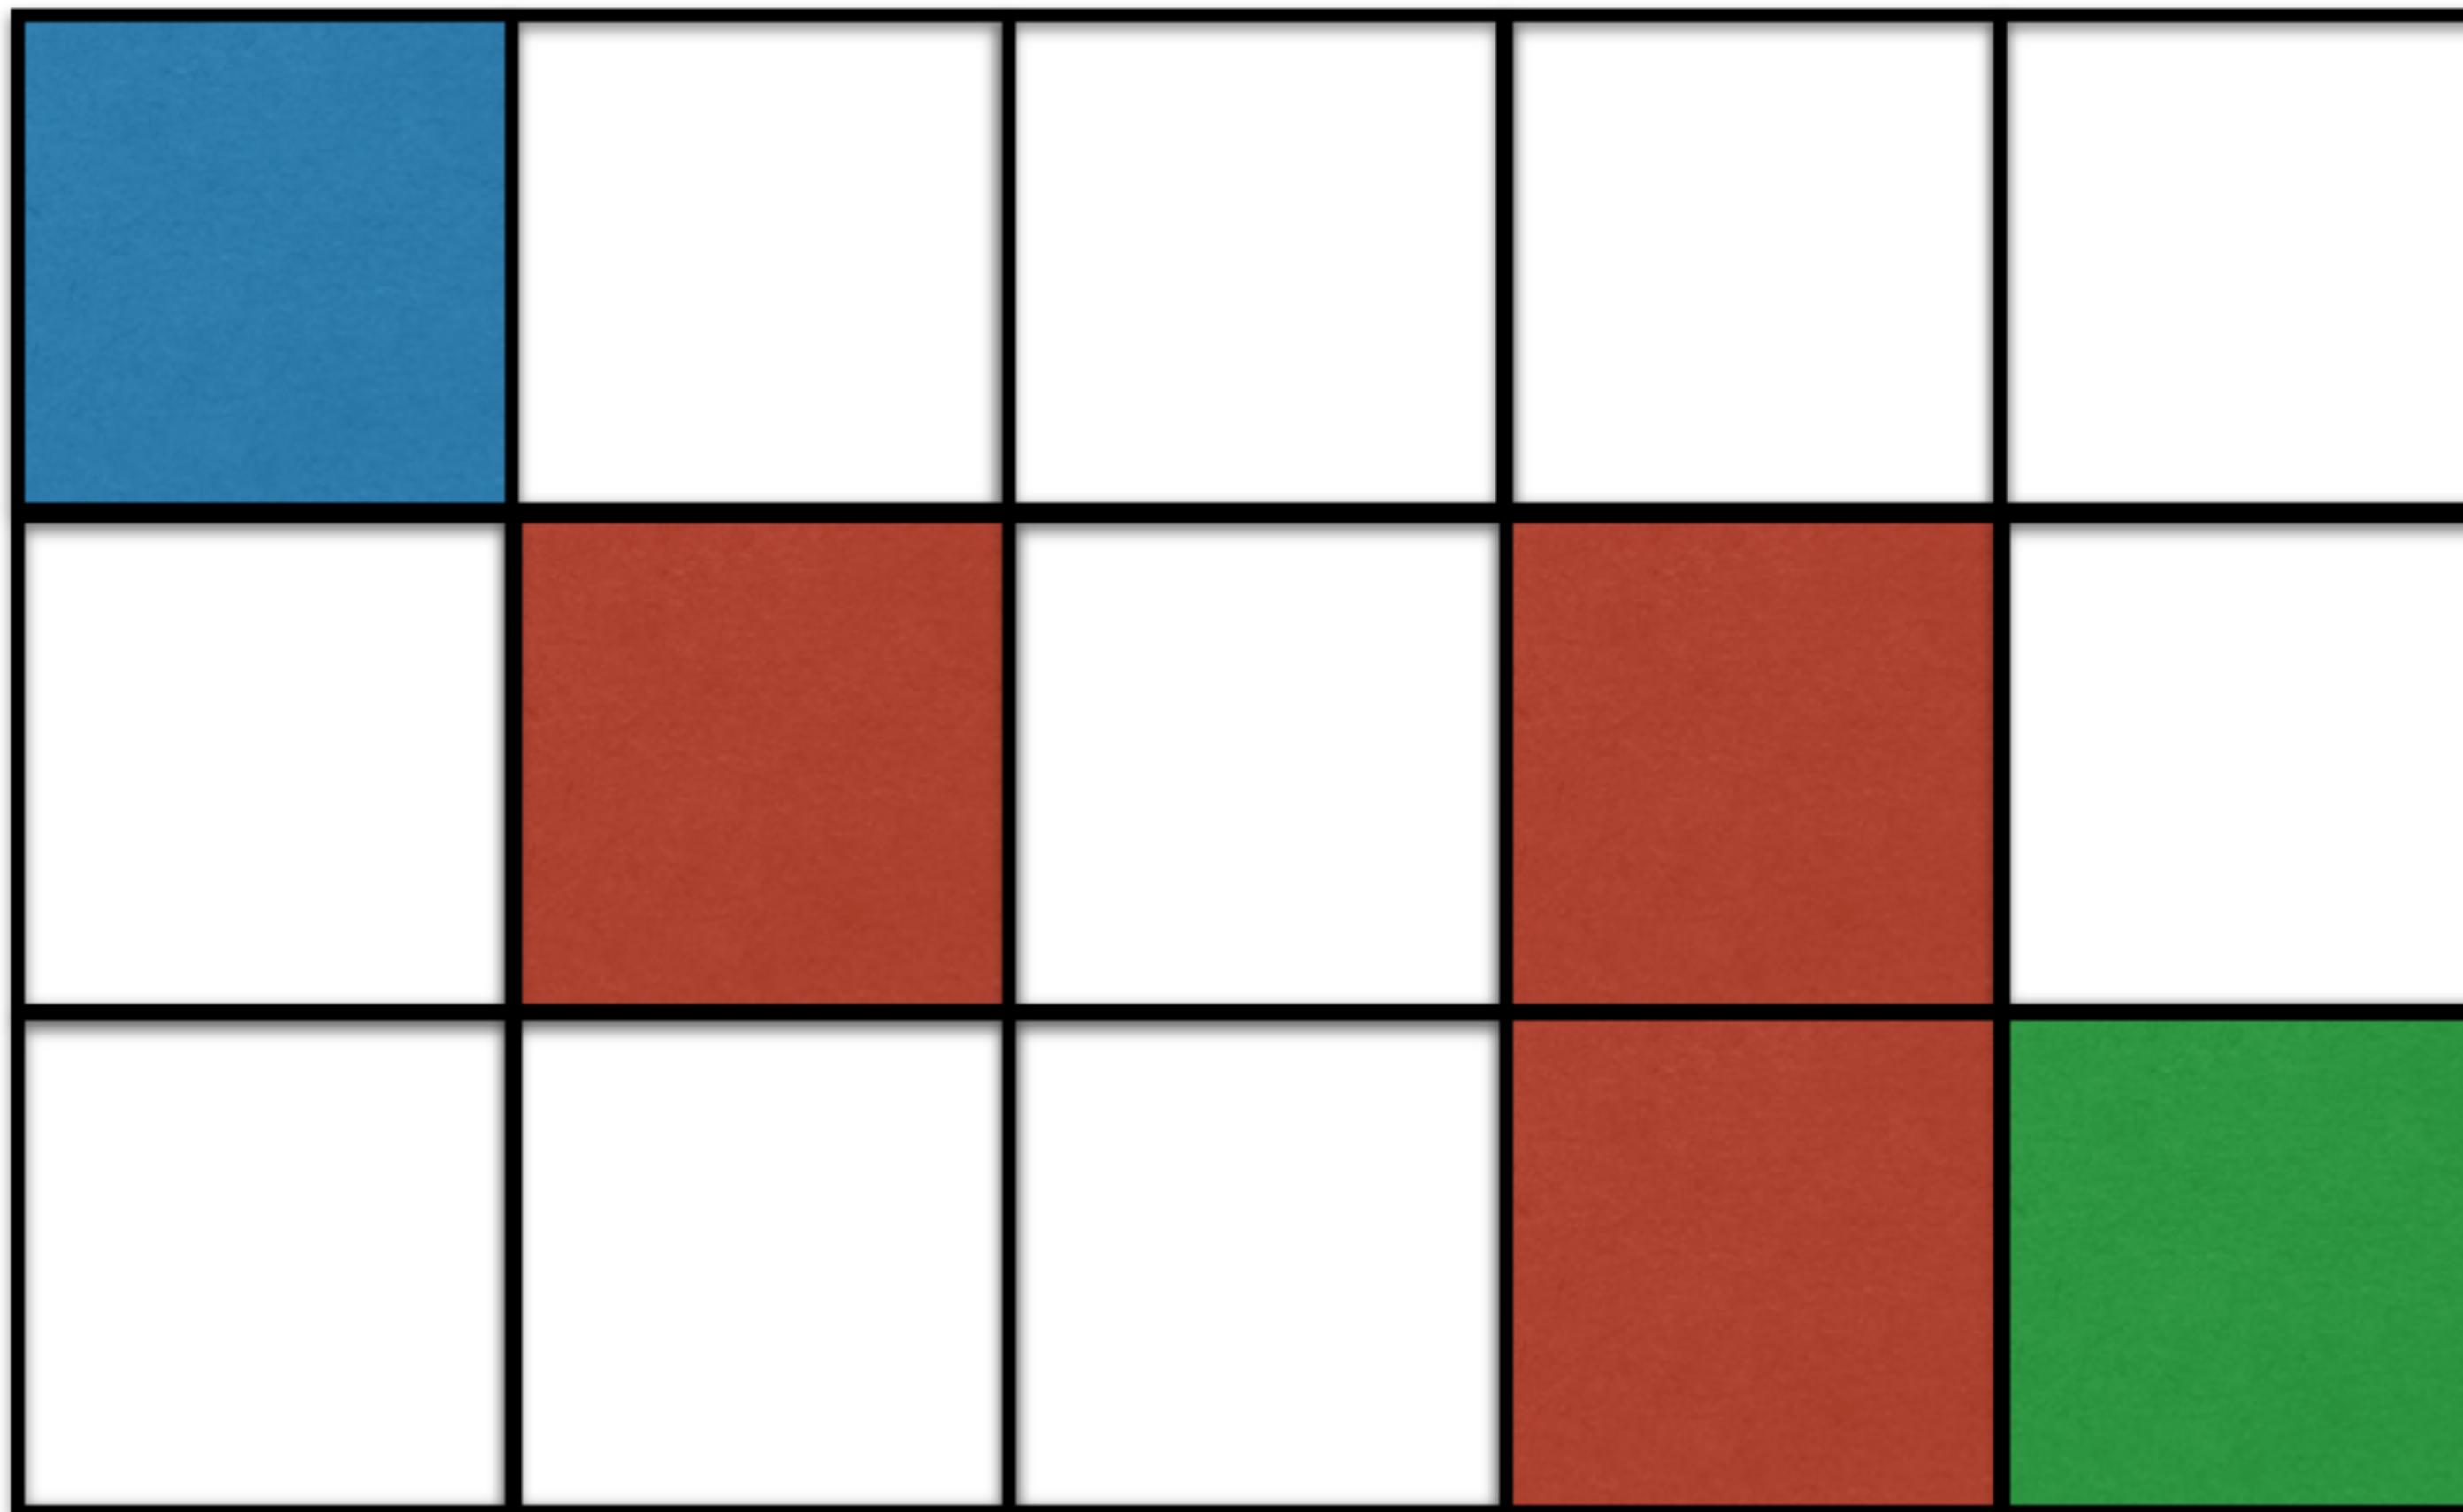
- What we just did was *policy evaluation*, determining the value of states given our policy  $\pi$ . Simply put, we figured out how good our policy is.
- But remember, what we are really interested in is the optimal policy  $\pi^*$ ! How do we find this?
- We need one more step - *policy iteration*.

# Gridworld Policy Iteration Example

---

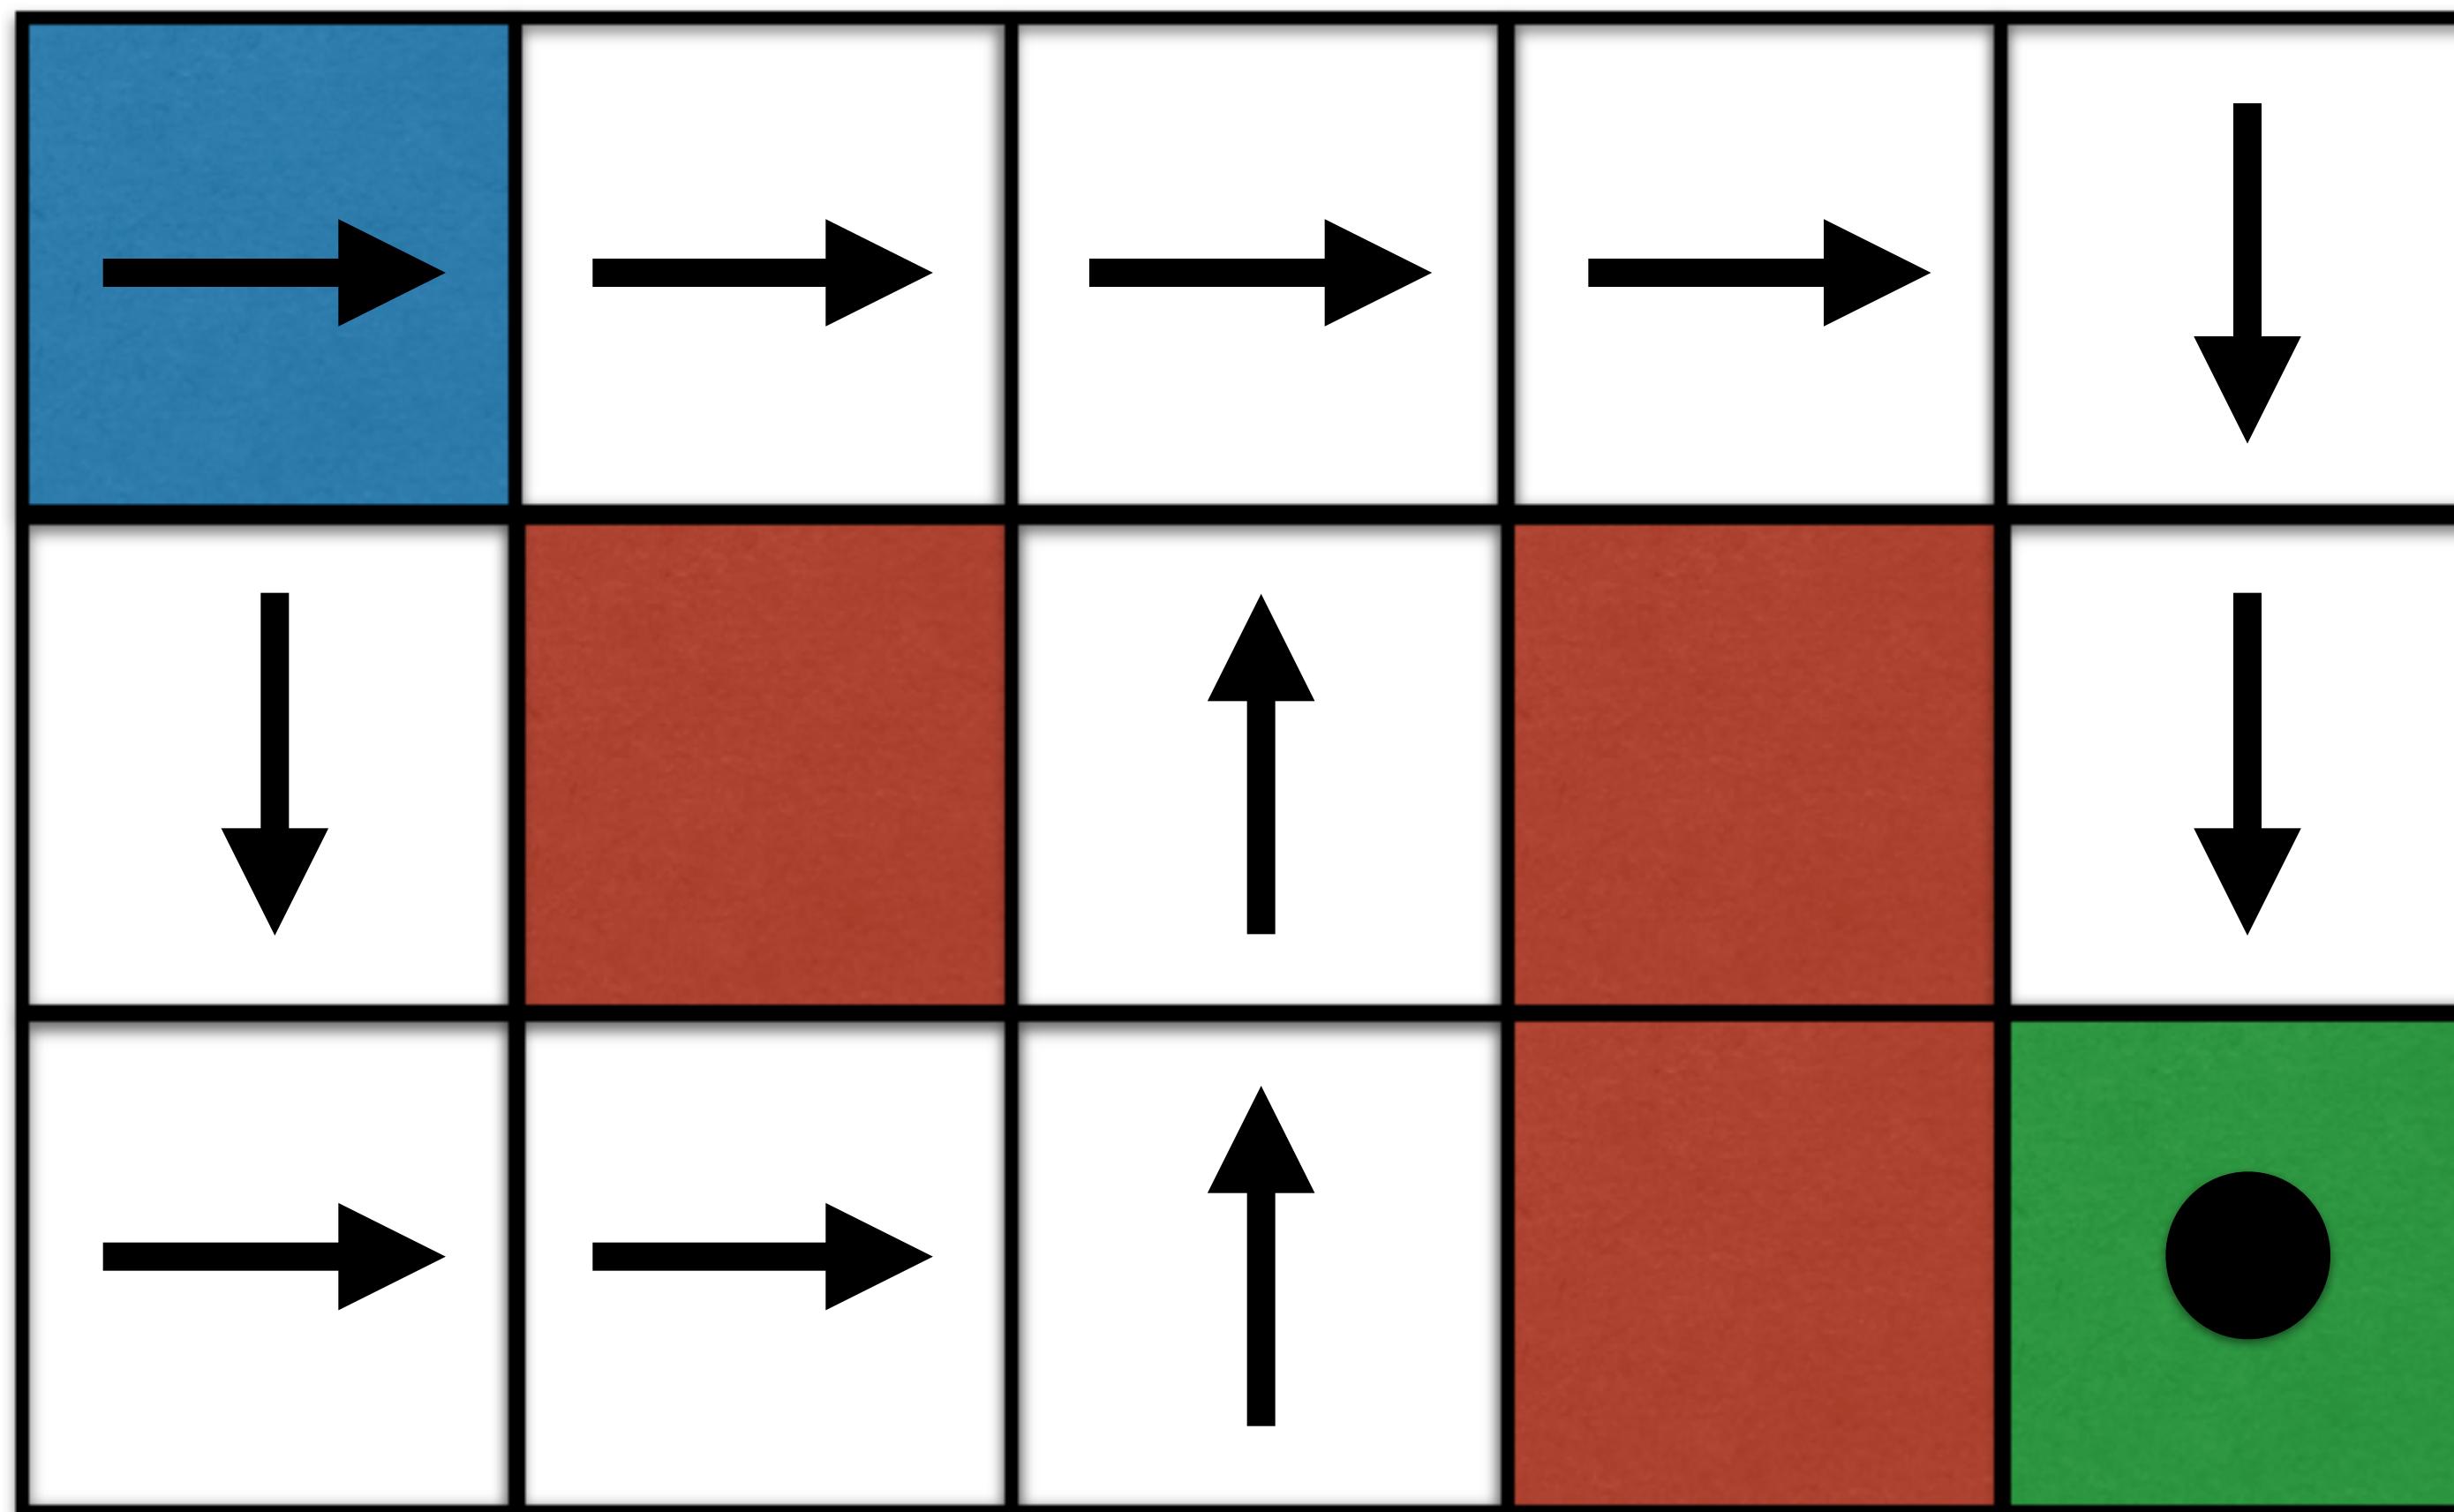
# Gridworld Policy Iteration Example

---



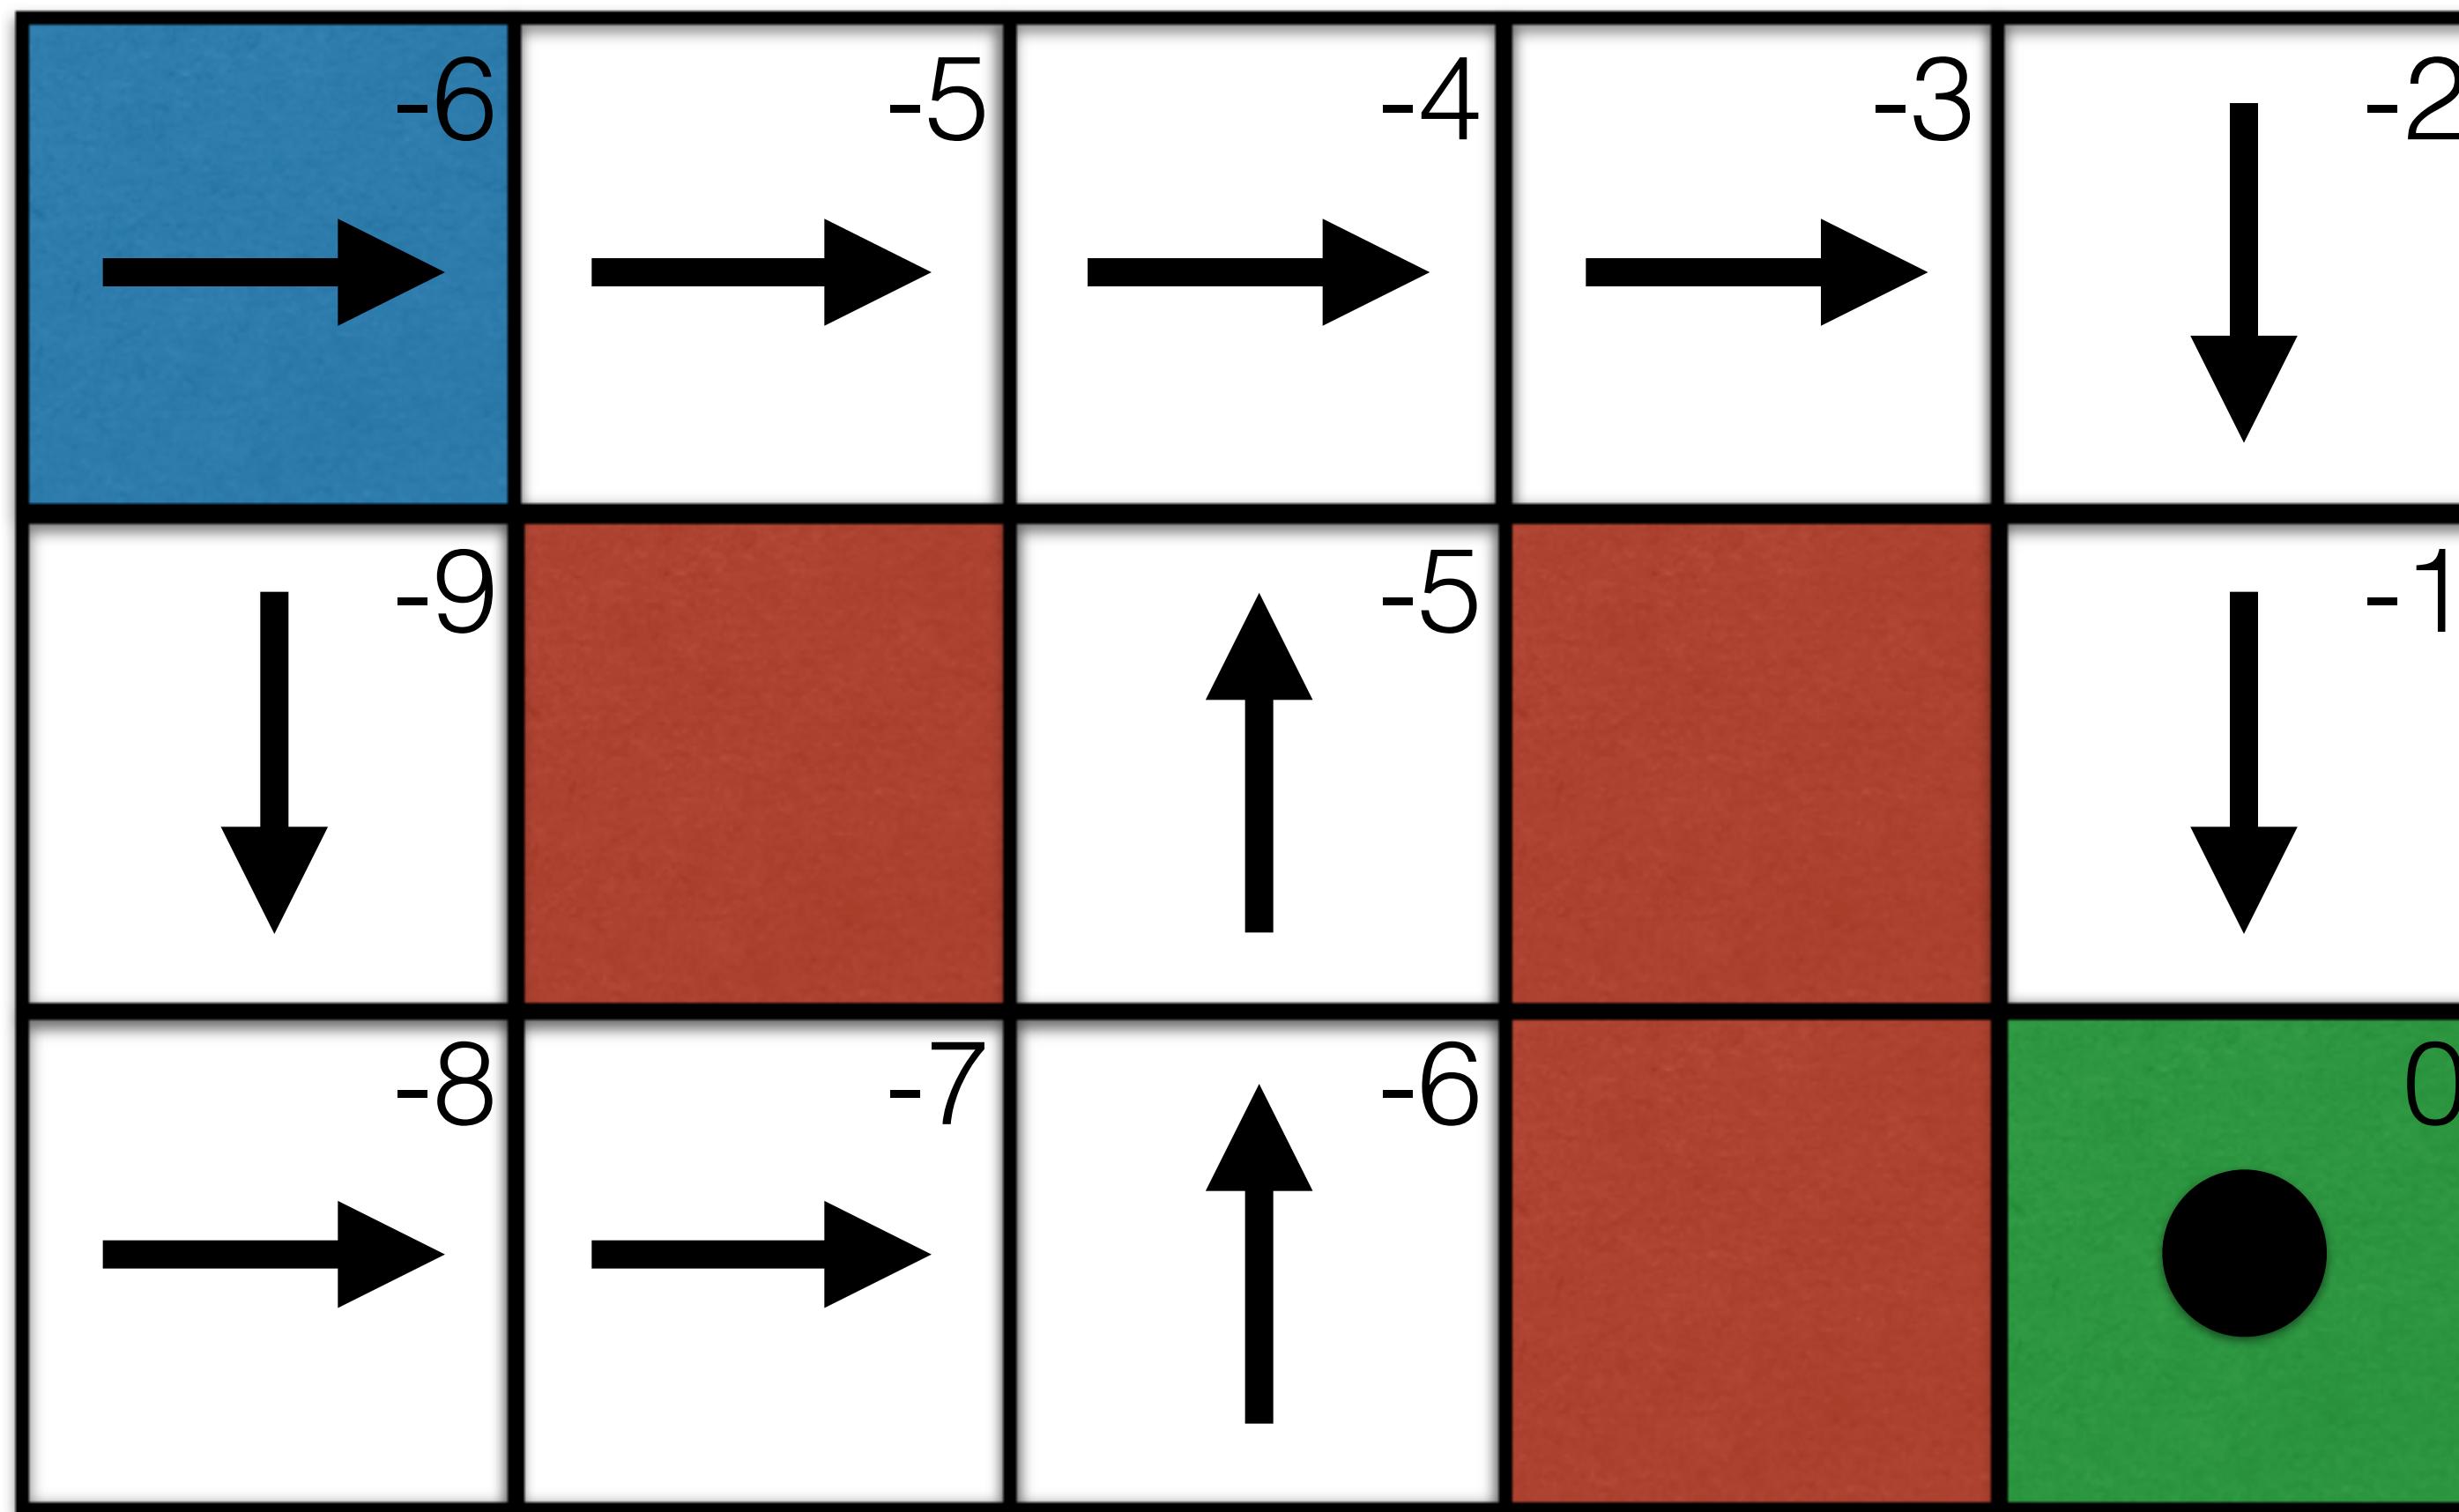
# Gridworld Policy Iteration Example

---



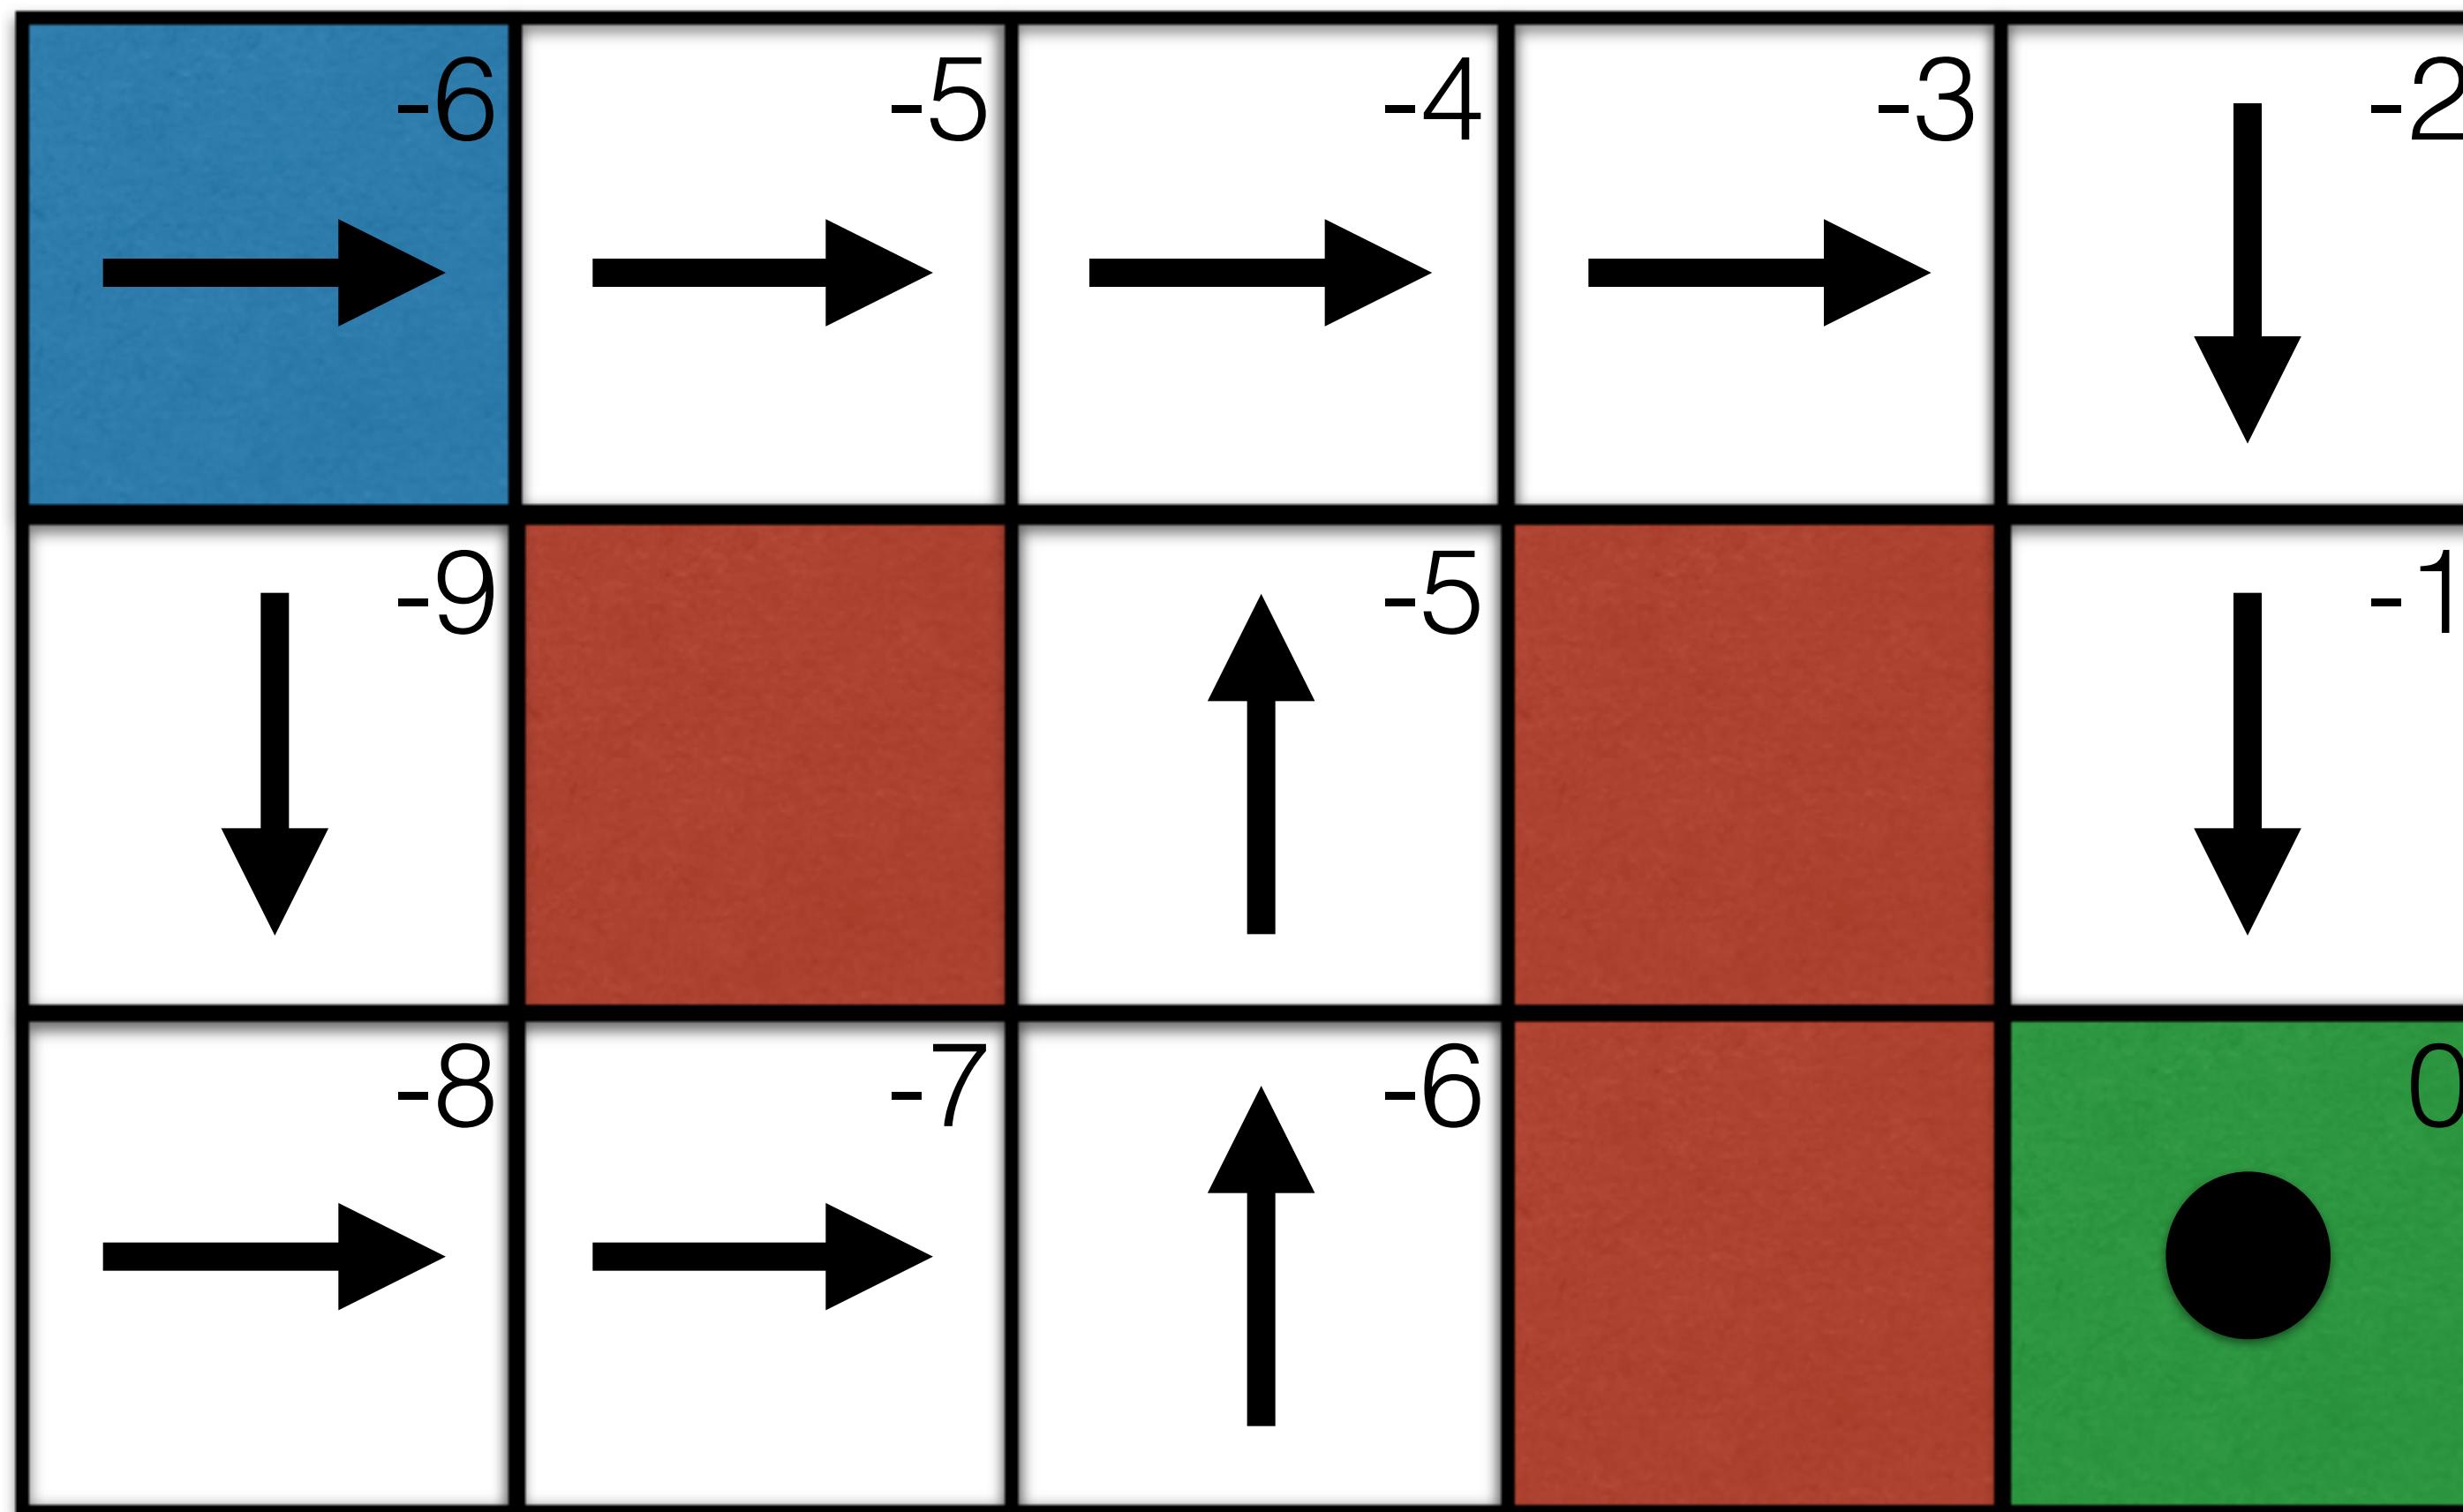
# Gridworld Policy Iteration Example

---



# Gridworld Policy Iteration Example

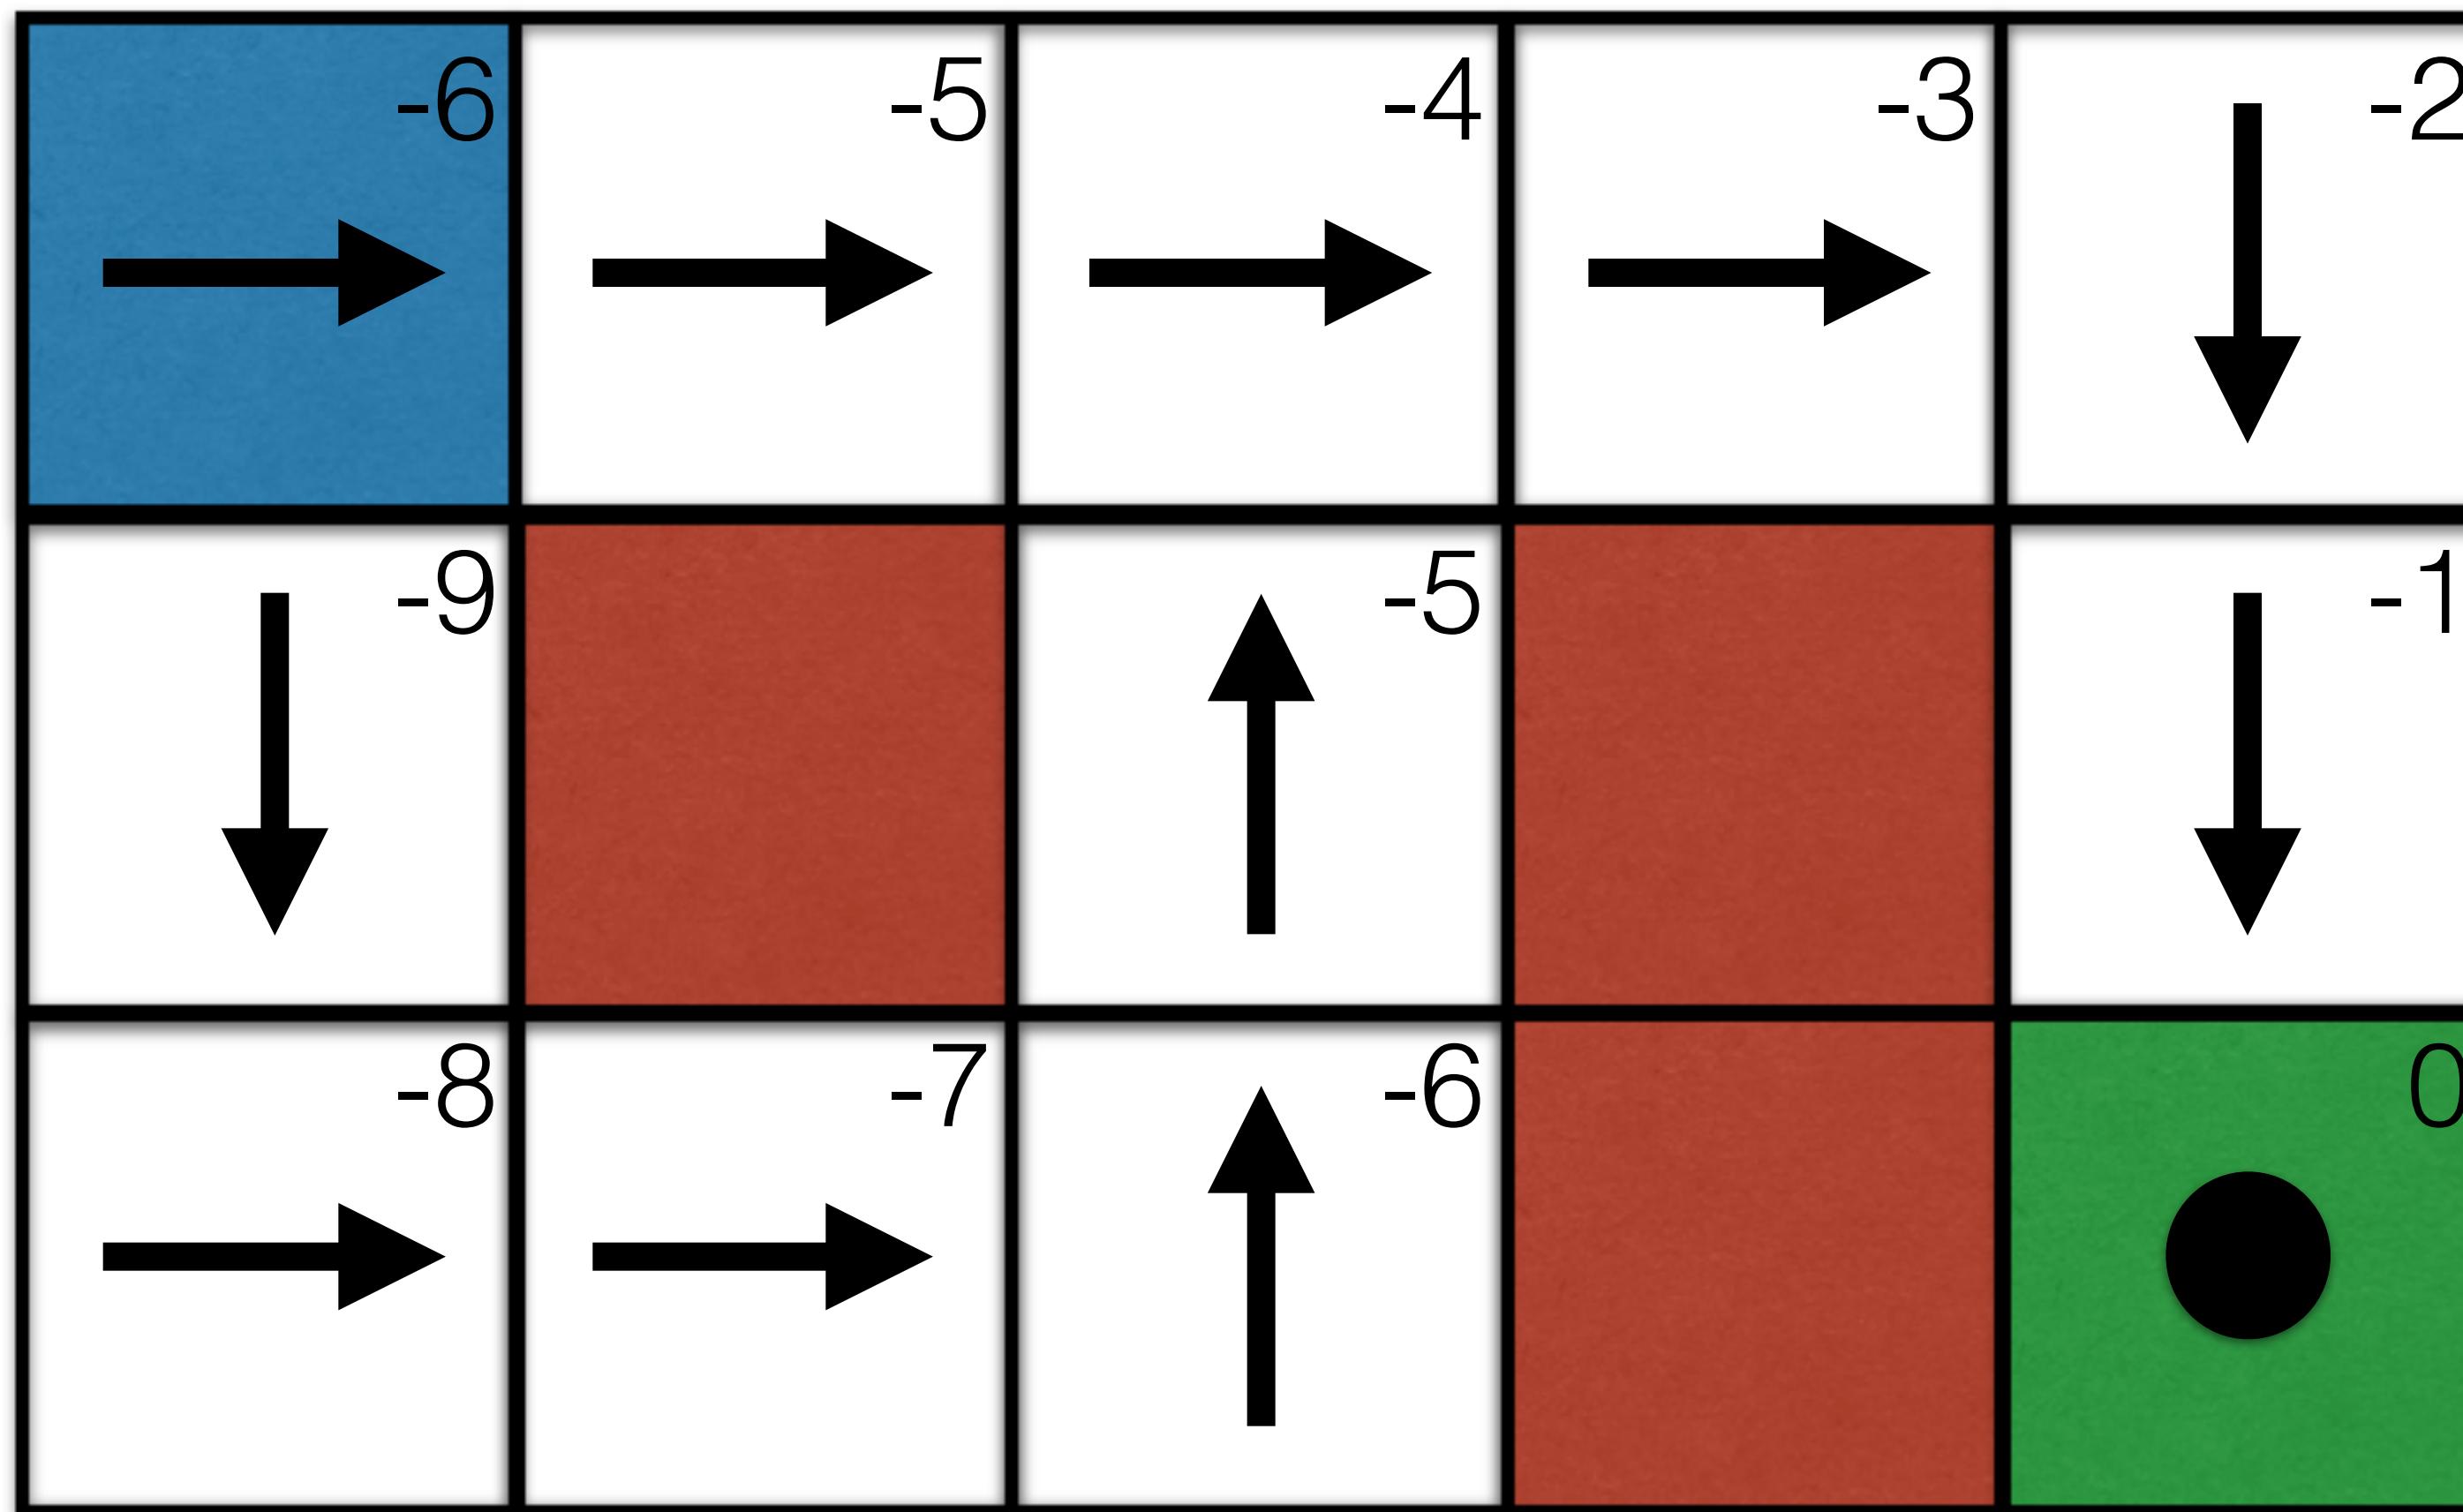
---



- Arrows denote the policy.

# Gridworld Policy Iteration Example

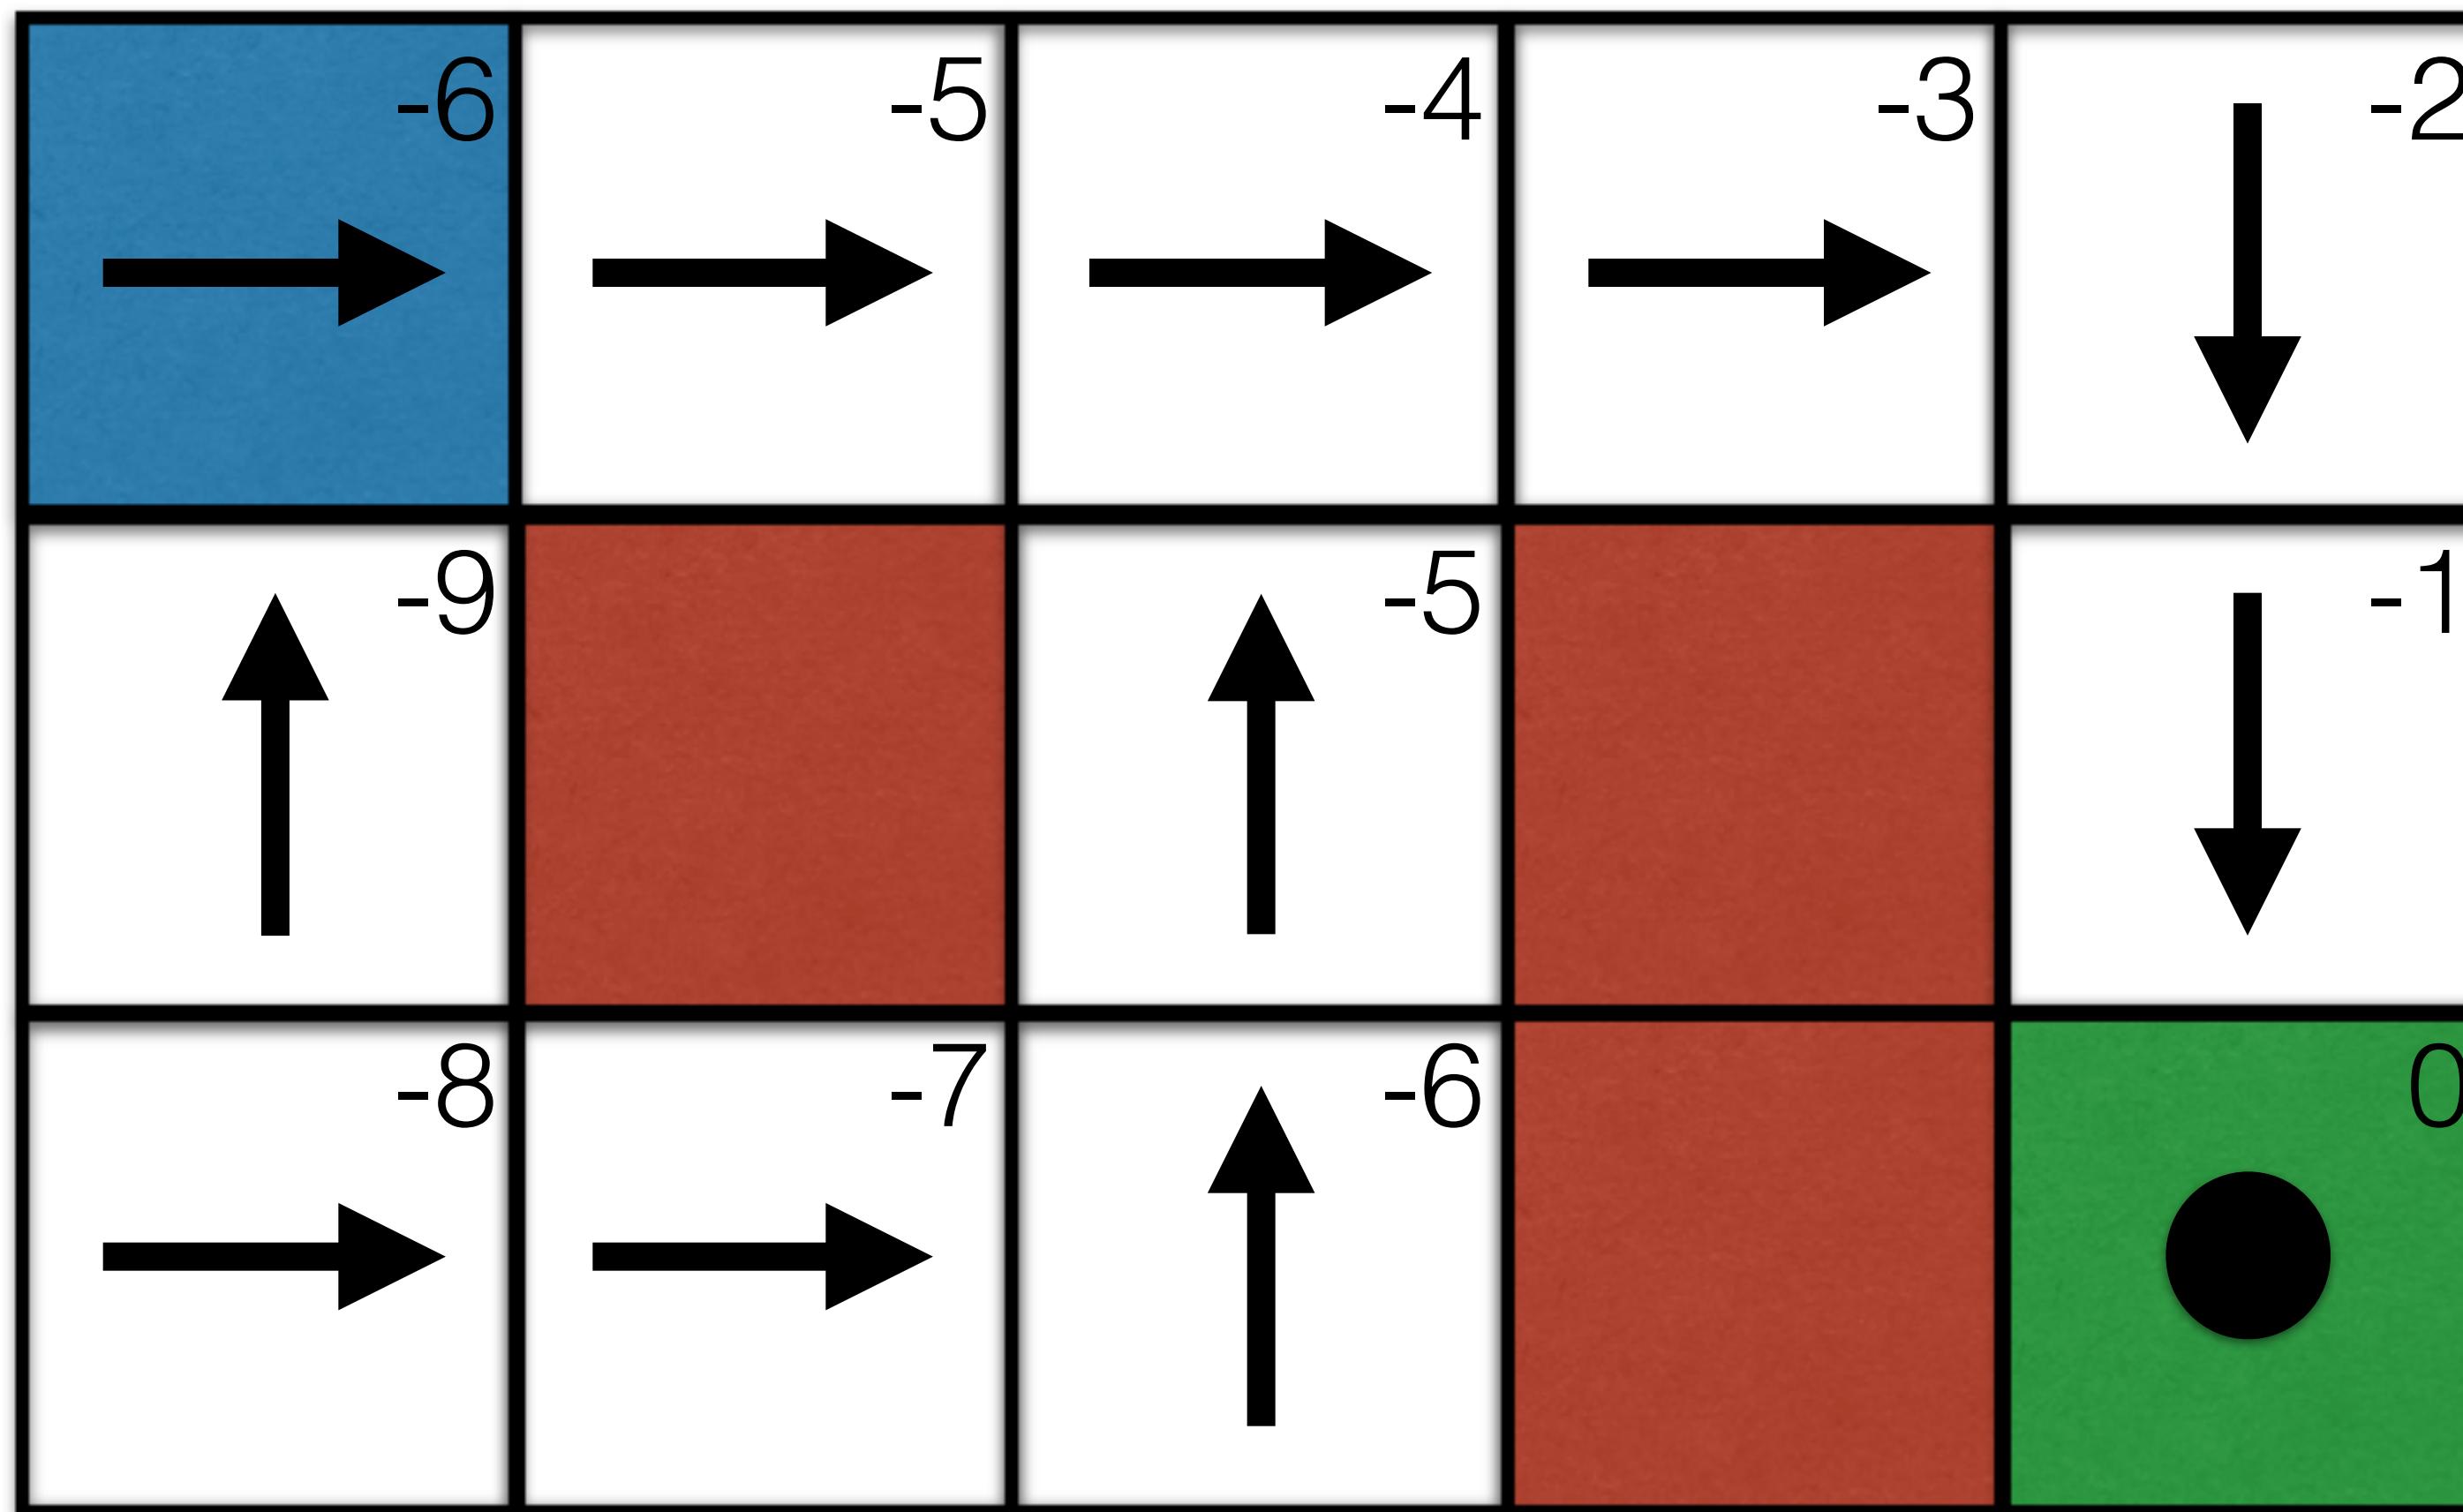
---



- Arrows denote the policy.
- Based on the value function, which action of the current policy should we change?

# Gridworld Policy Iteration Example

---



- Arrows denote the policy.
- Based on the value function, which action of the current policy should we change?

# Policy Iteration

---

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:
  - For every state  $s$ ,  $\pi'$  picks the action that leads to the next state  $s'$  with the highest value.

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:
  - For every state  $s$ ,  $\pi'$  picks the action that leads to the next state  $s'$  with the highest value.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:
  - For every state  $s$ ,  $\pi'$  picks the action that leads to the next state  $s'$  with the highest value.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- This is called *policy iteration*.

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:
  - For every state  $s$ ,  $\pi'$  picks the action that leads to the next state  $s'$  with the highest value.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- This is called *policy iteration*.

```
def new_policy(s):
    return max(actions,
               key=lambda a: sum([P(s, a, n_s) * V(n_s) for n_s in states]))
```

# Policy Iteration

---

- Now that we know  $V(s)$ , we improve our policy  $\pi$  to a new policy  $\pi'$  as follows:
  - For every state  $s$ ,  $\pi'$  picks the action that leads to the next state  $s'$  with the highest value.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- This is called *policy iteration*.

```
def new_policy(s):
    return max(actions,
               key=lambda a: sum([P(s, a, n_s) * V(n_s) for n_s in states]))
```

# Optimal Policy

---

# Optimal Policy

---

- So, to find the optimal policy:

# Optimal Policy

---

- So, to find the optimal policy:
  - Initialize some policy  $\pi$ .

# Optimal Policy

---

- So, to find the optimal policy:
  - Initialize some policy  $\pi$ .
  - Repeat:

# Optimal Policy

---

- So, to find the optimal policy:
  - Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

# Optimal Policy

---

- So, to find the optimal policy:
  - Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- If  $\pi == \pi'$ , return  $\pi$ .

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- If  $\pi == \pi'$ , return  $\pi$ .
  - Otherwise, set  $\pi$  equal to  $\pi'$ .

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- If  $\pi == \pi'$ , return  $\pi$ .
  - Otherwise, set  $\pi$  equal to  $\pi'$ .

- We can prove that this  $\pi$  we return is optimal, i.e.  $\pi == \pi^*$ !  
We won't do the math, though.

# Optimal Policy

---

- So, to find the optimal policy:

- Initialize some policy  $\pi$ .
  - Repeat:
    - Determine  $V(s)$  using *policy evaluation*.

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Find a better policy  $\pi'$  using *policy iteration*.

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- If  $\pi == \pi'$ , return  $\pi$ .
  - Otherwise, set  $\pi$  equal to  $\pi'$ .

- We can prove that this  $\pi$  we return is optimal, i.e.  $\pi == \pi^*$ !  
We won't do the math, though.

# Terminology So Far

---

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*Long-term*)

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*Long-term*)
  - Measures long-term expected reward.

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*Long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*Long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Policy iteration:

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Policy iteration:
  - Using our value function  $V$  to get a better policy  $\pi'$ .

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Policy iteration:
  - Using our value function  $V$  to get a better policy  $\pi'$ .

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Policy iteration:
  - Using our value function  $V$  to get a better policy  $\pi'$ .

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- Basically every RL algorithm is a combination of policy evaluation and policy iteration! Let's take a closer look at two such algorithms.

# Terminology So Far

---

- Reward function ( $R$ ): how good is a state? (*short-term*)
- Value function ( $V$ ): how good is a state? (*long-term*)
  - Measures long-term expected reward.
  - Depends on the current policy  $\pi$ .
- Policy evaluation:
  - Evaluating our current policy  $\pi$  to get a value function  $V$ .

$$V^\pi(s) = R(s) + \sum_{s'} P(s, \pi(s), s') V^\pi(s')$$

- Policy iteration:
  - Using our value function  $V$  to get a better policy  $\pi'$ .

$$\pi'(s) = \arg \max_a \sum_{s'} P(s, a, s') V^\pi(s')$$

- Basically every RL algorithm is a combination of policy evaluation and policy iteration! Let's take a closer look at two such algorithms.

# Value Iteration

---

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

- If  $V$  doesn't change, return the policy  $\pi$  that acts according to the maximum value of  $V$ .

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

- If  $V$  doesn't change, return the policy  $\pi$  that acts according to the maximum value of  $V$ .

$$\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') V(s')$$

# Value Iteration

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

- If  $V$  doesn't change, return the policy  $\pi$  that acts according to the maximum value of  $V$ .

$$\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') V(s')$$

- Again, we can show that this policy is optimal, i.e.  $\pi == \pi^*$ !  
Again, let's not do the math.

# Value Iteration

(demo)

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

- If  $V$  doesn't change, return the policy  $\pi$  that acts according to the maximum value of  $V$ .

$$\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') V(s')$$

- Again, we can show that this policy is optimal, i.e.  $\pi == \pi^*$ !  
Again, let's not do the math.

# Value Iteration

(demo)

---

- Value iteration is an algorithm that *combines* the policy evaluation and policy iteration steps into one single step.
- Repeat:
  - For all states  $s$ , determine  $V(s)$ , and set  $V(s)$  to its maximum possible value.

$$V(s) = R(s) + \max_a \sum_{s'} P(s, a, s') V(s')$$

- If  $V$  doesn't change, return the policy  $\pi$  that acts according to the maximum value of  $V$ .

$$\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') V(s')$$

- Again, we can show that this policy is optimal, i.e.  $\pi == \pi^*$ !  
Again, let's not do the math.

# Questions and Limitations

---

# Questions and Limitations

---

- What if there are way too many states and actions to try?

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.
- What if we don't know how the environment works?

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.
- What if we don't know how the environment works?
  - It's reasonable to think that the agent doesn't completely understand what the next possible states are from taking an action.

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.
- What if we don't know how the environment works?
  - It's reasonable to think that the agent doesn't completely understand what the next possible states are from taking an action.
  - In this case, we have to try different things in order to figure out our environment - this is called *exploration*.

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.
- What if we don't know how the environment works?
  - It's reasonable to think that the agent doesn't completely understand what the next possible states are from taking an action.
  - In this case, we have to try different things in order to figure out our environment - this is called *exploration*.
  - Sometimes, we also want to just keep doing what we know is good - this is called *exploitation*.

# Questions and Limitations

---

- What if there are way too many states and actions to try?
  - We have to find a way to only look at a *subset* of states and actions, and we also need to reasonably *approximate* their values.
- What if we don't know how the environment works?
  - It's reasonable to think that the agent doesn't completely understand what the next possible states are from taking an action.
  - In this case, we have to try different things in order to figure out our environment - this is called *exploration*.
  - Sometimes, we also want to just keep doing what we know is good - this is called *exploitation*.

# Rollout-based Policy Iteration

---

# Rollout-based Policy Iteration

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.

# Rollout-based Policy Iteration

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.

# Rollout-based Policy Iteration

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.
  - The *subset* of states that we look at are the states we encounter during our rollouts.

# Rollout-based Policy Iteration

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.
  - The *subset* of states that we look at are the states we encounter during our rollouts.
  - The *approximation* of the value of the states is by measuring our total reward over of course of our rollout.

# Rollout-based Policy Iteration

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.
  - The *subset* of states that we look at are the states we encounter during our rollouts.
  - The *approximation* of the value of the states is by measuring our total reward over of course of our rollout.
  - We balance *exploration* and *exploitation* by sometimes randomly selecting our action.

# Rollout-based Policy Iteration

(demo)

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.
  - The *subset* of states that we look at are the states we encounter during our rollouts.
  - The *approximation* of the value of the states is by measuring our total reward over of course of our rollout.
  - We balance *exploration* and *exploitation* by sometimes randomly selecting our action.

# Rollout-based Policy Iteration

(demo)

---

- In RBPI, instead of full policy evaluation, we run simulations, or *rollouts*, to approximate our value function.
- This addresses the limitations on the previous slide.
  - The *subset* of states that we look at are the states we encounter during our rollouts.
  - The *approximation* of the value of the states is by measuring our total reward over of course of our rollout.
  - We balance *exploration* and *exploitation* by sometimes randomly selecting our action.

# A Parting Thought

---

# A Parting Thought

---

- We've talked about three of the projects in this class, and how machine learning applies to them.

# A Parting Thought

---

- We've talked about three of the projects in this class, and how machine learning applies to them.
- What about the last one? How might machine learning apply to the Scheme project?

# A Parting Thought

---

- We've talked about three of the projects in this class, and how machine learning applies to them.
- What about the last one? How might machine learning apply to the Scheme project?